

Qubit Electronics

Martina Paola Calvi

June 11, 2026

Contents

Qubit Introduction	16
What is a qubit?	16
Computers and Computation	16
Models of classical computation	16
Examples	16
Models of computation	16
Von Neumann architecture	16
Moore's Law	16
Quantum history	16
Current status	17
Main quantum computer players	17
Data Representation	17
Classical computers	17
Quantum computers	17
Classical parallelism	17
Quantum parallelism	18
DiVincenzo Criteria	18
Requirements of a classical logic element	18
What makes a good qubit?	18
Qubit Figures of Merit	18
Coherence time	18
Gate speed	19
Gate fidelity	19
Bits and Qubits	19
Classical computers	19
Quantum computers	19
X gate	19
Circuits in space vs circuits in time	20
Two-Qubit Gates	20
Classical XOR	20
Quantum CNOT	20
Qubit States	21
Bra-ket notation	21
Measurement of Qubits	21
Measurement process	21
Measurement depends on a basis	21
State collapse	21
Born rule	22
Other common measurement bases	22
Bloch Sphere	22
Examples	22
Unitary Operations and Single-Qubit Gates	22

Single-qubit gates	22
Unitary matrices	23
Common Single-Qubit Gates	23
Pauli gates	23
Pauli matrix properties	23
Rotation Gates	24
Hadamard Gate	24
Action	24
Role	24
Phase Gate	24
Action	25
Role	25
Applying Multiple Gates	25
Density Matrix	25
Why it is needed	25
Pure state	25
General properties	25
Matrix form	26
Mixed state	26
Pure vs mixed states	26
Superconducting Qubits	26
1. Superconductivity Review	26
Cooper Pairs	26
The Macroscopic Wave Function	27
Key Electromagnetic Properties	27
The Meissner Effect	27
2. The Josephson Junction	28
Physical Definition	28
Coupled Superconductors Derivation	28
3. The Constitutive Josephson Relations	29
1. The Current-Phase Relation (First Josephson Relation)	29
2. The Voltage-Phase Relation (Second Josephson Relation)	29
4. Josephson Inductance and Nonlinearity	29
Physical Significance	30
5. Josephson Energy	30
6. Artificial Atoms vs. Natural Atoms	31
The Energy Spectrum Problem	31
Cryogenic Initialization	31
7. Harmonic Oscillator Quantization	31
8. Superconducting Qubit Hamiltonian	32
9. Charge Offset and Dephasing Noise	33
The Charge Noise Problem	33
10. The Transmon Regime	33
Suppression of Charge Dispersion	34
Taylor Series Expansion of the Potential	34

Dressed Multi-Level Energy Distribution	35
11. Architectural Families of Superconducting Qubits	35
1. The Transmon	35
2. The Flux Qubit	35
3. The Fluxonium	36
12. Core Mathematical Reference Sheet	36
Circuit QED	36
1. From Atomic Physics to Quantum Optics	36
Two-Level System Driven by a Field	36
Single-Photon Coupling and Vacuum Fluctuations	37
2. Cavity QED Foundations	37
The Fabry–Pérot Resonator	37
Quantization of the Electromagnetic Field	37
3. From Cavity QED to Circuit QED	38
Why Shift to Circuit QED?	38
Key Advantages of Circuit QED	38
Typical Implementation Layout	39
4. Modeling the Resonator as a Single Mode	39
One-Mode Approximation	39
Lumped-Element LC Analogue	39
Zero-Point Fluctuations	39
5. Total System Hamiltonian	40
Resonator Hamiltonian (H_r)	40
Transmon Qubit Hamiltonian (H_q)	40
Total Un-Truncated Expression	41
6. Two-Level System Approximation	41
Qubit Truncation	41
The Interaction Term (H_{int})	41
7. The Jaynes–Cummings Model	42
The Rotating Wave Approximation (RWA)	42
Physical Mechanics	42
8. The Dispersive Regime	42
Detuning Condition	42
Consequence: Dressed Frequencies	43
Physical Interpretation	43
9. Dispersive Qubit Readout	43
The Readout Principle	43
High-Fidelity Readout Parameters	44
Operational Parameters Table	44
Control vs. Readout Separation	44
10. Towards Xmon Qubits	44
Evolution of the Transmon	44
The Xmon Architecture	45
Magnetic Flux Tunability	45
Compact Resonator Layouts	45

11. Core Mathematical Reference Sheet	46
Qubit control	46
1. Deep Dive into Readout in cQED	46
The Jaynes-Cummings Hamiltonian	46
The Dispersive Regime ($\Delta \gg g, \kappa$)	46
2. The Purcell Effect and Purcell Filters	47
The Problem: Purcell Decay	47
The Solution: The Purcell Filter	47
3. Controlling the Qubit	47
Method A: Microwave Drive	48
Method B: Flux Control	48
4. Summary Table: Readout vs. Drive	48
5. Mathematical Derivation of Microwave Control	48
Step 1: Linearizing the Circuit Hamiltonian	48
Step 2: Quantization and Truncation	49
Truncation to a Two-Level System	49
6. Moving to the Rotating Frame	50
7. IQ Modulation and the Rotating Wave Approximation	50
Rotating Wave Approximation	51
X and Y Gates	51
Rotation Angle	51
8. Hardware: I/Q Modulation	51
Components	51
Frequency Multiplexing	52
9. What Is a Virtual Z-Gate?	52
10. Mathematics of Virtual Z-Gates	52
Example: Turning an X-Gate into a Y-Gate	52
Why Virtual Z-Gates Are Useful	53
Qubit Characterization	53
1. Closed vs. Open Quantum Systems	53
2. Types of Noise: Systematic vs. Stochastic	53
3. The Bloch Sphere Representation	54
Geometric breakdown	54
Reference frames	54
4. The Density Matrix and Real-World States	54
5. Decoherence	55
6. Longitudinal Relaxation (T_1)	55
Physics	55
How T_1 is measured	56
7. Pure Dephasing (T_ϕ)	56
Physics	56
8. Transverse Relaxation (T_2)	57
Why $\Gamma_1/2$ appears	57
9. Measuring T_2 : Ramsey vs. Spin Echo	57

Protocol A: Ramsey Measurement	57
Protocol B: Spin Echo	57
10. Takeaways for Transmon Qubits	57
11. Example Measurements	58
Time of Flight (TOF)	58
Resonator Spectroscopy	58
Resonator Punchout	58
Drive Spectroscopy	58
Rabi Oscillations	59
Lesson 1: Introduction to Transmission Line Theory	59
1. Defining the Regime: When Does a Wire Become a Transmission Line?	59
2. Distributed Parameter Representation	60
3. Deriving the Telegrapher Equations	60
Kirchhoff's Voltage Law (KVL)	60
Kirchhoff's Current Law (KCL)	61
Taking the Continuous Limit ($\Delta z \rightarrow 0$)	61
Decoupling into the Wave Equation	61
4. The Lossless Transmission Line Case	61
Propagation Velocity	62
5. Characteristic Impedance	62
Lesson 2: Boundary Conditions, Reflections, and Transients	63
1. Voltage and Current Reflection Coefficients at a Load	63
2. Transient Analysis and Lattice Diagrams	64
Step-by-Step Transient Example	64
3. Power, Return Loss, and Standing Wave Ratio (SWR)	65
Return Loss (RL)	65
Standing Wave Ratio (SWR)	66
Lesson 3: The Scattering Matrix (S-Matrix)	66
1. Core Concept of the S-Matrix	66
2. Mathematical Definition for a 2-Port Network	67
Physical Meaning of the Coefficients	67
3. Deriving S-Parameters from Voltages	67
Finding S_{11} and S_{21} (Drive Port 1, Terminate Port 2)	68
Finding S_{22} and S_{12} (Drive Port 2, Terminate Port 1)	68
Lesson 4: Analyzing Cascaded Lines and Impedance Junctions	68
1. S-Matrix of a Direct Characteristic Impedance Step Junction	68
Finding S_{22} and S_{12} (Driving from Line 2, with Line 1 Matched)	69
2. Complete Scattering Matrix Form	69
The Vector Network Analyzer (VNA)	70
Waves and S-Parameters	70

Individual S-Parameter Definitions	70
Logarithmic Metrics (decibels, dB)	71
Physical Meaning and Measurement Goals	71
VNA Basic Architecture and Signal Flow	71
Directional Couplers and Key RF Specs	73
Circulators and Isolators	73
Calibration: Purpose and Concepts	74
Standard Calibration Techniques	74
The 1-Port Mathematical Error Model	74
Measurement Accuracy and Limits	75
Specific Notes for Qubit/Resonator Characterization	75
Reference Metrics (Typical Baseline values)	75
Quick Glossary	76
Signal Representation	76
1. Why We Need Multiple Representations	76
2. Real vs. Complex (I/Q) Representation	77
Structural Properties of the I/Q Vector Space	78
3. Analytic Signal and Hilbert Transform	78
Frequency-Domain Operation	79
4. Euler's Formula and Complex Spectra	79
Mathematical Foundations	79
Real vs. Complex Signals in the Frequency Domain	80
Benefits of Complex Notation (Algebraic Simplicity)	80
5. Frequency Domain Examples & Symmetry	80
Example 1: Sum of Three Ideal Tones	80
Example 2: Sum of Three Tones with Phase Shifts	81
Key Symmetry Observations	81
6. Modulation Mechanics	82
Basic Concept of Modulation	82
Amplitude Modulation (AM)	82
Quadrature Modulation	82
7. I/Q Modulation and Demodulation Linearity	83
Upconversion Time-Domain Form (Transmitter)	83
I/Q Demodulation (Receiver Architecture)	83
Summary of Modulation vs. Demodulation	84
8. Positive/Negative Frequency and Image Rejection	85
9. Practical Notes for DSP and Hardware	85
Complex Baseband Sampling (Nyquist Criteria)	85
Filtering and Calibration	85
Application: Superconducting Qubit Control	86
10. Summary Reference Sheet	86
Core Formulation Guide	86
Identity Reference Checklist	86
1. The Basic RF Communication System Pipeline	86
Transmitter (TX) Operations	87

The Channel (The Qubit Context)	87
Receiver (RX) Operations	87
2. RF Front-End Architectures	88
Heterodyning (Direct Conversion / Zero-IF)	88
Superheterodyning	88
3. The Digital/Analog Interface & Sampling Rate Impact	89
The Nyquist Boundary	89
4. Analysis of the Three Radio Architectures	89
Architecture A: Direct RF Sampling (Almost All-Digital)	90
Architecture B: IF Sampling Radio	90
Architecture C: Baseband-Sampling Radio	90
5. Architectural Comparison	90
6. Advantages of Shifting to the Digital Domain	91
Radio Frequency System on Chip (RFSoc)	91
Introduction to Radio Frequency System on Chip (RFSoc)	92
1. Challenges in Qubit Control & Readout Systems	92
Current Commercial Solutions	92
Advantages of Open-Source Frameworks	93
2. What is an RFSoc?	93
Strategic Single-Chip Advantages	94
3. The Zynq SoC Evolution	94
Historical Milestones	94
The PS / PL Paradigm	94
4. RFSoc Processing System (PS) Architecture	95
Application Processing Unit (APU)	95
Real-Time Processing Unit (RPU)	95
Platform Management Unit (PMU)	95
5. The Programmable Logic (PL) Fabric	95
Key Functional Hardened Blocks	96
Device Performance and Speed Grades	96
6. PS-PL Interconnect Architecture	96
Memory-Mapped I/O (MMIO)	97
The Three Functional AXI4 Variants	97
7. The RF Data Converter (RFDC) Subsystem	97
RFDC Architecture and Internal Subsystems	97
8. Data Flow and PL-to-RFDC Interfacing	98
Transmit Path (RF-DAC Channel)	98
Receive Path (RF-ADC Channel)	98
9. Understanding Super Sample Rate (SSR)	99
The Speed Gap Problem	99
The SSR Solution	99
10. Multi-GSps Sampling and Nyquist Zone Operations	100
Direct Frequency Sampling (1st Nyquist Zone)	100
Under-Sampling Operations (2nd Nyquist Zone)	100
Comparison: RFSoc vs. Heterodyne Architectures	100

11. Overview of RFSoc Development Boards	101
ZCU208 Evaluation Board	101
Supporting Hardware Add-On Cards	101
Academic Boards: The RFSoc 4x2	101
12. Real-World Application: Two-Qubit Superconducting Architecture	101
Signal Architecture Functions	102
13. System Selection and Feature Matrix	103
14. RFSoc Family: Product Generations and Variants	103
Gen 1, Gen 2, Gen 3, and RFSoc DFE Overview	103
RFDC Tile Architecture	104
Device Speed Grades and Operating Temperature Ranges	104
FPGA Device Specifications	104
15. Processing System Architecture and Inter-Processor Coordination	104
Application Processing Unit (APU)	104
Real-Time Processing Unit (RPU)	105
Platform Management Unit (PMU)	105
16. Programmable Logic Fabric: Detailed Component Overview	105
Configurable Logic Blocks (CLBs)	105
Digital Signal Processing (DSP) Slices	105
Block RAM (BRAM)	106
UltraRAM (URAM)	106
17. Multi-Gigasample Bandwidth and Frequency Coverage	106
Direct Sampling in the 1st Nyquist Zone	106
Under-Sampling in Higher Nyquist Zones	106
18. Development Boards and Their Practical Integration	107
Mainstream Development Boards: ZCU208	107
Academic and Compact Boards: RFSoc 4x2	107
19. Real-World Quantum Control Setup Using RFSoc	108
Signal Routing and Circuit Components	108
System Advantages for Quantum Applications	109
Review of Digital Signal Processing	109
1. Analog-to-Digital Conversion (ADC)	109
2. Sampling Dynamics	109
The Effect of Sampling Rate Selection	110
3. Signal Characterization and Frequency Classification	110
4. The Nyquist-Shannon Sampling Theorem	110
Baseband Formulation	111
Bandpass Formulation	111
5. Nyquist Zones and Aliasing Mechanics	111
General Spectral Mapping Rule	111
6. Undersampling Applications in Quantum Control	112
Strategic Implementation	112
Downconversion Example	112
7. Sampling Jitter	113
8. Quantization Principles	113

Bit Resolution and Step Size	113
The Ideal Quantization Error Model	113
9. Dynamic Range and Spectral Purity	114
Dynamic Range (DR)	114
Frequency Spurs and SFDR	114
10. Digital Hardware Data Representation	115
The Q-Format Configuration	115
Structural Formats	115
11. Fixed-Point Arithmetic Math Operations	115
Addition and Subtraction	115
Multiplication	116
12. Digital-to-Analog Conversion (DAC) Mechanics	116
The Zero-Order Hold (ZOH) Effect	116
13. Reconstruction Filters and Oversampling Strategies	117
Reconstruction Filtering	117
Mitigation Strategies	117
14. Essential DSP Formula Reference Sheet	118
Multirate Operations	118
1. Limitations of Conventional ADC Systems	118
The Anti-Aliasing Constraint	118
Real-World Filtering Compromises	118
2. The Multirate Engineering Solution	119
Oversampling	119
The Computational Bottleneck	119
Multirate Operations	119
3. Core Multirate Mechanisms and Scenarios	120
Common Multirate Scenarios	120
4. Oversampled ADC and Decimation Pipelines	120
Decimation Framework	120
Frequency and Time Domain Decimation Example	121
5. Interpolated DAC Pipelines	121
Interpolation Framework	122
Time and Frequency Domain Interpolation Example	122
6. Structural Comparison Matrix	123
RF Data Converters	123
1. RF Data Converters: Analogue to Digital	123
Analogue to Digital Converter (ADC) Fundamentals	123
Superconducting Qubit Readout Context	124
2. Multi-GHz Nyquist Zones and Direct Sampling	124
First Nyquist Zone Direct Capture	124
Second Nyquist Zone Undersampling	125
Practical Frequency Boundaries	125
3. Analogue Front-End Filtering Strategies	125
4. Hardware Architecture: Interleaved Sub-ADC Tiles	126

The Interleaving Principle	126
Tile Structures: Dual vs. Quad Layouts	127
5. The Digital Downconverter (DDC) Pipeline	127
Generation 3 Frontend Additions: The DSA	128
Quadrature Modulation Correction (QMC)	128
Digital Complex Mixer and NCO Tuning	128
Programmable Decimation Engine	129
6. RF Data Converters: Digital to Analogue	129
Zero-Order Hold (ZOH) Distortions	129
Counteracting Amplitude Roll-Off	130
7. Dual-Zone DAC Modes: Normal vs. Mix-Mode	130
1. Normal ZOH Mode	130
2. Mix-Mode	130
8. The Digital Upconverter (DUC) Pipeline	131
Programmable Interpolation Engine	131
Digital Complex Modulation Mixer	132
Generation 3 Image Rejection (IMR) Filter	132
9. Comprehensive Architectural Layout Reference	132
Semiconductor Qubit Electronics: The Elzerman Method	133
1. The Elzerman Device Architecture	133
Heterostructure and Electrostatic Confinement	133
Charge Detection via Quantum Point Contact (QPC)	133
2. Parameter Tuning and Zeeman Splitting	134
Magnetic Field and Zeeman Energy	134
Thermal and Orbital Energy Relations	134
3. The Two-Level Pulse Readout Protocol	134
Stage 1: Empty the Dot	135
Stage 2: Inject and Wait	135
Stage 3: Readout (Spin-to-Charge Conversion)	135
4. Tuning the Read-Out Configuration	136
Mid-Gate (V_M) Bias vs. Plunger (V_P) Dynamics	136
Experimental Calibration via Color-Map Analysis	136
Finding the Optimal Operating Point	137
5. Measurements and Spin Relaxation (T_1)	137
Statistically Confirming Spin State Detection	137
Spin-Orbit Coupling and the T_1 Lifetime	137
6. Readout Fidelity and Optimization Error Analysis	138
Error Parameters (α and β)	138
Minimizing Total Error	138
Hardware Improvement Strategies	139
7. Advanced Verification: Randomized Benchmarking and QND	139
Randomized Benchmarking (RB)	139
Quantum Non-Demolition (QND) Readout	139

Coherent Manipulation and Characterization of Semiconductor

Spin Qubits	140
1. Core Qubit Operations: The Control Cycle	140
2. Electron Spin Resonance (ESR) and the Rabi Experiment	140
Electromagnetic Wave Generation	140
Magnetic Resonance Condition	141
The “Spin-Dance” Analogy	141
3. Rabi Oscillations and Environmental Decoherence	142
Controlling Rotation Speed	142
Decoherence and Environmental Noise	142
4. Case Study: Cryo-CMOS Qubit Control via Intel Horse Ridge	142
Core Architecture and Setup	143
Characterizing Rabi Performance via Horse Ridge	143
5. Dual-Axis Coherent Control: The Ramsey Experiment	144
The Multi-Pulse Ramsey Sequence	144
Physical Realization of Control Axes	144
Coherent Validation Metrics	145
Classical Electronics Interfaces for Quantum Processors	145
1. Controller Specifications across Qubit Modalities	145
Transmon Qubits vs. Semiconductor Spin Qubits	145
Critical Signal Integrity Vectors	145
2. Structural Benchmarking of Instrumentation Controllers	146
Transmon Controller Architecture (e.g., 5-Qubit Stabilizer Rig)	146
Semiconductor Spin Controller Architecture (e.g., 2-Qubit Silicon Rig)	147
3. The Scaling Bottleneck and Cryogenic Solutions	148
Physical Constraints of Room-Temperature Control	148
Cryogenic Hardware Alternatives	149
4. Deep Dive: Intel Horse Ridge Cryo-CMOS Controller	149
Architecture Overview	149
Digital Back-End Pulse Modulation	150
Scripting an Automated Rabi Sequence Block	151
Analog Front-End Modulator and Package	151
5. Comprehensive Modeling of a Semiconductor Spin-Qubit Processor	152
Performance Metrics	152
Generic Semiconductor Spin-Qubit Quantum Processor	152
Core Node Routing Functions	153
1. Qubit Encoding Physics: Spin vs. Charge	153
Spin Qubits	153
Charge Qubits	155
Hybrid (Singlet-Triplet) Qubits	155
2. Mathematical Framework of the Charge Qubit	155
Case 1: Zero Tunnel Coupling ($t_0 = 0$)	156
Case 2: Finite Tunnel Coupling ($t_0 > 0$)	156
3. Two-Qubit Coupling Mechanism & Exchange Interaction (J)	156
Physical Principle	156

Frequency Shifts	157
4. Hardware Implementation: The Veldhorst Architecture	157
Device Control Topology	157
Charge Stability & The “Diamond Diagram”	157
Readout via Single-Electron Transistor (SET)	157
5. Timeline of a CNOT Gate Operation	158
Phase 1: Initialization & Spin Alignment	158
Phase 2: Initial Target Rotation	158
Phase 3: Controlled Hybridization (The Entangling Step)	159
Phase 4: Conditional Rabi Oscillations	159
Phase 5: Target Finalization	159

Semiconductor Qubits: Materials, Fabrication, and Heterostructure Physics 159

1. Physical Encoding Across Quantum Dot Modalities	159
Classification of Quantum Dots and Qubits	160
2. Strategic Motivations for the Semiconductor Platform	160
3. High-Purity Atomic Layer Deposition: Molecular Beam Epitaxy (MBE)	160
Ultra-High Vacuum (UHV) Chamber Environment	161
Effusion Cells and Shutter Mechanics	161
4. Physics of the AlGaAs/GaAs Heterostructure and 2DEG Formation 161	
Heterostructure Composition and Energy Band Diagrams	161
Modulation Doping and Electrostatic Confinement	162
Gate Deconstruction into Quantum Dots	162
5. The Dephasing Bottleneck: Hyperfine Interaction	162
The Hyperfine Coupling Mechanism	163
Overhauser Field Fluctuations and Dephasing	163

Energy Band Engineering of the $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}/\text{i-GaAs}$ Heterojunction 163

1. Isolated Material Parameters before Contact ($t = 0$)	163
Material Properties & Mathematical Models	164
Bulk Band Alignment Comparison	164
2. Band Discontinuities at the Interface: The Electron Affinity Rule . 164	
3. Electrostatic Equilibrium, Band Bending, and Poisson’s Equation . 166	
Electrostatic Modeling via Poisson’s Equation	166
The Resulting Equilibrium Band Diagram	167
4. Alternate Model Formulations: Metal-Semiconductor Analogies . . 167	
Case A: Degenerate n -type to p -type Homojunction	167
Case B: Metal-Like Degenerate Layer to Intrinsic Substrate . . . 168	
Exercise 2: Deriving Interface Band Discontinuities (The Affinity Rule)	168
Exercise 3: Charge Density and Electrostatic Potential (Poisson Modeling)	169
Exercise 4: Comprehensive Band Engineering and 2DEG Analysis 169	

Exercise 5: Comparative Analysis of Metal-Semiconductor Analogs	170
Solutions & Detailed Explanations	170
Solution to Exercise 1: Boundary Conditions and Energy Parameters	170
Solution to Exercise 2: Interface Band Discontinuities	171
Solution to Exercise 3: Charge Density and Poisson Modeling	172
Solution to Exercise 4: Comprehensive Band Engineering and 2DEG Analysis	173
Solution to Exercise 5: Comparative Analysis of Metal-Semiconductor Analogs	174
Lecture 3: The Physics of Semiconductor Quantum Dots	175
1. The Hyperfine Interaction and Nuclear Spin Bath	175
The “Spin Bath” Problem	175
2. Bandgap Engineering and Heterostructure Mechanics	175
Bandgap Variation in $\text{Si}_{1-x}\text{Ge}_x$	176
Type II Band Alignment	176
3. Strain Engineering and Effective Mass Modification	176
The Challenge of Pure Silicon	176
The Role of Strain (LPCVD Deposition)	177
Heterostructure Band Engineering and Quantum Dot Confinement in Strained $\text{Si}_{1-x}\text{Ge}_x$ Systems	177
1. Relaxation, Lattice Constants, and Misfit Strain Mechanics	177
Bulk Material Relationships	177
Epitaxial Strain Generation	178
2. Band Structure Deformation and Effective Mass Alterations	178
Conduction Band Flattening under Tensile Strain	179
Valence Band Splitting under Compressive Strain	179
3. Structure A: Tensile Strained Silicon Quantum Wells	179
Mathematical Discontinuity Derivation	179
Resulting Band Edge Alignment	180
4. Structure B: Compressively Strained Germanium Quantum Wells	181
Mathematical Discontinuity Derivation	181
Resulting Band Edge Alignment	181
5. Analytical Practice Problems with Solutions	182
Problem 1: Discontinuity Sensitivity Analysis	182
Problem 2: Mechanical Strain Limit Check	183
Lecture 5: The Physics and Material Challenges of MOS Qubits	184
Section 1: Deepening the Physical Concepts	184
1. The Multi-Dimensional Schrödinger Equation	184
2. Physical Reality of Boundary Conditions	184
Section 2: Mathematical Derivations of the Exercises	185
Exercise 1: Energy Splitting Derivation (ΔE)	185
Exercise 2: Scaling Law Derivation (ΔE_{eV} vs L_{nm})	186
Section 3: Material Implementation Exercises	187

Exercise 3: Designing a Gallium Arsenide (GaAs) Qubit	187
Exercise 4: Designing a Silicon (Si) Qubit	188
Mechanical Conclusion	188
1. The MOS Quantum Dot Architecture & Electrostatic Confinement	188
2. The Microscopic Nightmare: Interface Defects and Oxide Charges	189
The Four Oxide Charges & Their Impact on Qubits:	189
3. Diagnostics via C-V (Capacitance-Voltage) Characterization	190
Poisson-Boltzmann Analysis (Page 12)	190
Reading Trap Density (N_{it}) from C-V Shifts (Page 13)	190
4. Scalability Barriers: Fabrication Variability & Control	191
Fabrication Variability	191
Control, Readout, and Coupling	191
Lecture 6: Qubit Fabrication & Device Architectures	192
1. Planar Silicon Qubits: The Multi-Layer Gate Stack	192
The Multi-Step Process Demystified:	192
2. Advanced Planar Design: FD-SOI Qubit Technology	193
Single vs. Double Quantum Dot Operating Conditions	193
The Power of UTBB (Ultra-Thin Body and BOX)	194
3. 3D Technologies: FinFET Qubit Architectures	194
Why FinFETs are Ideal for Large Qubit Arrays	194
1. The Single Quantum Dot & Coulomb Blockade Mechanics	195
The Constant Interaction (Capacitive) Model	195
2. SET Readout for a Quantum Dot	195
The Measurement Technique	196
3. The Double Quantum Dot (DQD) Capacitive Model & Stability	
Diagram	196
The DQD Stability Diagram	196
4. Spin Manipulation: The Zeeman Effect	196
5. DQD Operation Sequence & Pauli Spin Blockade (PSB)	197
Step 1: Initialization	197
Step 2: Tuning Energy Levels	197
Step 3: Spin-Selective Tunneling (The Blockade)	197
Step 4: Manipulation and Readout	198
1. Electrostatic Derivation of the Dot Potential (V_{DOT})	198
2. The Charging Energy Component	199
3. Deriving Tunneling Conditions (T_1 and T_2)	199
Condition T_1 : Source $\rightarrow$ Dot Tunneling	200
Condition T_2 : Dot $\rightarrow$ Drain Tunneling	200
4. The Coulomb Diamond Diagram Explained	201
5. High-Sensitivity Qubit Readout via Cross-Capacitive Shifting	201
6. Double Quantum Dot (DQD) Stability Contours	201

Qubit Introduction

What is a qubit?

A qubit, or quantum bit, is the fundamental unit of quantum information. It is a two-level quantum system that can exist in a superposition of states, unlike a classical bit which can only be in one of two states (0 or 1). A qubit can be represented as:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

where α and β are complex numbers that satisfy the normalization condition $|\alpha|^2 + |\beta|^2 = 1$.

Computers and Computation

Models of classical computation

Classical computation can be realized with different physical technologies: - Mechanical - Electrical - Optical - Biological

Examples

- Curta calculator
- Digi-Comp I
- Babbage's difference engine

Models of computation

Computing can be described using: - **Physical models:** mechanical, optical, electrical, biological - **Conceptual models:** - Turing machine - Cellular automata - Von Neumann architecture - Logic-in-memory

Von Neumann architecture

Main components: - CPU - Memory - Bus - I/O

Moore's Law

Historically, transistor density and performance improved rapidly over time, driving classical hardware scaling.

Quantum history

- In the early 1980s, Richard Feynman suggested that quantum systems should be simulated with quantum systems.
- This motivated the idea of a quantum computer.
- Important milestones in the 1990s:

- Shor’s algorithm
- First quantum error-correcting code

Current status

- Quantum processors with more than 100 qubits already exist.
- IBM Condor is reported to have more than 1000 qubits.

Main quantum computer players

Quantum computing is being developed by several major industrial and academic groups.

Data Representation

Classical computers

- Classical computers process information using **transistors**.
- Information is encoded in bits:
 - 0
 - 1

Quantum computers

- Quantum computers use **qubits**.
- A qubit can be in a superposition:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

- If $|\alpha|^2 = |\beta|^2 = 1/2$, then the qubit is measured as:
 - 50% $|0\rangle$
 - 50% $|1\rangle$

Classical parallelism

- N classical bits represent one N -bit state.
- Example for $N = 3$:
 - 000, 001, $\dots$, 111
- A classical function must be evaluated on each input separately.

Quantum parallelism

- N qubits describe a state with 2^N components.
- Example for $N = 3$:

$$|\psi\rangle = c_1|000\rangle + c_2|001\rangle + c_3|010\rangle + c_4|011\rangle + c_5|100\rangle + c_6|101\rangle + c_7|110\rangle + c_8|111\rangle$$

DiVincenzo Criteria

Requirements of a classical logic element

A classical logic element should:

- Scale well for complex systems
- Be characterizable and manufacturable on a large scale
- Represent classical information reliably
- Allow state initialization and measurement
- Provide gain and noise margins
- Be robust against noise and failure

Transistors satisfy these requirements.

What makes a good qubit?

A good qubit should:

- Be a scalable physical system
- Allow initialization of the input state
- Have long coherence times
- Support a universal set of gates
- Represent quantum information with high fidelity
- Allow conversion between stationary and flying qubits
- Allow faithful transmission of flying qubits between locations

Qubit Figures of Merit

Coherence time

Coherence time is the time before quantum information is lost due to interaction with the environment.

Main mechanisms:

- Energy relaxation: T_1
- Dephasing: T_ϕ

The relation is:

$$\frac{1}{T_2} = \frac{1}{T_\phi} + \frac{1}{2T_1}$$

Gate speed

- **Gate time:** time required to perform a quantum operation
- A useful figure of merit is the number of gates that can be performed within the qubit lifetime:

$$N_{\text{gates}} \sim \frac{T_2}{t_{\text{gate}}}$$

Gate fidelity

Gate fidelity measures how accurately a gate is performed.

- High fidelity is essential for reliable quantum computation.
 - It can be measured by:
 - Process tomography
 - Randomized benchmarking
-

Bits and Qubits

Classical computers

- Classical computers use Boolean logic
- Universal classical computation can be built from single- and two-bit gates such as NOT and AND

Quantum computers

- Quantum computers use quantum logic
- A qubit can be in superposition:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

X gate

The quantum NOT gate is the **X gate**:

- $|0\rangle \rightarrow |1\rangle$
- $|1\rangle \rightarrow |0\rangle$

For a general state:

$$\alpha|0\rangle + \beta|1\rangle \rightarrow \beta|0\rangle + \alpha|1\rangle$$

Circuits in space vs circuits in time

Classical logic

- Gates act on wires in space
- Input and output are at different physical locations

Quantum logic

- Gates act on qubits in time
 - The same qubit evolves during the operation
-

Two-Qubit Gates

Classical XOR

For two bits A and B :

A	B	$A \oplus B$
0	0	0
0	1	1
1	0	1
1	1	0

Quantum CNOT

The quantum analogue is the **CNOT gate**.

If the control qubit is in superposition, CNOT can create entanglement.

For example, if the control qubit is in a superposition and the target is $|0\rangle$:

$$|\psi_{\text{in}}\rangle = (|0\rangle + |1\rangle)_A |0\rangle_B$$

then:

$$|\psi_{\text{out}}\rangle = |0\rangle_A |0\rangle_B + |1\rangle_A |1\rangle_B$$

Qubit States

A qubit is described by:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

where:

- α and β are complex numbers
- normalization requires:

$$|\alpha|^2 + |\beta|^2 = 1$$

Bra-ket notation

- $|\psi\rangle$ is a ket
 - $\langle\psi|$ is its Hermitian conjugate, or bra
 - The inner product is written as $\langle\phi|\psi\rangle$
-

Measurement of Qubits

Measurement process

- Measurement is an active process
- The measurement apparatus interacts with the qubit
- Only partial information is obtained
- The full quantum state is not directly accessible

Measurement depends on a basis

Most measurements are performed in the computational basis:

$$\{|0\rangle, |1\rangle\}$$

State collapse

After measurement, the qubit collapses to one basis state:

- $|0\rangle$
- $|1\rangle$

This is called **Z-measurement**.

Born rule

The probability of an outcome is the squared modulus of its amplitude.

For $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$:

- $P(0) = |\alpha|^2$
- $P(1) = |\beta|^2$

Other common measurement bases

X basis

- $|+\rangle = (|0\rangle + |1\rangle)/\sqrt{2}$
- $|-\rangle = (|0\rangle - |1\rangle)/\sqrt{2}$

Y basis

- $|+i\rangle = (|0\rangle + i|1\rangle)/\sqrt{2}$
- $|-i\rangle = (|0\rangle - i|1\rangle)/\sqrt{2}$

Bloch Sphere

Any single-qubit pure state can be written as:

$$|\psi\rangle = \cos(\theta/2)|0\rangle + e^{i\phi} \sin(\theta/2)|1\rangle$$

where $\theta \in [0, \pi]$ and $\phi \in [0, 2\pi]$.

The Bloch vector is:

$$\mathbf{r} = (\sin \theta \cos \phi, \sin \theta \sin \phi, \cos \theta)$$

Examples

- $|0\rangle \rightarrow \mathbf{r} = (0, 0, 1)$
 - $|1\rangle \rightarrow \mathbf{r} = (0, 0, -1)$
 - $|+\rangle \rightarrow \mathbf{r} = (1, 0, 0)$
 - $|-\rangle \rightarrow \mathbf{r} = (-1, 0, 0)$
-

Unitary Operations and Single-Qubit Gates

Single-qubit gates

A single-qubit gate transforms one quantum state into another:

$$|\psi'\rangle = U|\psi\rangle$$

To preserve normalization, U must be **unitary**.

Unitary matrices

A matrix U is unitary if:

$$U^\dagger U = U U^\dagger = I$$

Properties:

- Preserves total probability
 - Preserves inner products
 - Is reversible
 - Does not destroy information
-

Common Single-Qubit Gates

Pauli gates

Pauli-X $X = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$

- $X|0\rangle = |1\rangle$
- $X|1\rangle = |0\rangle$

Interpretation: rotation around the x-axis by π .

Pauli-Y $Y = \begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix}$

Interpretation: rotation around the y-axis by π .

Pauli-Z $Z = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

•

$$Z|+\rangle = |-\rangle$$

Interpretation: rotation around the z-axis by π .

Pauli matrix properties

- $X^2 = Y^2 = Z^2 = I$
- All Pauli matrices are Hermitian
- Commutation relations:
 - $[X, Y] = 2iZ$
 - $[Y, Z] = 2iX$
 - $[Z, X] = 2iY$

- Anticommutation relations:
 - $\{X, Y\} = 0$
 - $\{Y, Z\} = 0$
 - $\{Z, X\} = 0$
-

Rotation Gates

Rotation around axis n :

$$R_n(\theta) = e^{-i\theta\sigma_n/2}$$

In particular:

- $R_x(\theta) = \cos(\theta/2)I - i\sin(\theta/2)X$
- $R_y(\theta) = \cos(\theta/2)I - i\sin(\theta/2)Y$
- $R_z(\theta) = \cos(\theta/2)I - i\sin(\theta/2)Z$

When $\theta = \pi$, these reduce to the corresponding Pauli operations.

Hadamard Gate

The Hadamard gate is:

$$H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$$

Action

- $H|0\rangle = |+\rangle$
-

$$H|1\rangle = |-\rangle$$

Role

- Creates superposition
 - Changes basis between Z and X bases
-

Phase Gate

The S gate is the phase gate:

$$S = \begin{bmatrix} 1 & 0 \\ 0 & i \end{bmatrix}$$

Action

- $S|+\rangle = |+i\rangle$
- $S|-\rangle = |-i\rangle$

Role

- Adds a 90° phase shift
 - Helps change between Z and Y bases
-

Applying Multiple Gates

Multiple gates are applied sequentially using matrix multiplication.

Matrix expressions act **right to left**, while circuit diagrams are read **left to right**.

Example:

X Y H $|0\rangle$

means:

1. Start in $|0\rangle$
 2. Apply H
 3. Apply Y
 4. Apply X
-

Density Matrix

Why it is needed

Real systems are not perfectly isolated, so pure states are not always sufficient. A more general description is given by the **density matrix** ρ .

Pure state

For a pure state:

$$\rho = |\psi\rangle\langle\psi|$$

General properties

- ρ is a 2×2 matrix for a qubit
- ρ is Hermitian
- $\text{tr}(\rho) = 1$

Matrix form

For a qubit:

$$\rho = \begin{pmatrix} \rho_{11} & \rho_{12} \\ \rho_{21} & \rho_{22} \end{pmatrix}$$

- Diagonal terms: populations of $|0\rangle$ and $|1\rangle$
- Off-diagonal terms: coherences
- Hermiticity implies:

$$\rho_{12} = \rho_{21}^*$$

Mixed state

If the qubit is in a probabilistic mixture:

$$\rho = \sum_i p_i |\psi_i\rangle\langle\psi_i|$$

Pure vs mixed states

A useful indicator is:

- $\text{tr}(\rho^2) = 1$ for a pure state
- $\text{tr}(\rho^2) < 1$ for a mixed state

Superconducting Qubits

1. Superconductivity Review

Superconductivity is a macroscopic quantum phenomenon where electrons form a collective state that flows without resistance.

Cooper Pairs

In 1957, Leon Cooper, John Bardeen, and Robert Schrieffer proposed the BCS theory, establishing that superconductivity arises from the formation of bound electron pairs called **Cooper pairs**.

At low temperatures, an electron traveling through a crystal lattice slightly distorts the surrounding positively charged ions. This local polarization creates an attractive potential that draws in a second electron. The resulting bound Cooper pair consists of two strongly correlated electrons with:

- Opposite momentum ($\vec{p}_1 = -\vec{p}_2$)
- Opposite spin ($\uparrow\downarrow$), forming a spin-singlet state with an overall spin $S = 0$

Because they have integer spin, Cooper pairs act as bosons. Below a material's **critical temperature** (T_c), they undergo a process akin to Bose-Einstein condensation, collapsing into a single, coherent macroscopic quantum state described by a unified wave function.

The Macroscopic Wave Function

The coherent quantum condensate can be expressed as a complex scalar order parameter:

$$\psi(\vec{x}, t) = \psi_0 e^{-j \frac{1}{\hbar} (Et - \vec{p} \cdot \vec{x})} = \sqrt{\rho} e^{j\theta(\vec{x}, t)}$$

Where:

- $\rho = \psi\psi^* = |\psi_0|^2$ is the physical local density of Cooper pairs.
- $\theta(\vec{x}, t)$ is the macroscopic quantum phase.

Key Electromagnetic Properties

1. **Zero Electrical Resistance:** Because the Cooper pairs are locked into a single macroscopic state, scattering from individual lattice defects is energetically suppressed, dropping DC electrical resistance to zero.
2. **Phase Coherence:** The phase θ is rigid across the entire superconductor, locking the material into a coherent quantum system.
3. **Critical Magnetic Field (H_c):** Applying an external magnetic field raises the free energy of the superconducting state. If the field exceeds the critical magnetic field H_c , superconductivity breaks down, and the material reverts to its normal resistive state. H_c follows a parabolic temperature dependence:

$$H_c(T) = H_c(0) \left[1 - \left(\frac{T}{T_c} \right)^2 \right]$$

The Meissner Effect

Discovered in 1933 by Walter Meissner and Robert Ochsenfeld, the Meissner effect shows that a superconductor does more than act as an ideal conductor. When cooled below T_c in the presence of a magnetic field, it actively expels all internal magnetic flux lines from its bulk.

The material sets up persistent surface screening currents that generate an internal field that cancels out the external field, causing it to behave as a perfect diamagnet with a magnetic susceptibility $\chi = -1$.

2. The Josephson Junction

Physical Definition

In 1962, Brian D. Josephson analyzed a system consisting of two superconducting layers (S_1, S_2) separated by a thin insulating barrier (I). This configuration forms an **SIS Josephson Junction**.

Superconductor 1 (S_1)	Insulator (I)	Superconductor 2 (S_2)
Wavefunction: ψ_1	Thickness	Wavefunction: ψ_2
Phase: θ_1	0.1 – 1 nm	Phase: θ_2

⇓

Cooper Pair Tunneling Loop

When the insulating barrier is thin enough (0.1 – 1 nm), the macroscopic wave functions of the two superconductors overlap within the barrier, allowing Cooper pairs to tunnel coherently across the gap without breaking apart.

Coupled Superconductors Derivation

Let the isolated states of the two superconductors be:

$$\psi_1 = \sqrt{\rho_1} e^{j\theta_1}, \quad \psi_2 = \sqrt{\rho_2} e^{j\theta_2}$$

The relative phase difference across the barrier is defined as:

$$\phi = \theta_2 - \theta_1$$

When the two systems are coupled across the thin barrier, their behavior is governed by a coupled set of Schrödinger equations:

$$j\hbar \frac{\partial \psi_1}{\partial t} = E_1 \psi_1 + K_S \psi_2$$

$$j\hbar \frac{\partial \psi_2}{\partial t} = E_2 \psi_2 + K_S \psi_1$$

Where K_S is the coupling coefficient determined by the barrier thickness and geometry. Substituting the wave functions into these equations isolates the dynamics of the Cooper pair densities (ρ_1, ρ_2) and the relative phase (ϕ).

3. The Constitutive Josephson Relations

Evaluating the coupled state equations yields the two fundamental Josephson relations that govern the behavior of the junction.

1. The Current-Phase Relation (First Josephson Relation)

The charge current I passing through the junction is directly proportional to the sine of the relative macroscopic phase difference:

$$I = I_c \sin \phi$$

Where I_c is the **critical current** of the junction—the maximum supercurrent the barrier can support before a voltage drops across it.

2. The Voltage-Phase Relation (Second Josephson Relation)

When a voltage V exists across the junction, the relative phase difference evolves linearly over time:

$$V = \frac{\Phi_0}{2\pi} \frac{d\phi}{dt} = \frac{\hbar}{2e} \frac{d\phi}{dt}$$

Where Φ_0 is the **magnetic flux quantum**, a fundamental physical constant given by:

$$\Phi_0 = \frac{h}{2e} \approx 2.0678 \times 10^{-15} \text{ Wb}$$

4. Josephson Inductance and Nonlinearity

The Josephson junction acts as a non-linear inductor. We can show this by differentiating the current-phase relation with respect to time:

$$\frac{dI}{dt} = I_c \cos \phi \left(\frac{d\phi}{dt} \right)$$

Substituting the voltage-phase relation ($\frac{d\phi}{dt} = \frac{2\pi V}{\Phi_0}$) into this equation yields:

$$\frac{dI}{dt} = I_c \cos \phi \left(\frac{2\pi V}{\Phi_0} \right)$$

Rearranging this into the standard form for an inductor ($V = L \frac{dI}{dt}$) isolates the **Josephson Inductance** (L_J):

$$V = \left[\frac{\Phi_0}{2\pi I_c \cos \phi} \right] \frac{dI}{dt} \implies L_J(\phi) = \frac{\Phi_0}{2\pi I_c \cos \phi}$$

By defining the characteristic constant inductance factor as $L_{J0} = \frac{\Phi_0}{2\pi I_c}$, the total parametric inductance becomes:

$$L_J(\phi) = \frac{L_{J0}}{\cos \phi}$$

Physical Significance

Unlike a standard wire inductor where L is a fixed value, the Josephson inductance depends nonlinearly on the phase variable ϕ (and therefore on the current passing through it). As $\phi \rightarrow \frac{\pi}{2}$, the inductance grows toward infinity. This clean, dissipationless nonlinearity is what enables the development of superconducting qubits.

5. Josephson Energy

The potential energy stored in the junction is calculated by integrating the electrical power over time:

$$E = \int I \cdot V dt$$

Substituting the two Josephson relations into this integral:

$$E = \int (I_c \sin \phi) \left(\frac{\Phi_0}{2\pi} \frac{d\phi}{dt} \right) dt = \frac{I_c \Phi_0}{2\pi} \int \sin \phi d\phi$$

Evaluating this integral yields the potential energy profile of the junction:

$$E(\phi) = E_J(1 - \cos \phi)$$

Where the characteristic **Josephson Energy** (E_J) scale is defined as:

$$E_J = \frac{I_c \Phi_0}{2\pi}$$

This cosine potential energy profile replaces the standard parabolic potential of a traditional linear inductor, introducing the anharmonicity needed to form an artificial atom.

6. Artificial Atoms vs. Natural Atoms

The Energy Spectrum Problem

- **Natural Atoms:** Systems like Hydrogen or Rubidium have a naturally non-linear potential energy profile. This creates an unequally spaced energy spectrum, allowing an external laser to address a specific transition (e.g., $|0\rangle \rightarrow |1\rangle$) without driving higher-order transitions.
- **Standard LC Circuits:** A standard lumped-element inductor connected to a capacitor forms a harmonic oscillator. Its potential energy profile is perfectly parabolic ($E \propto \Phi^2$), producing an equally spaced ladder of energy states ($E_n = \hbar\omega_0(n + 1/2)$).

[Image comparing a harmonic parabolic potential energy well with equally spaced levels to an anharmonic cosine potential well with unequally spaced levels]

If you attempt to use a standard LC circuit as a qubit, an external microwave pulse meant to transition the system from $|0\rangle \rightarrow |1\rangle$ will also drive transitions from $|1\rangle \rightarrow |2\rangle$ and $|2\rangle \rightarrow |3\rangle$. This causes the quantum state to leak into higher energy levels, preventing the system from acting as a two-level qubit.

Cryogenic Initialization

To initialize an artificial atom into its ground state $|0\rangle$, thermal energy must be significantly lower than the qubit's transition energy:

$$k_B T \ll \hbar\omega_{01}$$

For a typical qubit transition frequency $\omega_{01}/2\pi \approx 5$ GHz, the equivalent thermal transition boundary is $T \approx 240$ mK. Superconducting quantum processors are housed inside dilution refrigerators that reach baseline temperatures of ~ 10 mK. At these temperatures, thermal fluctuations are suppressed, and the system cools reliably into its ground state.

7. Harmonic Oscillator Quantization

Before analyzing the non-linear case, we can review the canonical quantization of a standard linear LC oscillator. The classical Hamiltonian of the circuit is expressed using charge Q and flux Φ :

$$H = \frac{Q^2}{2C} + \frac{\Phi^2}{2L}$$

Through canonical quantization, the continuous variables become non-commuting operators ($\hat{Q}, \hat{\Phi}$) that satisfy the commutation relation:

$$[\hat{\Phi}, \hat{Q}] = j\hbar$$

We define the standard creation ($\hat{a}^\dagger$) and annihilation ($\hat{a}$) operators to express the charge and flux operators in terms of the circuit's characteristic impedance $Z_0 = \sqrt{\frac{L}{C}}$:

$$\hat{\Phi} = \sqrt{\frac{\hbar Z_0}{2}}(\hat{a}^\dagger + \hat{a})$$

$$\hat{Q} = j\sqrt{\frac{\hbar}{2Z_0}}(\hat{a}^\dagger - \hat{a})$$

This transforms the linear Hamiltonian into its quantized form, featuring equally spaced energy rungs and a residual zero-point energy:

$$\hat{H} = \hbar\omega_r \left(\hat{a}^\dagger \hat{a} + \frac{1}{2} \right)$$

8. Superconducting Qubit Hamiltonian

By replacing the linear inductor with a non-linear Josephson junction, the parabolic potential energy term $\frac{\hat{\Phi}^2}{2L}$ is replaced by the junction's cosine potential energy profile:

$$\hat{H} = \frac{\hat{Q}^2}{2C} - E_J \cos \left(2\pi \frac{\hat{\Phi}}{\Phi_0} \right)$$

To simplify the math, we transition to dimensionless variables:

1. The Reduced Phase Operator ($\hat{\phi}$):

$$\hat{\phi} = \frac{2\pi\hat{\Phi}}{\Phi_0}$$

2. The Excess Cooper-Pair Number Operator ($\hat{n}$):

$$\hat{n} = \frac{\hat{Q}}{2e}$$

The canonical commutation relation for these dimensionless operators preserves the uncertainty principle:

$$[\hat{\phi}, \hat{n}] = j$$

Substituting these operators into the Hamiltonian yields the standard expression for a superconducting qubit:

$$\hat{H} = 4E_c \hat{n}^2 - E_J \cos \hat{\phi}$$

Where E_c represents the **Charging Energy** required to add a single Cooper pair to the circuit's capacitance:

$$E_c = \frac{e^2}{2C}$$

9. Charge Offset and Dephasing Noise

In practical circuits, stray ambient electric fields or intentional gate connections introduce a background charge offset. This transforms the kinetic charging term of the Hamiltonian:

$$\hat{H} = 4E_c (\hat{n} - n_g)^2 - E_J \cos \hat{\phi}$$

Where $n_g = \frac{C_g V_g}{2e}$ represents the continuous external dimensionless offset charge.

[Image plotting energy levels of a charge qubit versus offset charge n_g showing high charge dispersion curves]

The Charge Noise Problem

When the charging energy is comparable to or larger than the Josephson energy ($E_J \lesssim E_c$), the system operates in the **Cooper-Pair Box** regime. In this regime, the system's energy levels depend heavily on the value of n_g .

Small background fluctuations in n_g caused by moving trap charges in the substrate shift the qubit's transition frequency over time. This variation causes the quantum phase to drift, leading to rapid **charge-noise dephasing** and short coherence times (T_2).

10. The Transmon Regime

The **Transmon Qubit** resolves this charge-noise sensitivity by shunting the Josephson junction with a large parallel capacitor. This significantly increases the total capacitance C , lowering the characteristic charging energy E_c .

COOPER-PAIR BOX		TRANSMON REGIME	
	• EJ / EC ~ 1		• EJ / EC >> 1 (~50)
	• High Anharmonicity		• Lower Anharmonicity
	• High Charge Dispersion		• Flat Energy Bands
	• Short Coherence Time		• Long Coherence Time

The transmon is defined by engineering this ratio to be large:

$$\frac{E_J}{E_c} \gg 1 \quad (\text{Typical Target: } \frac{E_J}{E_c} \approx 50)$$

Suppression of Charge Dispersion

As the ratio $\frac{E_J}{E_c}$ increases, the charge dispersion (the dependence of the energy levels on n_g) decreases exponentially:

$$\epsilon_m \propto \exp\left(-\sqrt{8E_J/E_c}\right)$$

This flattens the energy bands relative to n_g . Because the energy bands are flat, fluctuations in background charge no longer shift the qubit's transition frequency, protecting the system from charge-noise dephasing.

Taylor Series Expansion of the Potential

Because $E_J \gg E_c$, the phase variable ϕ is tightly confined near zero. This allows us to approximate the cosine potential using a Taylor series expansion:

$$\cos \hat{\phi} \approx 1 - \frac{\hat{\phi}^2}{2} + \frac{\hat{\phi}^4}{24} - \dots$$

Substituting this expansion back into the Hamiltonian:

$$\hat{H} \approx 4E_c(\hat{n} - n_g)^2 + \frac{1}{2}E_J\hat{\phi}^2 - \frac{1}{24}E_J\hat{\phi}^4$$

- **Harmonic Base** ($\frac{1}{2}E_J\hat{\phi}^2$): Combined with the kinetic charging term, this forms a baseline harmonic oscillator profile with a plasma frequency $\omega_p \approx \sqrt{8E_cE_J}$.
- **Anharmonic Perturbation** ($-\frac{1}{24}E_J\hat{\phi}^4$): This higher-order term introduces a weak negative nonlinearity, pulling down the higher energy levels.

Dressed Multi-Level Energy Distribution

Transforming this expanded relation into standard creation and annihilation operators yields a weakly anharmonic multi-level spectrum:

$$E_n \approx \hbar\omega_p \left(n + \frac{1}{2} \right) - \frac{E_c}{2} (n^2 + n)$$

The transition energies between consecutive levels are no longer equal:

- $E_{10} = E_1 - E_0 = \hbar\omega_p - E_c$
- $E_{21} = E_2 - E_1 = \hbar\omega_p - 2E_c$

The difference between these transition energies is the **Anharmonicity** (α):

$$\alpha = E_{21} - E_{10} = -E_c$$

For a typical transmon, $E_c \approx 300$ MHz. This provides enough anharmonicity to cleanly address the $|0\rangle \rightarrow |1\rangle$ transition using short microwave pulses ($\sim 10 - 20$ ns) without driving the system into higher-order states.

11. Architectural Families of Superconducting Qubits

Superconducting qubits can be engineered into different families by varying the balance between charging energy, Josephson energy, and inductive configurations.

[Image comparing schematic circuit diagrams for a Charge, Transmon, Flux, and Fluxonium qubit]

1. The Transmon

- **Ratio:** $E_J/E_c \gg 1$
- **Design:** Uses a large shunted capacitance to minimize charge noise at the cost of working with weak anharmonicity.

2. The Flux Qubit

- **Ratio:** $E_J/E_c \sim 10$
- **Design:** Uses a superconducting loop interrupted by multiple junctions. The system is biased with an external magnetic flux ($\Phi_{\text{ext}} = \frac{\Phi_0}{2}$), storing information in the direction of persistent circulating supercurrents (clockwise vs. counter-clockwise). This provides high anharmonicity but increases sensitivity to magnetic flux noise.

3. The Fluxonium

- **Ratio:** Includes a large shunting array of Josephson junctions that act as a superinductance.
- **Design:** The large inductance bypasses low-frequency charge noise while maintaining high anharmonicity and a large flux-tuning range, helping protect the qubit's coherence.

12. Core Mathematical Reference Sheet

Parameter Metric	Mathematical Formulation
First Josephson Relation (Current)	$I = I_c \sin \phi$
Second Josephson Relation (Voltage)	$V = \frac{\Phi_0}{2\pi} \frac{d\phi}{dt} = \frac{\hbar}{2e} \frac{d\phi}{dt}$
Variable Josephson Inductance	$L_J(\phi) = \frac{\Phi_0}{2\pi I_c \cos \phi} = \frac{L_{J0}}{\cos \phi}$
Junction Potential Energy Well	$E(\phi) = E_J(1 - \cos \phi)$ (where $E_J = \frac{I_c \Phi_0}{2\pi}$)
Canonical Quantization Operator	$[\hat{\phi}, \hat{n}] = j$
Total Qubit Circuit Hamiltonian	$\hat{H} = 4E_c(\hat{n} - n_g)^2 - E_J \cos \hat{\phi}$ (where $E_c = \frac{e^2}{2C}$)
Transmon Anharmonicity Index	$\alpha = E_{21} - E_{10} = -E_c$

Circuit QED

1. From Atomic Physics to Quantum Optics

Quantum optics historically began with atomic physics, studying how isolated atoms interact with freely propagating or weakly confined electromagnetic fields.

Two-Level System Driven by a Field

A physical quantum system with a complex Hilbert space can often be truncated to a **two-level system (TLS)** when the driving field is near-resonant with a specific transition. This interaction between light and matter is dominated by the electric dipole coupling, where the atom's internal charge distribution shifts in response to an external electric field:

$$H_{\text{int}} = -\vec{d} \cdot \vec{E}(t)$$

Where $\vec{d} = q\vec{r}$ is the electric dipole moment operator. When driven by a resonant classical microwave or optical field:

$$\vec{E}(t) = \vec{E}_0 \cos(\omega_{01}t)$$

The field oscillates exactly at the transition frequency $\omega_{01} = \frac{E_1 - E_0}{\hbar}$ between the ground state $|0\rangle$ and excited state $|1\rangle$. This classical drive induces coherent **Rabi oscillations**, swapping the population between states at a frequency $\Omega_R \propto \vec{d} \cdot \vec{E}_0$.

Single-Photon Coupling and Vacuum Fluctuations

In free space, the interaction between a single photon and an atom is incredibly weak because the photon's energy is spread over a large spatial volume. The fundamental question of cavity quantum electrodynamics (QED) is: *Can a single photon, or even pure vacuum fluctuations (zero-point energy), produce a measurable effect on an atom?*

The answer becomes yes when the electromagnetic field is strongly confined within a small volume V . The root-mean-square vacuum electric field fluctuation E_{zpf} inside a volume V is given by:

$$E_{\text{zpf}} = \sqrt{\frac{\hbar\omega}{2\epsilon_0 V}}$$

As the cavity volume V decreases, E_{zpf} increases. This boosts the light-matter coupling strength $g = \vec{d} \cdot \vec{E}_{\text{zpf}}$, allowing a single photon to alter the quantum state of the atom before it decays.

2. Cavity QED Foundations

The Fabry–Pérot Resonator

In atomic cavity QED, light confinement is achieved using a **Fabry–Pérot cavity**, which consists of two highly reflective parallel mirrors separated by a macroscopic distance L .

Boundary conditions enforce that the electric field must vanish at the mirror surfaces. Consequently, only specific wavelengths can survive inside the cavity due to constructive interference. This yields the strict resonance condition:

$$2L = m\lambda \implies f_m = \frac{mc}{2L}$$

Where m is an integer mode number, λ is the wavelength, and c is the speed of light in the medium.

Quantization of the Electromagnetic Field

The cavity quantizes the continuous spectrum of free-space light into discrete, isolated resonant modes. Each mode m behaves mathematically as an independent quantum harmonic oscillator characterized by:


- Its resonance frequency $\omega_m = 2\pi f_m$.
- Its photon loss linewidth κ_m , which defines how quickly photons leak out through the imperfect mirrors.

The mode spacing (Free Spectral Range, $\text{FSR} = \frac{c}{2L}$) must be significantly larger than the linewidth κ to ensure the modes remain isolated and do not overlap.

3. From Cavity QED to Circuit QED

Why Shift to Circuit QED?

While atomic cavity QED offers an excellent platform for fundamental physics, it faces scalability and coupling limitations. **Circuit QED (cQED)** maps these exact same quantum optical principles onto a solid-state superconducting chip, swapping physical atoms for superconducting circuits and mirrors for transmission lines.



Key Advantages of Circuit QED

- **Massive Artificial Dipole Moments:** Natural atoms have small dipole moments bounded by the Bohr radius ($d \sim ea_0$). Superconducting qubits are macroscopic circuits where Cooper pairs tunnel across micrometric junctions. This gives them artificial dipole moments that are 10^4 to 10^5 times larger than natural atoms.
- **Ultra-Tight Field Confinement:** Instead of a 3D optical cavity, cQED uses 1D coplanar waveguide (CPW) micro-resonators. The cross-sectional mode area drops from square micrometers to square nanometers, confining the electric field and maximizing E_{zpf} .
- **Deep Strong Coupling** ($g \gg \kappa, \gamma$): Due to the massive dipole moment and small mode volume, the coupling rate g easily outpaces the qubit decay rate γ and cavity loss rate κ . This places the system deep within the **strong coupling regime**.
- **Engineered Architecture:** Features like transition frequencies, coupling constants, and decay rates are completely customizable using standard lithographic cleanroom processing.

Typical Implementation Layout

- **The Cavity:** A superconducting coplanar transmission-line resonator (often made of Niobium or Aluminum) acting as a microwave cavity.
- **The Artificial Atom:** A transmon qubit placed at an electric field anti-node (maximum) inside the resonator.

This integrated layout allows for coherent control of the qubit using microwave pulses, as well as high-fidelity state readout by monitoring how the qubit alters the resonator's properties.

4. Modeling the Resonator as a Single Mode

One-Mode Approximation

A physical coplanar waveguide resonator acts as a distributed-element transmission line with an infinite series of harmonic modes ($\omega_r, 2\omega_r, 3\omega_r, \dots$). However, if the mode spacing is much larger than the linewidth of each mode:

$$\kappa_m \ll |\omega_{m+1} - \omega_m|$$

And if the system's drive frequencies are restricted near the fundamental frequency ω_r , we can drop the higher-order modes. This reduces the complex transmission line to a single-mode approximation.

Lumped-Element LC Analogue

Under this approximation, the fundamental mode of the resonator maps directly to an equivalent lumped-element parallel LC resonant circuit:

$$\omega_r = \frac{1}{\sqrt{L_r C_r}}$$

Where L_r and C_r are the equivalent total inductance and capacitance of the fundamental mode. The loss rate κ (the full-width at half-maximum of the resonator's spectral line) is determined by its internal and external Quality Factor Q :

$$\kappa = \frac{\omega_r}{Q}$$

Zero-Point Fluctuations

Quantizing the LC circuit reveals that the charge $\hat{Q}$ and flux $\hat{\Phi}$ operators act as canonical conjugate variables ($[\hat{\Phi}, \hat{Q}] = j\hbar$). Even in its absolute ground state

$|0\rangle$, the resonator maintains zero-point energy, which translates to a fluctuating vacuum voltage across the capacitor:

$$V_{\text{zpf}} = \omega_r \Phi_{\text{zpf}} = \sqrt{\frac{\hbar \omega_r}{2C_r}}$$

In typical superconducting chips, this zero-point voltage fluctuation is approximately $V_{\text{zpf}} \sim 1 \mu\text{V}$. While this value sounds small, when applied across a sub-micron gate gap, it creates a local electric field strong enough to drive artificial atoms.

5. Total System Hamiltonian

The complete Hamiltonian of a circuit QED system combines the energy of the isolated resonator, the isolated transmon qubit, and the capacitive interaction zone between them:

$$H = H_r + H_q + H_{\text{int}}$$

Resonator Hamiltonian (H_r)

Expressed using its continuous conjugate lumped variable operators (Charge $\hat{Q}_r$ and Flux $\hat{\Phi}_r$):

$$H_r = \frac{\hat{Q}_r^2}{2C_r} + \frac{\hat{\Phi}_r^2}{2L_r}$$

By introducing the standard boson creation ($\hat{a}^\dagger$) and annihilation ($\hat{a}$) operators, where $\hat{a}^\dagger \hat{a} = \hat{n}$ represents the photon number operator, the quantized harmonic oscillator simplifies to:

$$H_r = \hbar \omega_r \left(\hat{a}^\dagger \hat{a} + \frac{1}{2} \right)$$

Transmon Qubit Hamiltonian (H_q)

A transmon is a non-linear Josephson oscillator defined by its charging energy E_c and Josephson tunneling energy E_J :

$$H_q = 4E_c \hat{n}^2 - E_J \cos \hat{\varphi}$$

Where $\hat{n}$ is the Cooper-pair number operator and $\hat{\varphi}$ is the gauge-invariant phase operator across the junction. The non-linear $\cos \hat{\varphi}$ potential introduces

anharmonicity, shifting the higher energy levels so the $0 \rightarrow 1$ transition can be isolated from the $1 \rightarrow 2$ transition.

Total Un-Truncated Expression

Combining these terms yields the total combined system Hamiltonian before any two-level truncation:

$$H = \left(\frac{\hat{Q}_r^2}{2C_r} + \frac{\hat{\Phi}_r^2}{2L_r} \right) + (4E_c \hat{n}^2 - E_J \cos \hat{\varphi}) + H_{\text{int}}$$

6. Two-Level System Approximation

Qubit Truncation

Because the transmon's anharmonicity isolates the ground state $|0\rangle$ and first excited state $|1\rangle$, we can truncate the infinite-dimensional transmon Hilbert space down to a simple two-level system. Using the standard Pauli spin operators $(\sigma_x, \sigma_y, \sigma_z)$, the isolated qubit Hamiltonian simplifies to:

$$H_{\text{q}} \approx \frac{\hbar\omega_{01}}{2} \sigma_z$$

Where $\sigma_z = |1\rangle\langle 1| - |0\rangle\langle 0|$.

The Interaction Term (H_{int})

The physical link between the transmon and the resonator is a coupling capacitor C_g . The interaction energy depends on the voltage across the resonator and the charge on the qubit:

$$H_{\text{int}} = -\hat{q}_{\text{qubit}} \cdot \hat{V}_{\text{resonator}}$$

Expressing the resonator voltage via creation/annihilation operators yields $\hat{V}_r = V_{\text{zpf}}(\hat{a}^\dagger + \hat{a})$. Similarly, the qubit's charge operator maps to the atomic raising and lowering operators (σ_+, σ_-) . This gives the interaction term its standard form:

$$H_{\text{int}} = \hbar g (\hat{a}^\dagger + \hat{a})(\sigma_- + \sigma_+)$$

Where g represents the fundamental vacuum Rabi coupling strength. Expanding this product yields four distinct interaction terms:

$$H_{\text{int}} = \hbar g (\hat{a}^\dagger \sigma_- + \hat{a} \sigma_+ + \hat{a}^\dagger \sigma_+ + \hat{a} \sigma_-)$$

7. The Jaynes–Cummings Model

The Rotating Wave Approximation (RWA)

To simplify the interaction expression, we transform the system into an uncoupled rotating interaction frame. In this frame, the terms oscillate at different rates:

- **Coherent Exchange Terms** ($\hat{a}\sigma_+$ and $\hat{a}^\dagger\sigma_-$): These terms oscillate at the difference frequency ($\omega_{01} - \omega_r$). Near resonance, this difference is small, meaning these terms oscillate slowly and exert a measurable net force on the system.
- **Counter-Rotating Terms** ($\hat{a}\sigma_-$ and $\hat{a}^\dagger\sigma_+$): These terms oscillate rapidly at the sum frequency ($\omega_{01} + \omega_r$). Over measurable timescales, these rapid oscillations average out to zero.

Rejection of these fast counter-rotating terms is called the **Rotating Wave Approximation (RWA)**. Dropping them yields the iconic **Jaynes–Cummings Hamiltonian**:

$$H_{JC} = \hbar\omega_r\hat{a}^\dagger\hat{a} + \frac{\hbar\omega_{01}}{2}\sigma_z + \hbar g (\hat{a}^\dagger\sigma_- + \hat{a}\sigma_+)$$

Physical Mechanics

The surviving interaction terms describe the coherent, bi-directional exchange of energy between the artificial atom and the confined field:

- $\hbar g\hat{a}^\dagger\sigma_-$: The qubit drops from the excited state to the ground state ($|1\rangle \rightarrow |0\rangle$) and emits a single microwave photon ($\hat{a}^\dagger$) into the resonator.
- $\hbar g\hat{a}\sigma_+$: The qubit absorbs a single microwave photon ($\hat{a}$) from the resonator and transitions from the ground state to the excited state ($|0\rangle \rightarrow |1\rangle$).

This coherent energy exchange occurs at the **Vacuum Rabi Frequency** $2g$. The Jaynes-Cummings model provides the standard description for both cavity and circuit QED within the weak-coupling, near-resonant regime ($\Delta \rightarrow 0$).

8. The Dispersive Regime

Detuning Condition

When the qubit's transition frequency and the resonator's base frequency are intentionally detuned far apart, the energy difference is defined as:

$$\Delta = \omega_{01} - \omega_r$$

The system enters the **dispersive regime** when this detuning significantly exceeds the coupling and loss rates:

$$|\Delta| \gg g, \kappa, \gamma$$

Consequence: Dressed Frequencies

Because $|\Delta| \gg g$, there is not enough energy to directly swap excitations between the qubit and the resonator. Real photon absorption and emission are suppressed. Instead, virtual photon exchanges induce a **state-dependent frequency shift**.

Applying a second-order perturbation (Schrieffer–Wolff transformation) to the Jaynes-Cummings Hamiltonian drops the direct exchange terms and yields the effective Dispersive Hamiltonian:

$$H_{\text{disp}} \approx \hbar (\omega_r + \chi \sigma_z) \hat{a}^\dagger \hat{a} + \frac{\hbar \omega'_{01}}{2} \sigma_z$$

Where $\chi = \frac{g^2}{\Delta}$ is the **dispersive shift coefficient**, and $\omega'_{01} = \omega_{01} + \chi$ represents the Lamb-shifted qubit frequency.

Physical Interpretation

Grouping the terms highlights how the resonator's effective frequency depends directly on the state of the qubit:

$$H_{\text{disp}} = \hbar \underbrace{(\omega_r + \chi \sigma_z)}_{\omega_{r,\text{eff}}} \hat{a}^\dagger \hat{a} + \frac{\hbar \omega'_{01}}{2} \sigma_z$$

$$\omega_{r,\text{eff}} = \begin{cases} \omega_r + \chi, & \text{if qubit is in } |1\rangle (\sigma_z = +1) \\ \omega_r - \chi, & \text{if qubit is in } |0\rangle (\sigma_z = -1) \end{cases}$$

This state-dependent frequency splitting forms the foundation of non-destructive **dispersive readout**.

9. Dispersive Qubit Readout

The Readout Principle

To measure the qubit, a weak microwave tone is sent into the resonator at its bare resonance frequency ω_r .

Because the resonator's actual resonance frequency shifts down to $\omega_r - \chi$ (if $|0\rangle$) or up to $\omega_r + \chi$ (if $|1\rangle$), the transmission curve shifts across the frequency axis.

As a result, the reflected or transmitted microwave signal experiences a phase shift $\Delta\theta \approx \pm 2\arctan(2\chi/\kappa)$ depending on the qubit state.

This phase shift is measured using room-temperature electronics (such as the I/Q demodulation framework discussed previously) to determine the qubit state.

High-Fidelity Readout Parameters

To achieve high-fidelity readout, the state-dependent frequency shift must be larger than the cavity’s linewidth ($\chi > \kappa$), allowing the two response curves to be clearly distinguished. However, κ cannot be too small, or photons will take too long to exit the cavity, slowing down the readout process.

Operational Parameters Table

Parameter Metric	3D Waveguide Cavity Layout	2D Coplanar Waveguide Layout
Coupling Strength ($g/2\pi$)	50 – 150 MHz	100 – 300 MHz
Cavity Decay Rate ($\kappa/2\pi$)	1 – 10 kHz (High- Q storage)	5 – 10 MHz (Fast Readout)
Qubit Relaxation Time (T_1)	Up to 100 – 200 μ s	50 – 150 μ s
Readout Mechanism Type	Dispersive (Nondestructive)	Dispersive (Nondestructive)

Control vs. Readout Separation

This framework completely separates qubit control from qubit readout:

- **Qubit Control Line:** Driven near ω_{01} to induce Rabi rotations and execute single-qubit gates.
- **Qubit Readout Line:** Driven near ω_r to measure the qubit’s state via the resonator’s phase shift.

This spectral separation allows independent optimization of control and measurement paths.

10. Towards Xmon Qubits

Evolution of the Transmon

The standard transmon qubit consists of a Cooper-pair box shunted by a large capacitor to reduce its sensitivity to charge noise. Early versions used interdigital finger capacitors, which had limited coupling options and larger electric field footprints.

The Xmon Architecture

The **Xmon** variant rearranges this shunted capacitance into a cross-shaped geometry.

This cross-shaped design provides distinct architectural advantages:

- **Multi-Port Connectivity:** The four arms of the cross provide independent connectivity points. One arm couples to a readout resonator, another to a drive line, a third to an adjacent qubit for multi-qubit gates, and the fourth can be used for a dedicated flux-tuning line.
- **Scalable Footprint:** The orthogonal layout simplifies the design of 2D planar grids for large-scale quantum processors.

Magnetic Flux Tunability

By replacing the single Josephson junction with a parallel pair forming a **SQUID loop (Superconducting Quantum Interference Device)**, the qubit's transition frequency can be tuned using an external magnetic flux Φ_{ext} .

The external flux alters the effective Josephson energy E_J according to the relation:

$$E_J(\Phi_{\text{ext}}) = 2E_{J0} \left| \cos \left(\frac{\pi \Phi_{\text{ext}}}{\Phi_0} \right) \right|$$

Where $\Phi_0 = \frac{h}{2e}$ is the magnetic flux quantum. Because the qubit frequency scales as $\omega_{01} \approx \sqrt{8E_c E_J}$, changing the local magnetic flux shifts the qubit's frequency. This allows processors to quickly tune qubits into resonance for two-qubit gates or detune them to prevent unwanted cross-talk.

Compact Resonator Layouts

To package multiple qubits onto a single chip, straight coplanar waveguide resonators are wound into compact **meander line shapes**. These meanders reduce the physical area of each readout channel, enabling higher packing densities on superconducting quantum processing units (QPUs).

11. Core Mathematical Reference Sheet

Quantum Metric / Model	Mathematical Formulation
Electric Dipole Interaction	$H_{\text{int}} = -\vec{d} \cdot \vec{E}(t)$
Vacuum Field Fluctuation	$E_{\text{zpf}} = \sqrt{\frac{\hbar\omega}{2\epsilon_0 V}}$
Lumped Resonator Frequency	$\omega_r = \frac{1}{\sqrt{L_r C_r}}$
Jaynes-Cummings Expression	$H_{\text{JC}} = \hbar\omega_r \hat{a}^\dagger \hat{a} + \frac{\hbar\omega_{01}}{2} \sigma_z + \hbar g (\hat{a}^\dagger \sigma_- + \hat{a} \sigma_+)$
Dispersive Regime Shift	$H_{\text{disp}} \approx \hbar(\omega_r + \chi\sigma_z) \hat{a}^\dagger \hat{a} + \frac{\hbar\omega'_{01}}{2} \sigma_z \quad \left(\text{where } \chi = \frac{g^2}{\Delta}\right)$
Flux-Tuned Josephson Energy	$E_J(\Phi_{\text{ext}}) = 2E_{J0} \cos(\pi\Phi_{\text{ext}}/\Phi_0) $

Qubit control

1. Deep Dive into Readout in cQED

The core challenge in circuit Quantum Electrodynamics (cQED) is measuring a qubit without destroying its fragile quantum state.

This is done by mapping the qubit state onto a microwave resonator physically coupled to it.

The Jaynes-Cummings Hamiltonian

The interaction between a two-level atom (qubit) and a quantized electromagnetic field (resonator) is described by the Jaynes-Cummings Hamiltonian:

$$H = \omega_r a^\dagger a + \frac{\omega_{01}}{2} \sigma_z - g (a^\dagger \sigma_- + a \sigma_+)$$

- $\omega_r a^\dagger a$: energy of the photons in the resonator
- $\frac{\omega_{01}}{2} \sigma_z$: energy of the qubit
- $g(a^\dagger \sigma_- + a \sigma_+)$: interaction term, with coupling strength g

The Dispersive Regime ($\Delta \gg g, \kappa$)

To avoid direct energy exchange between the qubit and the resonator, the system is operated in the dispersive regime, where they are highly detuned.

The detuning is:

$$\Delta = \omega_{01} - \omega_r$$

When $\Delta \gg g$, the Hamiltonian can be approximated as:

$$H \approx (\omega_r + \chi\sigma_z) a^\dagger a + \frac{\omega_{01}}{2} \sigma_z$$

where:

$$\chi \approx \frac{g^2}{\Delta}$$

This means the resonator frequency depends on the qubit state:

- if the qubit is in $|0\rangle$, the cavity responds at $\omega_r - \chi$
- if the qubit is in $|1\rangle$, the cavity responds at $\omega_r + \chi$

By sending a microwave probe tone at ω_r through the feedline, the reflected or transmitted signal acquires a phase shift. Measuring that phase shift reveals the qubit state.

2. The Purcell Effect and Purcell Filters

The Problem: Purcell Decay

Because the resonator is coupled to a feedline with decay rate κ , the qubit can leak energy into the environment through the resonator. This unwanted relaxation is called Purcell decay:

$$\gamma_p = \left(\frac{g}{\Delta}\right)^2 \kappa$$

A larger κ improves readout speed, but also increases γ_p . This creates a trade-off between measurement speed and qubit lifetime.

The Solution: The Purcell Filter

A Purcell filter is placed between the resonator and the transmission line.

It acts as a frequency-selective bandpass filter:

- at ω_r : transparent, allowing readout photons to pass
- at ω_q : opaque, suppressing qubit relaxation

This allows fast readout without strongly reducing qubit coherence.

3. Controlling the Qubit

To perform quantum logic gates, the qubit state must be rotated on the Bloch sphere. For a transmon, there are two main control methods.

Method A: Microwave Drive

This is used for X and Y rotations.

- A classical microwave voltage $V_d(t)$ is sent through a control line
- The line is coupled to the transmon by a small capacitor C_d
- The oscillating electric field couples to the qubit dipole moment and drives transitions

The coupling capacitor must be small to avoid strong dissipation into the $50\ \Omega$ control environment.

Method B: Flux Control

This is used to change the qubit frequency and implement Z gates or tune qubits into resonance.

A flux-tunable transmon uses a SQUID loop instead of a single Josephson junction. A magnetic flux Φ changes the Josephson energy E_J , and therefore the qubit frequency:

$$\omega_q \propto \sqrt{8E_c E_J}$$

4. Summary Table: Readout vs. Drive

Feature	Readout (Measurement)	Drive (Control)
Objective	Determine whether the qubit is in $ 0\rangle$ or $ 1\rangle$	Rotate the qubit state on the Bloch sphere
Coupling type	Dispersive capacitive coupling to a resonator	Direct capacitive or inductive coupling to a control line
Key physics	Jaynes-Cummings dispersive shift (χ)	Dipole transitions or SQUID tuning
Protection mechanism	Purcell filter	Small coupling capacitor C_d

5. Mathematical Derivation of Microwave Control

Step 1: Linearizing the Circuit Hamiltonian

Consider a classical transmon qubit connected to a microwave drive line via a small coupling capacitor C_d . The drive line applies a time-dependent voltage $V_d(t)$.

Near the ground state, the Josephson junction can be linearized and treated approximately as a harmonic oscillator. The Hamiltonian is:

$$H = \frac{Q^2}{2C_\Sigma} + \frac{\Phi^2}{2L} + \frac{C_d}{C_\Sigma} V_d(t) Q$$

where:

- Q and Φ are charge and flux operators
- $C_\Sigma = C + C_d$ is the total capacitance
- the third term is the driving Hamiltonian

Step 2: Quantization and Truncation

The charge operator is written as:

$$Q = -iQ_{\text{zpf}}(a - a^\dagger)$$

where:

$$Q_{\text{zpf}} = \sqrt{\frac{\hbar}{2Z_0}}, \quad Z_0 = \sqrt{\frac{L}{C_\Sigma}}$$

Substituting into the Hamiltonian gives:

$$H = \omega_q a^\dagger a - \frac{C_d}{C_\Sigma} V_d(t) iQ_{\text{zpf}}(a - a^\dagger)$$

Truncation to a Two-Level System

Because the transmon is weakly anharmonic, only the lowest two levels are kept:

$$a \rightarrow \sigma_-, \quad a^\dagger \rightarrow \sigma_+$$

Using:

$$\sigma_x = \sigma_+ + \sigma_-, \quad \sigma_y = i(\sigma_- - \sigma_+)$$

the Hamiltonian becomes:

$$H = -\frac{\omega_q}{2} \sigma_z + \Omega V_d(t) \sigma_y$$

where:

$$\Omega = \frac{C_d}{C_\Sigma} Q_{\text{zpf}}$$

This shows that the drive line naturally produces rotations about the y axis.

6. Moving to the Rotating Frame

The Hamiltonian

$$H = -\frac{\omega_q}{2} \sigma_z + \Omega V_d(t) \sigma_y$$

is written in the laboratory frame, where the qubit precesses rapidly around the z axis.

To simplify control, we apply the unitary transformation:

$$U(t) = e^{i\omega_q t \sigma_z / 2}$$

In the rotating frame:

$$U(t) \sigma_y U^\dagger(t) = \cos(\omega_q t) \sigma_y - \sin(\omega_q t) \sigma_x$$

This means that a y -drive in the lab frame becomes a combination of x and y rotations in the rotating frame.

7. IQ Modulation and the Rotating Wave Approximation

A realistic drive signal is:

$$V_d(t) = V_0 s(t) \sin(\omega_d t + \phi)$$

Using trigonometric identities, this can be written as:

$$V_d(t) = V_0 s(t) [I \sin(\omega_d t) - Q \cos(\omega_d t)]$$

where:

$$I = \cos \phi, \quad Q = \sin \phi$$

Rotating Wave Approximation

After transforming to the rotating frame, terms appear at:

1. $\omega_d - \omega_q$
2. $\omega_d + \omega_q$

The high-frequency sum terms average to zero, so they are dropped. This is the Rotating Wave Approximation.

If the drive is on resonance, $\delta\omega = \omega_q - \omega_d = 0$, then:

$$H'_d = -\frac{\Omega}{2} V_0 s(t) (I\sigma_x + Q\sigma_y)$$

X and Y Gates

- *X* gate: set $\phi = 0$, so $I = 1$, $Q = 0$
- *Y* gate: set $\phi = \pi/2$, so $I = 0$, $Q = 1$

Rotation Angle

The total rotation angle is:

$$\Theta(t) = -\Omega V_0 \int_0^t s(t') dt'$$

A perfect π -pulse is obtained by calibrating the pulse duration or amplitude so that the pulse area equals π .

8. Hardware: I/Q Modulation

The drive signal is produced using I/Q modulation:

$$V_d(t) = V_0 s(t) [I \sin(\omega_d t) - Q \cos(\omega_d t)]$$

Components

1. **Local Oscillator (LO)**: generates a high-frequency microwave tone
2. **Arbitrary Waveform Generator (AWG)**: outputs low-frequency envelopes $I(t)$ and $Q(t)$
3. **IQ Mixer**: combines the LO with the AWG outputs to create the final microwave drive

The resulting drive frequency is:

$$\omega_d = \omega_{\text{LO}} + \omega_{\text{AWG}}$$

Frequency Multiplexing

Multiple qubits or resonators can be addressed by sending several intermediate frequencies through the same physical line.

9. What Is a Virtual Z-Gate?

A physical Z gate rotates the qubit state around the z axis, changing its relative phase.

Traditionally, this was done using a flux pulse. However, flux pulses add noise, calibration errors, and distortion.

A **virtual Z-gate** performs the same phase update with no physical pulse. Instead, the software updates the reference phase of future control pulses.

10. Mathematics of Virtual Z-Gates

From the rotating-frame drive Hamiltonian:

$$H'_d = -\frac{\Omega}{2} V_0 s(t) (I\sigma_x + Q\sigma_y)$$

a virtual Z_ϕ gate rotates the control coordinates:

$$I' = I \cos \phi - Q \sin \phi$$

$$Q' = I \sin \phi + Q \cos \phi$$

Example: Turning an X-Gate into a Y-Gate

Suppose an X -gate is followed by a virtual Z -gate of $\phi = \pi/2$.

Before the virtual gate:

- $I = 1$
- $Q = 0$

After the transformation:

$$I' = (1) \cos\left(\frac{\pi}{2}\right) - (0) \sin\left(\frac{\pi}{2}\right) = 0$$

$$Q' = (1) \sin\left(\frac{\pi}{2}\right) + (0) \cos\left(\frac{\pi}{2}\right) = 1$$

So the pulse is effectively shifted from the I channel to the Q channel.

Why Virtual Z-Gates Are Useful

Feature	Physical Z-Gate	Virtual Z-Gate
Execution time	Finite	Instantaneous
Fidelity	Limited by flux noise	Effectively perfect
Hardware required	Fast flux line	Software phase update
Decoherence during gate	Yes	No

Virtual Z-gates are standard in modern superconducting quantum processors.

Qubit Characterization

1. Closed vs. Open Quantum Systems

The fundamental challenge in quantum computing is moving from an ideal theoretical system to a real physical system.

- **Closed systems (ideal):** the qubit is perfectly isolated. Its evolution is deterministic and fully described by its Hamiltonian.
- **Open systems (realistic):** the qubit interacts with its environment. These uncontrolled interactions introduce noise and lead to loss of quantum information.

2. Types of Noise: Systematic vs. Stochastic

Environmental noise is typically divided into two categories.

Noise Type	Characteristics	Example	Solution
Systematic noise	Non-random, reproducible errors caused by fixed control or readout flaws	A microwave pulse consistently over-rotates a qubit because of poor calibration	Calibration or hardware improvement
Stochastic noise	Random, unpredictable fluctuations from parameters coupled to the qubit	Thermal noise, charge noise, flux noise, fluctuating magnetic or electric fields	Hard to eliminate; causes decoherence

3. The Bloch Sphere Representation

The Bloch sphere is a geometric tool used to visualize the state of a two-level quantum system.

Geometric breakdown

- **Pure states:** the qubit is fully known and lies on the surface of the sphere, with $|\vec{a}| = 1$
- **Z-axis:** connects the north pole $|0\rangle$ and south pole $|1\rangle$
- **X-Y plane:** represents superposition states

A general qubit state can be written as:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

and mapped to spherical coordinates as:

$$\begin{pmatrix} \sin \theta \cos \phi \\ \sin \theta \sin \phi \\ \cos \theta \end{pmatrix}$$

with:

$$|\psi\rangle = \cos\left(\frac{\theta}{2}\right)|0\rangle + e^{i\phi}\sin\left(\frac{\theta}{2}\right)|1\rangle$$

Reference frames

- **Laboratory frame:** the physical frame. A superposition state precesses around the z -axis at frequency ω_{01}
 - **Rotating frame:** rotates at the qubit frequency ω_{01} , making the Bloch vector appear stationary
-

4. The Density Matrix and Real-World States

In noisy systems, a state vector $|\psi\rangle$ is not always enough. We use the **density matrix** $\hat{\rho}$.

- **Pure state:**

$$\hat{\rho} = |\psi\rangle\langle\psi|$$

and

$$\text{Tr}(\hat{\rho}^2) = 1$$

- **Mixed state:** stochastic noise forces the qubit into a statistical mixture of states. The Bloch vector has length $|\vec{a}| < 1$ and lies inside the Bloch sphere. In both pure and mixed cases, the density matrix is normalized so that $\text{Tr}(\hat{\rho}) = 1$.

A density matrix can be written in terms of the Bloch vector $\vec{a}$ as:

$$\hat{\rho} = \frac{1}{2} (I + \vec{a} \cdot \vec{\sigma})$$

This compact form describes the qubit state, but it does not by itself include the action of environmental noise.

To account for weak coupling to a Markovian bath, one can use open-system models such as the **Bloch-Redfield** formalism, which adds dissipative terms to the density-matrix evolution.

5. Decoherence

Decoherence is the process by which a qubit loses its quantum properties due to interaction with the environment.

In superconducting qubits such as transmons, decoherence is usually split into:

1. **Longitudinal relaxation** (T_1): loss of energy
2. **Pure dephasing** (T_ϕ): loss of phase information

Together, these determine the transverse relaxation time T_2 .

6. Longitudinal Relaxation (T_1)

Also called energy relaxation or depolarization, this is the process in which the qubit decays from $|1\rangle$ to $|0\rangle$.

Physics

- Caused by transverse noise, i.e. fluctuations along x or y
- This noise couples to the qubit transition frequency ω_q

The total relaxation rate is:

$$\Gamma_1 \equiv \frac{1}{T_1} = \Gamma_{1\downarrow} + \Gamma_{1\uparrow}$$

Because superconducting qubits operate in very cold dilution refrigerators, the thermal excitation rate is negligible:

$$\Gamma_{1\uparrow} = \Gamma_{1\downarrow} e^{-\frac{\hbar\omega_q}{k_B T}} \approx 3 \times 10^{-7} \Gamma_{1\downarrow}$$

So T_1 is dominated by downward decay.

How T_1 is measured

1. Apply an X_π pulse to prepare the qubit in $|1\rangle$
2. Wait for a variable delay time τ
3. Measure the qubit along the z axis
4. Repeat many times and fit the decay

The excited-state population follows:

$$P_{|1\rangle}(\tau) \propto e^{-\tau/T_1}$$

At $\tau = T_1$, the population has decayed to about 36.8% of its initial value.

7. Pure Dephasing (T_ϕ)

Pure dephasing is the loss of relative phase in a superposition state, without energy loss.

Physics

- Associated with noise that causes the qubit frequency to fluctuate, inducing precession along the z -axis of the state vector
- Caused by **longitudinal noise**, i.e. fluctuations along the z axis, which add a time-dependent shift $\delta\omega(t)$ to the qubit frequency
- On resonance, in the rotating frame, the ideal H' is zero; longitudinal noise makes the Bloch vector precess forward or backward in the rotating frame depending on $\delta\omega(t)$
- It is represented by the rate Γ_ϕ , which describes depolarization in the x - y plane of the Bloch sphere

In the rotating frame, the Bloch vector wanders around the equator. When averaged over many repetitions, the phase randomization shrinks the vector toward the center of the sphere.

- Dephasing is a stochastic effect, and the transverse relaxation parameter captures it
- Ensemble averaging over random phase shifts makes the average state vector shorter than 1 in length
- This averaging causes decay of transverse coherence
- Unlike energy relaxation, pure dephasing is not a resonant phenomenon; noise at any frequency can modify the qubit frequency and cause dephasing

8. Transverse Relaxation (T_2)

The total decay of superposition coherence is given by:

$$\Gamma_2 = \frac{1}{T_2^*} = \frac{\Gamma_1}{2} + \Gamma_\phi = \frac{1}{2T_1} + \frac{1}{T_\phi}$$

Why $\Gamma_1/2$ appears

Energy relaxation is also a phase-breaking process. Once the qubit falls to $|0\rangle$, any phase information on the equator is lost. That is why $1/(2T_1)$ contributes to the total decoherence rate.

9. Measuring T_2 : Ramsey vs. Spin Echo

Two main protocols are used to measure coherence.

Protocol A: Ramsey Measurement

1. Apply an $X_{\pi/2}$ pulse to place the qubit on the equator
2. Wait for a time τ with a small detuning $\delta\omega$
3. Apply a second $X_{\pi/2}$ pulse
4. Read out the qubit

The result is Ramsey fringes, which decay with a characteristic time T_2^* .

This method is very sensitive to slow, low-frequency noise.

Protocol B: Spin Echo

This is the same as Ramsey, but an X_π pulse is inserted halfway through the waiting time.

- The π pulse reverses the phase accumulation
- Slow noise that acted during the first half is refocused in the second half

As a result, the coherence time measured with spin echo is typically larger than T_2^* .

10. Takeaways for Transmon Qubits

- For transmon qubits, pure dephasing T_ϕ is often relatively small
- In practice, coherence is usually limited mainly by energy relaxation T_1

- A central goal of quantum engineering is to maximize both T_1 and T_2 so that quantum circuits can run before information decays
-

11. Example Measurements

These experiments are commonly used in the qubit characterization workflow.

Time of Flight (TOF)

- First calibration step in the characterization chain
- Used to determine the propagation delay between readout pulse generation and the arrival of the corresponding signal at the acquisition input in the RFSoc
- Applied to all subsequent measurements

Resonator Spectroscopy

- Used to identify the resonator's fundamental frequency
- Obtained by sweeping the readout tone across one or more frequency ranges and measuring the transmitted signal amplitude
- In this experiment the qubit is decoupled from the resonator, so high readout power is used
- No drive pulse is generated, therefore the qubit is in principle completely unaffected
- Sweeps are typically started over a broad frequency range and then narrowed around the peak

Resonator Punchout

- Used to identify the optimal readout configuration in the dispersive regime
- Consists of a cartesian sweep over readout gain and frequency, providing a map of the device power response
- At high readout power the Josephson junction nonlinearity is not observed, revealing the bare resonance of the resonator
- As the power is reduced, the coupling becomes significant and the resonance shifts toward the dispersive readout frequency
- The dark absorption region at low power corresponds to the dressed resonance

Drive Spectroscopy

- Also called qubit spectroscopy
- A two-tone experiment: a drive pulse is sent through the qubit control line and is immediately followed by the readout tone
- The goal of the drive frequency sweep is to calibrate the drive pulse to the qubit transition frequency

- A high number of shots is required because the signal-to-noise ratio is very low

Rabi Oscillations

- Used to identify the drive amplitude required to calibrate the π pulse
- Three fundamental parameters must be tuned: drive frequency, drive duration, and drive gain
- The experiment can be implemented as a drive-gain sweep or a drive-duration sweep
- The sweeping pulse is modulated at the previously identified transition frequency and is sent to the qubit to produce a full excitation
- For finer characterization, a two-dimensional cartesian sweep over both drive duration and gain can be used
- Once the π pulse is calibrated, the $\pi/2$ pulse is obtained by halving the pulse energy

Lesson 1: Introduction to Transmission Line Theory

1. Defining the Regime: When Does a Wire Become a Transmission Line?

In traditional lumped-element circuit theory (e.g., standard Kirchhoff's laws), we assume that a change in voltage or current at one point in a circuit is felt everywhere instantaneously. This assumption holds true only when the physical size of the system, L , is significantly smaller than the wavelength, λ , of the signal propagating through it.

The wavelength of an electromagnetic signal depends on its frequency f and the speed of light c within the medium:

$$\lambda = \frac{c}{f}$$

To determine how to model a circuit, we analyze its **electrical length**, defined by the ratio $\frac{L}{\lambda}$:

- **Quasistatic / Lumped Parameter Regime** ($\frac{L}{\lambda} \ll 1$): The circuit elements are much smaller than the wavelength. Time delays across the physical components are negligible, allowing us to use standard lumped element models (individual resistors, capacitors, and inductors).
- **Resonance / Distributed Parameter Regime** ($\frac{L}{\lambda} \sim 1$): The signal wavelength is comparable to or smaller than the physical dimensions of

the circuit. The phase of the voltage and current varies across the physical length of the wire. The wire can no longer be treated as a single ideal node; it must be modeled as a **Transmission Line (TL)**.

- **Optical Regime** ($\frac{\ell}{\lambda} \gg 1$): The wavelength is vastly smaller than the system dimensions, and the system behaves according to geometric optics.

Example: Consider a high-frequency system operating at 5 GHz. The wavelength in a vacuum is $\lambda = \frac{3 \times 10^8 \text{ m/s}}{5 \times 10^9 \text{ Hz}} = 6 \text{ cm}$. In specialized systems like a quantum computing cryostat operating at 4 K, a signal line of just a few centimeters is comparable to λ . Therefore, distributed effects dominate, and transmission line theory must be applied.

2. Distributed Parameter Representation

When a line has a physical width w and length ℓ such that $w \ll \ell$, the voltage $V(z, t)$ and current $i(z, t)$ become continuous functions of both **time** (t) and **position** (z) along the line.

To model this mathematically, we divide the continuous transmission line into infinitesimally small differential sections of length Δz . Each segment is modeled using a combination of four primary distributed primary constants, specified **per unit length**:

- R_0 : Series resistance per unit length (Ω/m), representing conductor losses.
- L_0 : Series inductance per unit length (H/m), representing the magnetic field buildup around the conductors.
- C_0 : Shunt capacitance per unit length (F/m), representing the electric field coupling between the conductors.
- G_0 : Shunt conductance per unit length (S/m), representing dielectric leakage losses between the conductors.

3. Deriving the Telegrapher Equations

By applying Kirchhoff's Voltage Law (KVL) and Kirchhoff's Current Law (KCL) to a single differential slice Δz of the transmission line, we can mathematically formalize how voltage and current drop across an infinitesimal distance.

Kirchhoff's Voltage Law (KVL)

The difference in voltage between the output and input of the segment is equal to the voltage drop across the series impedance elements ($R_0 \Delta z$ and $L_0 \Delta z$):

$$\Delta V(z, t) = V(z + \Delta z, t) - V(z, t) = -R_0 \Delta z \cdot i(z, t) - L_0 \Delta z \cdot \frac{\partial i(z, t)}{\partial t}$$

Kirchhoff's Current Law (KCL)

The difference in current between the output and input of the segment equals the current leaking through the shunt admittance elements ($G_0\Delta z$ and $C_0\Delta z$):

$$\Delta i(z, t) = i(z + \Delta z, t) - i(z, t) = -G_0\Delta z \cdot V(z + \Delta z, t) - C_0\Delta z \cdot \frac{\partial V(z + \Delta z, t)}{\partial t}$$

Taking the Continuous Limit ($\Delta z \rightarrow 0$)

Dividing both expressions by Δz and taking the limit as $\Delta z \rightarrow 0$ yields the first-order differential equations known as the **Telegrapher Equations**:

$$\begin{aligned}\frac{\partial V(z, t)}{\partial z} &= -R_0 \cdot i(z, t) - L_0 \cdot \frac{\partial i(z, t)}{\partial t} \\ \frac{\partial i(z, t)}{\partial z} &= -G_0 \cdot V(z, t) - C_0 \cdot \frac{\partial V(z, t)}{\partial t}\end{aligned}$$

Decoupling into the Wave Equation

To solve for $V(z, t)$ and $i(z, t)$ independently, we differentiate the first Telegrapher Equation with respect to z and substitute the second. This decouples the relationships into second-order partial differential equations:

$$\begin{aligned}\frac{\partial^2 V(z, t)}{\partial z^2} &= R_0 G_0 \cdot V(z, t) + (L_0 G_0 + R_0 C_0) \frac{\partial V(z, t)}{\partial t} + L_0 C_0 \frac{\partial^2 V(z, t)}{\partial t^2} \\ \frac{\partial^2 i(z, t)}{\partial z^2} &= L_0 C_0 \frac{\partial^2 i(z, t)}{\partial t^2} + (L_0 G_0 + R_0 C_0) \frac{\partial i(z, t)}{\partial t} + R_0 G_0 \cdot i(z, t)\end{aligned}$$

4. The Lossless Transmission Line Case

For high-frequency components or highly optimized structures, losses can often be neglected ($R_0 = 0$, $G_0 = 0$). The wave equations simplify significantly:

$$\begin{aligned}\frac{\partial^2 V(z, t)}{\partial z^2} &= L_0 C_0 \frac{\partial^2 V(z, t)}{\partial t^2} \\ \frac{\partial^2 i(z, t)}{\partial z^2} &= L_0 C_0 \frac{\partial^2 i(z, t)}{\partial t^2}\end{aligned}$$

These represent classic **second-order linear wave equations**. The general solutions are combinations of arbitrary wave shapes traveling in opposite directions:

- **Forward-traveling wave (+z direction):** $g_+(z, t) = g(t - \sqrt{L_0 C_0} z)$
- **Backward-traveling/Reflected wave (-z direction):** $g_-(z, t) = f(t + \sqrt{L_0 C_0} z)$

Propagation Velocity

To track how fast a fixed point on the wave packet moves, consider a forward-traveling wave at time t_0 , and again after a short time step δt . For the wave profile to remain identical, the position must shift by δz_+ , satisfying:

$$t_0 - \sqrt{L_0 C_0} z = (t_0 + \delta t) - \sqrt{L_0 C_0} (z + \delta z_+)$$

$$\delta t = \sqrt{L_0 C_0} \delta z_+ \implies \delta z_+ = \frac{\delta t}{\sqrt{L_0 C_0}}$$

Taking the limit yields the phase velocity V :

$$V = \left| \lim_{\delta t \rightarrow 0} \frac{\delta z}{\delta t} \right| = \frac{1}{\sqrt{L_0 C_0}}$$

5. Characteristic Impedance

The total voltage and current everywhere along the line are the linear superpositions of their forward and backward components:

$$V(z, t) = V_+(z, t) + V_-(z, t) = g(t - \sqrt{L_0 C_0} z) + f(t + \sqrt{L_0 C_0} z)$$

$$i(z, t) = i_+(z, t) + i_-(z, t)$$

Substituting these wave solutions back into the primary lossless Telegrapher relationship ($\frac{\partial V}{\partial z} = -L_0 \frac{\partial i}{\partial t}$) links the voltage amplitudes to the current amplitudes:

$$i(z, t) = \sqrt{\frac{C_0}{L_0}} g(t - \sqrt{L_0 C_0} z) - \sqrt{\frac{C_0}{L_0}} f(t + \sqrt{L_0 C_0} z)$$

The ratio of voltage to current for a single traveling wave is a fundamental constant of the line called the **Characteristic Impedance** (Z_0):


$$Z_0 = \sqrt{\frac{L_0}{C_0}} \quad [\Omega]$$

This gives us the standard conventions for directional current:

$$i_+(z, t) = \frac{V_+(z, t)}{Z_0}$$

$$i_-(z, t) = -\frac{V_-(z, t)}{Z_0}$$

Note on Sign Convention: The negative sign for $i_-(z, t)$ indicates that the backward-traveling current wave physically moves in the reverse direction (towards the source, $-z$).



Lesson 2: Boundary Conditions, Reflections, and Transients

1. Voltage and Current Reflection Coefficients at a Load

When a transmission line is terminated by an arbitrary load resistor R_L at position $z = \ell$, the total voltage-to-current ratio at that specific point must satisfy Ohm's law for that load:

$$\left. \frac{V(z, t)}{i(z, t)} \right|_{z=\ell} = R_L$$

Substituting the forward and backward traveling components into this boundary condition:

$$\frac{V_+(\ell, t) + V_-(\ell, t)}{\frac{1}{Z_0}V_+(\ell, t) - \frac{1}{Z_0}V_-(\ell, t)} = R_L$$

Rearranging terms to find the ratio of the reflected wave amplitude to the incident wave amplitude yields the **Voltage Reflection Coefficient** (Γ_{Lv}):

$$\Gamma_{Lv} = \frac{V_-(\ell, t)}{V_+(\ell, t)} = \frac{R_L - Z_0}{R_L + Z_0}$$

Similarly, calculating the ratio for current components yields the **Current Reflection Coefficient** (Γ_{Li}):

$$\Gamma_{Li} = \frac{i_-(\ell, t)}{i_+(\ell, t)} = -\frac{R_L - Z_0}{R_L + Z_0} = -\Gamma_{Lv}$$

2. Transient Analysis and Lattice Diagrams

When a source steps cleanly into an uncharged transmission line, the signal takes a finite transit time $T = \frac{\ell}{V}$ to reach the load. If the load is mismatched ($R_L \neq Z_0$), a reflection is generated, which travels back to the source. If the source impedance is also mismatched, it reflects *again*.

To visualize these multiple step-reflections over time and position, we construct a **Lattice Diagram**.

Step-by-Step Transient Example

Consider a circuit where an ideal step voltage source V with zero internal resistance ($R_S = 0$) connects to a transmission line of characteristic impedance Z_0 at $t = 0$. The line is terminated with a load $R_L = 3Z_0$.

1. Calculate Reflection Coefficients:

- **At the Load ($z = \ell$):** $\Gamma_L = \frac{3Z_0 - Z_0}{3Z_0 + Z_0} = \frac{1}{2}$
- **At the Source ($z = 0$):** $\Gamma_S = \frac{R_S - Z_0}{R_S + Z_0} = \frac{0 - Z_0}{0 + Z_0} = -1$

2. Tracking Steps Over Time Periods:

Time Interval	Wave Activity	Net Voltage $V(z)$	Net Current $i(z)$
$0 < t < T$	Incident wave $V_+ = V$ travels forward.	$V = V$	$i = \frac{V}{Z_0}$
$T \leq t < 2T$	Reaches load, reflects $V_- = \Gamma_L V_+ = \frac{V}{2}$ backward.	$V = V + \frac{V}{2} = \frac{3V}{2}$	$i = \frac{V}{Z_0} - \frac{V}{2Z_0} = \frac{V}{2Z_0}$
$2T < t < 3T$	Reaches source, reflects $V_+ = \Gamma_S V_- = -\frac{V}{2}$ forward.	$V = \frac{3V}{2} - \frac{V}{2} = V$	$i = \frac{V}{2Z_0} - \frac{V}{2Z_0} = 0$
$3T \leq t < 4T$	Reaches load, reflects $V_- = \Gamma_L V_+ = -\frac{V}{4}$ backward.	$V = V - \frac{V}{4} = \frac{3V}{4}$	$i = 0 + \frac{V}{4Z_0} = \frac{V}{4Z_0}$

Time Interval	Wave Activity	Net Voltage $V(z)$	Net Current $i(z)$
$4T < t < 5T$	Reaches source, reflects $V_+ = \Gamma_S V_- = \frac{V}{4}$ forward.	$V = V + \frac{V}{4} = V$	$i = \frac{V}{4Z_0} + \frac{V}{4Z_0} = \frac{V}{2Z_0}$

Over time, these discrete stepping values bounce back and forth, eventually damping out and converging to the steady-state DC distribution values.

3. Power, Return Loss, and Standing Wave Ratio (SWR)

In the sinusoidal steady state, net time-averaged power $P(z)$ at any position along the line is calculated using the complex voltage and current:

$$P(z) = \frac{1}{2} \text{Re}\{V(z) \cdot I^*(z)\}$$

Expressing this in terms of forward and backward waves (using characteristic admittance $Y_0 = \frac{1}{Z_0}$):

$$P(z) = \frac{1}{2} \text{Re}\{[V^+(z) + V^-(z)] \cdot Y_0^* \cdot [V^+(z) - V^-(z)]^*\}$$

For a lossless line, Y_0 is purely real. Expanding and simplifying eliminates cross-terms, leaving:

$$P(z) = \frac{1}{2} Y_0 |V^+(z)|^2 - \frac{1}{2} Y_0 |V^-(z)|^2 = P_{\text{incident}} - P_{\text{reflected}}$$

Factoring out the incident power gives the net power delivered in terms of the reflection magnitude:

$$P(z) = \frac{|V^+(z)|^2}{2Z_0} (1 - |\Gamma|^2)$$

Return Loss (RL)

Return Loss scales the ratio of reflected power to incident power logarithmically in decibels (dB):

$$RL = -10 \log_{10} \left| \frac{P_{\text{reflected}}}{P_{\text{incident}}} \right| = -10 \log_{10} |\Gamma|^2 = -20 \log_{10} |\Gamma| \quad [\text{dB}]$$

- **$RL = 0$ dB:** Total reflection ($|\Gamma| = 1$), which occurs with purely reactive loads (no power absorbed).
- **$RL \rightarrow \infty$ dB:** Perfect impedance match ($|\Gamma| = 0$), meaning zero power is reflected back.

Standing Wave Ratio (SWR)

When continuous sinusoidal waves reflect, the forward and backward waves interfere, creating a static spatial pattern of local voltage maxima and minima along the line. The ratio of these peaks and troughs is the Standing Wave Ratio:

$$SWR = \frac{|V|_{\max}}{|V|_{\min}} = \frac{|V^+| + |V^-|}{|V^+| - |V^-|} = \frac{1 + |\Gamma|}{1 - |\Gamma|}$$

- **Matched condition** ($|\Gamma| = 0$): $SWR = 1$ (often written as 1 : 1).
- **Total reflection** ($|\Gamma| = 1$): $SWR \rightarrow \infty$.

Engineering Benchmark: A common threshold for an acceptable impedance match is delivering at least 90% of incident power to the load ($P_{\text{reflected}}/P_{\text{incident}} \leq 0.1$). This equals a Return Loss of 10 dB:

$$|\Gamma|^2 = 0.1 \implies |\Gamma| \approx 0.316$$

$$SWR = \frac{1 + 0.316}{1 - 0.316} \approx 1.92 \approx 2 \quad (2 : 1 \text{ ratio})$$

Lesson 3: The Scattering Matrix (S-Matrix)

1. Core Concept of the S-Matrix

At high frequencies, measuring absolute voltages and currents directly at a specific internal circuit location becomes difficult due to phase sensitivity and parasitic probe interactions. To bypass this, we treat multi-port networks as a “black box” and analyze them using **Power Waves**, which quantify the incoming and outgoing power waves relative to known reference transmission lines.

For each individual port i within an N -port system, we define a characteristic reference impedance Z_{ri} (and corresponding reference admittance $Y_{ri} = \frac{1}{Z_{ri}}$). We define normalized power wave variables a_i and b_i :

$$a_i = \sqrt{Y_{ri}} V_i^+ = \frac{V_i^+}{\sqrt{Z_{ri}}} \quad (\text{Incident wave entering port } i)$$

$$b_i = \sqrt{Y_{ri}} V_i^- = \frac{V_i^-}{\sqrt{Z_{ri}}} \quad (\text{Reflected/outgoing wave leaving port } i)$$

This normalizes the net active power equation down to:

$$P_i = \frac{1}{2}|a_i|^2 - \frac{1}{2}|b_i|^2$$

2. Mathematical Definition for a 2-Port Network

For a standard two-port system, the linear relationship between the total incoming wave vectors $[a]$ and outgoing wave vectors $[b]$ is governed by a 2×2 matrix called the **Scattering Matrix** ($[S]$):

$$\begin{bmatrix} b_1 \\ b_2 \end{bmatrix} = \begin{bmatrix} S_{11} & S_{12} \\ S_{21} & S_{22} \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \end{bmatrix}$$

The general matrix element equation is defined as:

$$S_{ij} = \left. \frac{b_i}{a_j} \right|_{a_k=0 \text{ for } k \neq j}$$

To satisfy the mathematical condition $a_k = 0$, port k must not receive any incoming wave reflections. This is achieved by terminating port k with a load matching its own reference characteristic impedance ($R_L = Z_{rk}$).

Physical Meaning of the Coefficients

- S_{11} : Input reflection coefficient at Port 1, assuming Port 2 is perfectly matched ($a_2 = 0$).
 - S_{22} : Output reflection coefficient at Port 2, assuming Port 1 is perfectly matched ($a_1 = 0$).
 - S_{21} : Forward transmission gain/loss from Port 1 to Port 2.
 - S_{12} : Reverse isolation transmission from Port 2 to Port 1.
-

3. Deriving S-Parameters from Voltages

To calculate specific matrix entries, we convert the power wave variables back into explicit forward and backward traveling voltage components:

Finding S_{11} and S_{21} (Drive Port 1, Terminate Port 2)

By terminating Port 2 with Z_{r2} , any wave exiting Port 2 is fully absorbed, meaning $a_2 = 0$.


$$S_{11} = \left. \frac{b_1}{a_1} \right|_{a_2=0} = \frac{V_1^- / \sqrt{Z_{r1}}}{V_1^+ / \sqrt{Z_{r1}}} = \frac{V_1^-}{V_1^+} = \Gamma_1$$
$$S_{21} = \left. \frac{b_2}{a_1} \right|_{a_2=0} = \frac{V_2^+ / \sqrt{Z_{r2}}}{V_1^+ / \sqrt{Z_{r1}}} = \sqrt{\frac{Z_{r1}}{Z_{r2}}} \left(\frac{V_2^+}{V_1^+} \right)$$

Finding S_{22} and S_{12} (Drive Port 2, Terminate Port 1)

By terminating Port 1 with Z_{r1} , any wave exiting Port 1 is absorbed, meaning $a_1 = 0$.

$$S_{22} = \left. \frac{b_2}{a_2} \right|_{a_1=0} = \frac{V_2^- / \sqrt{Z_{r2}}}{V_2^+ / \sqrt{Z_{r2}}} = \frac{V_2^-}{V_2^+} = \Gamma_2$$
$$S_{12} = \left. \frac{b_1}{a_2} \right|_{a_1=0} = \frac{V_1^+ / \sqrt{Z_{r1}}}{V_2^+ / \sqrt{Z_{r2}}} = \sqrt{\frac{Z_{r2}}{Z_{r1}}} \left(\frac{V_1^+}{V_2^+} \right)$$

Example Application (Amplifiers): An ideal forward amplifier requires high forward gain ($S_{21} > 0$ dB) and excellent reverse isolation to prevent signals from feeding back to the input ($S_{12} \approx 0$).



Lesson 4: Analyzing Cascaded Lines and Impedance Junctions

1. S-Matrix of a Direct Characteristic Impedance Step Junction

Consider a direct step junction connection between two different transmission lines. The left line has characteristic impedance Z_1 , and the right line has characteristic impedance Z_2 . We treat this junction interface as a 2-port network.

To solve for the matrix elements, we evaluate the boundary conditions at the junction interface, denoted as point A.

Finding S_{22} and S_{12} (Driving from Line 2, with Line 1 Matched)

We drive a wave from the right side via Line 2 ($a_2 = V_A^+$). Line 1 is perfectly matched to its own impedance Z_1 , meaning no reflections return from the left ($a_1 = 0$).

The reflection seen at junction interface A from this direction is determined by the impedance step change:

$$S_{22} = \left. \frac{b_2}{a_2} \right|_{a_1=0} = \frac{V_A^-}{V_A^+} = \Gamma_A^+ = \frac{Z_1 - Z_2}{Z_1 + Z_2}$$

The total voltage at interface point A is the sum of the incident and reflected waves from Line 2:

$$V_A = V_A^+ + V_A^- = V_A^+(1 + \Gamma_A^+)$$

Substituting Γ_A^+ gives:

$$V_A = V_A^+ \left(1 + \frac{Z_1 - Z_2}{Z_1 + Z_2} \right) = V_A^+ \left(\frac{2Z_1}{Z_1 + Z_2} \right)$$

Since voltage must be continuous across the junction, this total boundary voltage V_A passes directly into Line 1 as the outgoing wave V_A^- (which maps to the power wave variable b_1).

Now, calculating the reverse transmission parameter S_{12} :

$$S_{12} = \left. \frac{b_1}{a_2} \right|_{a_1=0} = \left(\frac{V_A^-}{\sqrt{Z_1}} \right) \cdot \left(\frac{\sqrt{Z_2}}{V_A^+} \right) = \sqrt{\frac{Z_2}{Z_1}} \left(\frac{V_A^-}{V_A^+} \right)$$

Substituting the voltage transmission ratio calculated above yields:

$$S_{12} = \sqrt{\frac{Z_2}{Z_1}} \cdot \left(\frac{2Z_1}{Z_1 + Z_2} \right) = \frac{2\sqrt{Z_1 Z_2}}{Z_1 + Z_2}$$

2. Complete Scattering Matrix Form

Performing the same analysis for a signal driving from the opposite direction (from Line 1 to Line 2) completes the full system matrix:

$$S = \frac{1}{Z_1 + Z_2} \begin{bmatrix} Z_2 - Z_1 & 2\sqrt{Z_1 Z_2} \\ 2\sqrt{Z_1 Z_2} & Z_1 - Z_2 \end{bmatrix}$$

From this result, we can deduce two fundamental structural properties of the junction:

- $S_{11} \neq S_{22}$ (**Asymmetry**): Because the diagonal entries are not equal ($Z_2 - Z_1 \neq Z_1 - Z_2$), the device is **not symmetric**. The reflection coefficient changes sign depending on which side the wave approaches from.
- $S_{12} = S_{21}$ (**Reciprocity**): Because the off-diagonal transmission elements are completely identical, the network is **reciprocal**. The passive transmission efficiency is equal in both directions.

Here is an upgraded, highly structured version of your content. The explanations have been expanded for clarity, the markdown formatting has been optimized for readability, and the compact math has been converted into proper LaTeX formatting.

The Vector Network Analyzer (VNA)

A **Vector Network Analyzer (VNA)** measures both the **amplitude** and **phase** of traveling high-frequency waves. Unlike a Spectrum Analyzer, which only measures signal power, a VNA stimulates the **Device Under Test (DUT)** with a known signal and measures how the device modifies that signal. This allows it to fully characterize linear networks across a frequency spectrum using Scattering Parameters (S-parameters). It is a foundational tool for characterizing RF components, filters, antennas, resonators, and superconducting quantum bits (qubits).

Waves and S-Parameters

When high-frequency signals travel along a transmission line, they behave as waves. At any given port, we define complex traveling-wave amplitudes:

- a_n : The incident (incoming) wave at port n .
- b_n : The reflected/outgoing wave from port n .

For a standard two-port network, these waves are related by the scattering matrix (S -matrix):

$$\begin{pmatrix} b_1 \\ b_2 \end{pmatrix} = \begin{pmatrix} S_{11} & S_{12} \\ S_{21} & S_{22} \end{pmatrix} \begin{pmatrix} a_1 \\ a_2 \end{pmatrix}$$

Individual S-Parameter Definitions

- S_{11} (**Input Reflection Coefficient**): Measured as $\frac{b_1}{a_1}$ when no signal enters port 2 ($a_2 = 0$). It quantifies how much power bounces back from

the input.

- S_{21} (**Forward Transmission Coefficient**): Measured as $\frac{b_2}{a_1}$ when $a_2 = 0$. It quantifies how much signal passes through the device from input to output (gain or loss).
- S_{22} (**Output Reflection Coefficient**): Measured as $\frac{b_2}{a_2}$ when $a_1 = 0$.
- S_{12} (**Reverse Transmission Coefficient**): Measured as $\frac{b_1}{a_2}$ when $a_1 = 0$.

Logarithmic Metrics (decibels, dB)

Because RF power spans many orders of magnitude, S-parameters are typically viewed in decibels:

- **Return Loss (RL)**: Measures how well a port is matched to the transmission line. A higher return loss means less reflected power (which is ideal).

$$RL = -20 \log_{10} |S_{11}|$$

- **Insertion Loss (IL)**: Measures the attenuation of the signal as it passes through the device.


$$IL = -20 \log_{10} |S_{21}|$$

Physical Meaning and Measurement Goals

- **Impedance Mismatch (S_{11})**: If the input impedance of the DUT does not match the characteristic impedance of the system (usually 50Ω), energy is reflected. S_{11} maps this mismatch.
 - **Transmission (S_{21})**: Reveals the bandwidth, gain of amplifiers, attenuation of filters, or the sharp energy dips/peaks associated with resonators.
 - **Phase vs. Frequency**: The phase slope tells us the **group delay** (the time it takes for a signal to pass through the device). Nonlinear phase shifts cause **dispersion**, distorting complex modulated signals.
 - **Quantum Readout**: In superconducting quantum computing, a microwave resonator is coupled to a qubit. The state of the qubit ($|0\rangle$ or $|1\rangle$) slightly shifts the resonator's resonant frequency. By monitoring the tiny phase or amplitude shifts of a tone sent through (S_{21}) or bounced off (S_{11}) the resonator, the VNA determines the qubit state.
-

VNA Basic Architecture and Signal Flow

A VNA functions by generating a coherent stimulus and routing it via internal directional couplers to separate receivers.



1. **Source Generation:** A highly stable Phase-Locked Loop (PLL) synthesizer generates an RF signal. A variable attenuator controls its power level before it hits a power splitter.
2. **Reference Path:** One branch of the power splitter sends a portion of the raw signal directly to a reference receiver (RX_REF). This provides a phase and amplitude benchmark.
3. **Test Path & Forward Coupler:** The main signal is routed through an internal switch to Port 1. A directional coupler taps into any reflected waves coming back from the DUT, routing them to the RX_TEST1 receiver to calculate b_1 .
4. **Transmission & Reverse Coupler:** The wave passes through the DUT and emerges into Port 2. A second directional coupler taps this transmitted wave, routing it to the RX_TEST2 receiver to calculate b_2 .
5. **Signal Processing:** The VNA's Digital Signal Processor (DSP) compares the voltages at the test receivers against the reference receiver to calculate absolute magnitude and phase ratios.
6. **Direction Reversal:** The internal electronic switch (SW1) automatically flips the signal direction, injecting power into Port 2 and terminating Port 1, allowing the system to instantly measure S_{22} and S_{12} .

Why Coherence Matters: The source and all internal downconversion mixers share a local oscillator (LO) clock. This phase-coherence ensures that the phase relationship between the reference and test paths remains rock-solid during a frequency sweep.

Directional Couplers and Key RF Specs

A directional coupler is a passive 4-port device that samples power flowing in one specific direction while ignoring power flowing in the opposite direction.

- **Coupling Factor (dB):** The ratio of the throughput mainline power to the sampled coupled port power.

$$\text{Coupling}_{\text{dB}} = 10 \log_{10} \left(\frac{P_{\text{main}}}{P_{\text{coupled}}} \right) = 20 \log_{10} \left(\frac{V_{\text{main}}}{V_{\text{coupled}}} \right)$$

A -20 dB coupler samples 1% of the mainline power, leaving 99% to continue toward the DUT. * **Directivity (dB):** The ability of the coupler to isolate forward-traveling waves from reverse-traveling waves. If a coupler has poor directivity, a portion of the forward incident wave will leak into the reflection receiver, creating a massive systematic error in your S_{11} measurements. High directivity is the hallmark of a high-quality VNA. * **Insertion Loss:** The inevitable loss of signal power along the mainline path caused by transmission line attenuation and the power extracted by the coupled port.

Circulators and Isolators

- **What they are:** Non-reciprocal ferrite devices. A **circulator** has three ports; power entering Port 1 can only exit Port 2, power entering Port 2 can only exit Port 3, and power entering Port 3 can only exit Port 1.
 - **What an Isolator is:** A circulator with Port 3 terminated by a $50 \, \Omega$ dummy load. Any signal traveling backward into Port 2 gets safely routed to Port 3 and dumped as heat.
 - **Quantum Application:** Used at cryogenic temperatures (~ 10 mK) to protect fragile qubits. They allow the weak readout signals to travel out from the qubit toward the amplifiers, while preventing thermal noise and amplifier reflections (backaction) from traveling down into the quantum processor.
-

Calibration: Purpose and Concepts

A VNA inherently measures everything relative to its internal receivers. However, the cables, adapters, and fixtures connecting the VNA to your device introduce massive phase shifts (delays), reflections, and attenuation. **Calibration mathematically shifts the “measurement plane” from inside the instrument directly to the terminals of the DUT.**

[VNA Internal] ===== (Long Coaxial Cable) ===== [Cal Plane] [DUT]

Standard Calibration Techniques

- **SOL (Short, Open, Load):** Used for 1-port measurements (S_{11}). It utilizes three standards: a perfect short circuit (total reflection, 180° phase shift), a perfect open circuit (total reflection, 0° phase shift), and a precision $50\ \Omega$ load (zero reflection).
- **SOLT (Short, Open, Load, Through):** The classic 2-port calibration method. It adds a seamless “Through” connection between Port 1 and Port 2 to correct for transmission errors.
- **E-Cal (Electronic Calibration):** An electronic module containing solid-state states that automatically switch between various impedance states. It replaces manual mechanical standards, removing human error and reducing connector wear.

The 1-Port Mathematical Error Model

A 1-port calibration maps three systematic complex error terms across frequency:

1. **Directivity (E_d):** Leakage from the source directly into the reflection receiver.
2. **Source Match (E_s):** Reflections from the VNA port bouncing back toward the DUT.
3. **Reflection Tracking (E_r):** The frequency response (loss and delay) of the internal paths and external cables.

The VNA uses these terms to map the raw measured reflection (S_{11m}) to the true actual reflection (Γ_{true}):

$$S_{11m} = E_d + \frac{E_r \Gamma_{\text{true}}}{1 - E_s \Gamma_{\text{true}}}$$

By measuring three known calibration standards, the VNA solves this algebraic system for E_d , E_s , and E_r at every frequency point, allowing it to isolate and calculate Γ_{true} during runtime.

Measurement Accuracy and Limits

- **Dynamic Range:** The difference between the maximum output power of the source and the VNA’s noise floor. Reducing the **IF Bandwidth (Intermediate Frequency BW)** digitally filters out broadband noise, vastly increasing dynamic range at the expense of slower sweep speeds.
 - **Phase Stability:** Phase values are highly vulnerable to physical temperature swings and cable flexing. Even minor bending of an RF cable changes its physical length slightly, causing massive phase drift at gigahertz frequencies.
 - **Averaging:** Taking multiple trace sweeps and averaging them suppresses uncorrelated random white noise, which is vital when attempting to resolve sub-decibel changes in low-power experiments.
-

Specific Notes for Qubit/Resonator Characterization

- **The Single-Photon / Linear Regime:** Superconducting resonators are highly non-linear if overdriven. To measure their true quantum properties, you must attenuate the VNA source signal significantly (often using -60 dB to -80 dB of attenuation inside the dilution refrigerator) so that the average photon population inside the resonator is $\langle n \rangle \approx 1$.
 - **Thermalization:** Cryogenic attenuators and isolators do more than route signals—they physically anchor the coaxial center conductors to the various temperature stages (4 K, 100 mK, 10 mK), stripping away room-temperature blackbody photons before they hit the qubit.
 - **De-embedding:** Since you cannot place physical calibration standards inside a 10 mK dilution refrigerator, you calibrate at room temperature. You then use **de-embedding** software matrices to mathematically subtract the known electrical delays and propagation losses of the long cryogenic lines leading down to your sample holder.
-

Reference Metrics (Typical Baseline values)

Metric	Typical Value Range	Context
Coupling Strength ($g/2\pi$)	50 to 300 MHz	Interaction speed between a qubit and a readout resonator.
Resonator Linewidth ($\kappa/2\pi$) [3D]	1 to 10 kHz	Ultra-high quality factor macro-cavities (low loss).

Metric	Typical Value Range	Context
Resonator Linewidth ($\kappa/2\pi$) [2D]	5 to 10 MHz	On-chip coplanar waveguide resonators optimized for fast readout.
Zero-Point Fluctuations (V_{zpf})	$\sim 1 \mu\text{V}$	The fundamental quantum voltage fluctuations present in a 50Ω coplanar resonator even when completely empty of photons.

Quick Glossary

- **Return Loss:** A metric of port matching. A higher return loss (e.g., 30 dB) means less signal is bouncing back; a lower return loss (e.g., 3 dB) indicates a severe impedance mismatch.
- **Insertion Loss:** The drop in signal power across a system. An insertion loss of 3 dB means exactly half of the input power was lost/absorbed within the device.
- **Directivity:** The metric of a directional coupler's quality. It dictates how cleanly it isolates forward waves from reverse waves.
- **Isolation:** The measure of leakage in non-reciprocal devices like circulators or switches. High isolation prevents leakage into unwanted channels.
- **De-embedding:** A post-processing mathematical technique that uses matrix algebra to eliminate the parasitic effects of fixtures, bond wires, or cables that cannot be physically disconnected during calibration.

Here is the completely unified, comprehensive reference manual. This version explicitly preserves every introductory explanation, step-by-step point, and practical example from your original text, seamlessly weaving them into the deep mathematical, physical, and structural breakdowns.

Signal Representation

1. Why We Need Multiple Representations

Signals are physical quantities (such as time-varying voltages or currents) generated, transmitted, processed, and measured in different domains. A single representation rarely suits all engineering tasks.

Physical Domain		Mathematical Representation		Engineering Application
Time Domain	→	$s(t) = A(t) \cos(\omega t + \phi)$	→	Power/Amplifier Design
Frequency Domain	→	$S(\omega) = \mathcal{F}\{s(t)\}$	→	Spectrum/Filter Design
Complex Baseband	→	$u(t) = I(t) + jQ(t)$	→	DSP/SDR Algorithms

- **Time-Domain:** Essential for physical hardware design. Real voltages in engineering and communications are time-varying waveforms that can be positive or negative. Power amplifiers, Analog-to-Digital Converters (ADCs), and transmission media respond directly to these instantaneous values.
- **Frequency-Domain:** Necessary for regulatory compliance (spectral masks), channel allocation, and filter design. It isolates the distribution of energy across the spectrum, clarifying parameters like absolute bandwidth.
- **Complex/Baseband Domain:** Different views (time, frequency, complex/baseband, phasor) clarify bandwidth, phase, and modulation while simplifying design and DSP implementation. By decoupling the information-bearing envelope from the high-frequency carrier, it drastically reduces computational complexity.
- **Core Modern Applications:** This dual-representation paradigm is foundational to:
 - **Software Defined Radios (SDR):** Where modulation, filtering, and synchronization occur entirely in software via complex numeric arrays.
 - **Quantum Computing (Microwave Qubit Control):** Where microwave pulses must be shaped with nanosecond precision in both amplitude and phase to manipulate the state transitions of superconducting qubits.

2. Real vs. Complex (I/Q) Representation

A real narrowband carrier or physical passband signal localized around a carrier frequency $\omega_c = 2\pi f_c$ can be expressed in its time-domain real form as:

$$x(t) = A(t) \cos(\omega_c t + \phi(t))$$

Where $A(t)$ represents the instantaneous amplitude modulation and $\phi(t)$ represents the instantaneous phase modulation. In physical receivers, these signals are transformed into a complex form to make analysis and hardware design easier to handle. Using the trigonometric identity $\cos(\alpha + \beta) = \cos \alpha \cos \beta - \sin \alpha \sin \beta$, we can expand this expression:

$$x(t) = A(t) \cos(\phi(t)) \cos(\omega_c t) - A(t) \sin(\phi(t)) \sin(\omega_c t)$$

By defining the **In-phase** ($I(t)$) and **Quadrature** ($Q(t)$) components as:

- $I(t)$: The in-phase component, scaling the cosine carrier: $I(t) = A(t) \cos(\phi(t))$
- $Q(t)$: The quadrature component, scaling the sine carrier: $Q(t) = A(t) \sin(\phi(t))$

The real passband signal simplifies to a linear combination of two orthogonal carrier waves:

$$x(t) = I(t) \cos(\omega_c t) - Q(t) \sin(\omega_c t)$$

This complex notation makes quadrature modulators and demodulators straightforward to describe. Using Euler's formula, we wrap these components into an equivalent complex (analytic) representation that utilizes a **Complex Baseband Signal / Complex Envelope** $u(t) = I(t) + jQ(t)$:

$$x(t) = \text{Re} \{u(t)e^{j\omega_c t}\} = \text{Re} \{(I(t) + jQ(t))e^{j\omega_c t}\}$$

Structural Properties of the I/Q Vector Space

- **Orthogonality:** The carriers $\cos(\omega_c t)$ and $\sin(\omega_c t)$ are orthogonal (90° apart in phase) over a carrier period $T_c = \frac{2\pi}{\omega_c}$:

$$\int_0^{T_c} \cos(\omega_c t) \sin(\omega_c t) dt = 0$$

Because they are completely orthogonal, they permit independent amplitude and phase control without interfering with each other. * **Geometric Mapping:** This structural independence allows $I(t)$ and $Q(t)$ to map directly onto a two-dimensional Cartesian plane (Constellation Diagram), functioning as the physical implementation of a complex signal.

3. Analytic Signal and Hilbert Transform

For a real-valued physical signal $s(t)$, the **Analytic Signal** $s_a(t)$ is a complex-valued extension used for advanced processing. It eliminates all negative-frequency content while preserving the positive-frequency spectrum. It is defined mathematically as:

$$s_a(t) = s(t) + j\mathcal{H}\{s(t)\}$$

Where $\mathcal{H}\{\cdot\}$ is the **Hilbert Transform**, defined as the convolution of $s(t)$ with the impulse response $h(t) = \frac{1}{\pi t}$:

$$\mathcal{H}\{s(t)\} = s(t) * \frac{1}{\pi t} = \frac{1}{\pi} \int_{-\infty}^{\infty} \frac{s(\tau)}{t - \tau} d\tau$$

Frequency-Domain Operation

In the frequency domain, the Hilbert transform acts as an ideal -90° phase shifter for positive frequencies and a $+90^\circ$ phase shifter for negative frequencies:

$$H(\omega) = \mathcal{F}\left\{\frac{1}{\pi t}\right\} = -j\text{sgn}(\omega) = \begin{cases} -j = e^{-j\pi/2}, & \omega > 0 \\ 0, & \omega = 0 \\ +j = e^{j\pi/2}, & \omega < 0 \end{cases}$$

Consequently, the Fourier transform of the analytic signal $S_a(\omega)$ becomes:

$$S_a(\omega) = S(\omega) + j[-j\text{sgn}(\omega)]S(\omega) = S(\omega)[1 + \text{sgn}(\omega)] = \begin{cases} 2S(\omega), & \omega > 0 \\ S(0), & \omega = 0 \\ 0, & \omega < 0 \end{cases}$$

As a result, $s_a(t)$ contains only positive-frequency content. Any narrowband real signal can be written as the real part of this analytic signal, linking back to the complex envelope:

$$s(t) = \text{Re}\{s_a(t)\} = \text{Re}\{u(t)e^{j\omega_c t}\}$$

4. Euler's Formula and Complex Spectra

Mathematical Foundations

The foundation of this complex representation relies directly on Euler's formula:

- $e^{j\omega t} = \cos(\omega t) + j \sin(\omega t)$
- $e^{-j\omega t} = \cos(\omega t) - j \sin(\omega t)$

From these expressions, we derive the standard exponential identities for basic sinusoids:

- $2 \cos(\omega t) = e^{j\omega t} + e^{-j\omega t}$
- $2j \sin(\omega t) = e^{j\omega t} - e^{-j\omega t}$

Real vs. Complex Signals in the Frequency Domain

- **One-Sided Spectrum:** A real signal is usually illustrated practically using a one-sided spectrum, where only positive frequencies are displayed because the negative side contains duplicate information.
- **Two-Sided Spectrum:** A complex signal possesses a two-sided spectrum with independent components on both the positive and negative frequency axes.
- **Conceptual Interpretation:** The terms “positive” and “negative frequency” are best understood physically as **positive and negative complex exponentials** rotating in opposite directions in the complex plane.

Benefits of Complex Notation (Algebraic Simplicity)

- **Real Trigonometric Products:** Multiplying real signals yields unwieldy expansions:

$$\cos A \cos B = \frac{1}{2} [\cos(A + B) + \cos(A - B)]$$

This splits the energy and generates undesired high-frequency image terms that require extra hardware filtering to remove. * **Complex Exponentials:** Complex multiplication simplifies directly to exponent addition:

$$e^{jA} e^{jB} = e^{j(A+B)}$$

Multiplication becomes a clean frequency translation without generating stray image sidebands. * **Frequency Shift Theorem:** If $S(\omega)$ is the spectrum of a baseband signal $u(t)$, then multiplying it by a complex carrier shifts its entire spectrum cleanly along the frequency axis:

$$\mathcal{F}\{u(t)e^{j\omega_c t}\} = U(\omega - \omega_c)$$

5. Frequency Domain Examples & Symmetry

Example 1: Sum of Three Ideal Tones

Consider a multi-tone real signal defined by:

$$s_1(t) = 10 \cos(2\pi 100t) + \cos(2\pi 200t) + 4 \cos(2\pi 300t)$$

- **Time-Domain View:** In the time domain, plotting $s_1(t)$ yields a complex, overlapping waveform where individual frequency content is difficult to isolate visually.
- **Frequency-Domain View:** In the frequency domain, the signal is easy to interpret as three discrete sinusoidal components. Using the exponential

expansion $\cos(\omega t) = \frac{e^{j\omega t} + e^{-j\omega t}}{2}$, the complex two-sided spectrum displays symmetric component pairs at ± 100 Hz, ± 200 Hz, and ± 300 Hz. These symmetric components combine to form the real cosine signal, and the ideal imaginary spectrum is completely zero.

Example 2: Sum of Three Tones with Phase Shifts

Let us introduce arbitrary phase shifts to the standard cosine tones from the first example:

$$s_2(t) = 10 \cos(2\pi 100t + \pi/4) + \cos(2\pi 200t + \pi/6) + 4 \cos(2\pi 300t)$$

- **Time-Domain Changes:** In the time domain, the phase shifts alter the alignment of the peaks, changing the shape of the overall waveform. However, the *magnitude spectrum* remains exactly the same as it was for $s_1(t)$.
- **Phase Spectrum Changes:** The phase spectrum changes to reflect the added offsets. Expanding this with Euler's formula reveals complex coefficients:

$$10 \cos\left(2\pi 100t + \frac{\pi}{4}\right) = 5e^{j\pi/4}e^{j2\pi 100t} + 5e^{-j\pi/4}e^{-j2\pi 100t}$$

Evaluating the complex spectral coefficients c_n for each frequency step reveals non-zero real and imaginary parts:

$$\begin{aligned} \bullet \quad c_{100} &= 5e^{j\pi/4} = \frac{5\sqrt{2}}{2} + j\frac{5\sqrt{2}}{2}, & c_{-100} &= 5e^{-j\pi/4} = \frac{5\sqrt{2}}{2} - j\frac{5\sqrt{2}}{2} \\ \bullet \quad c_{200} &= 0.5e^{j\pi/6} = \frac{\sqrt{3}}{4} + j\frac{1}{4}, & c_{-200} &= 0.5e^{-j\pi/6} = \frac{\sqrt{3}}{4} - j\frac{1}{4} \\ \bullet \quad c_{300} &= 2e^{j0} = 2, & c_{-300} &= 2e^{-j0} = 2 \end{aligned}$$

Key Symmetry Observations

- **Conjugate Symmetry:** For all real-valued physical signals, the two-sided complex spectrum is strictly **conjugate symmetric** ($S(-\omega) = S^*(\omega)$). This means:
- The real part of the spectrum is **even-symmetric** ($\text{Re}\{S(-\omega)\} = \text{Re}\{S(\omega)\}$)
- The imaginary part of the spectrum is **odd-symmetric** ($\text{Im}\{S(-\omega)\} = -\text{Im}\{S(\omega)\}$)
- **Phase Representation:** Because phase information is naturally embedded within the relationship between these real and imaginary components, a separate phase spectrum is not always mathematically required.
- **Utility:** Complex two-sided spectra are especially useful in quadrature communications because they make phase shifts and directional frequency rotations easy to track and interpret.

6. Modulation Mechanics

Basic Concept of Modulation

A standard sinusoidal carrier wave can be mathematically described by its basic parameters:

$$v(t) = A \sin(2\pi ft + \phi)$$

Modulation alters one or more of these core parameters—**amplitude**, **frequency**, or **phase**—of a high-frequency carrier wave according to a lower-frequency message signal $m(t)$. This shifts the information to a higher frequency band so it can propagate effectively through physical transmission media such as air, optical fiber, or coaxial cables.

Amplitude Modulation (AM)

In standard Amplitude Modulation, the carrier's amplitude varies linearly with the message signal. Let a message signal be $m(t) = A_m \cos(2\pi f_m t)$ and the unmodulated carrier be $c(t) = A_c \cos(2\pi f_c t)$. The combined AM signal is expressed as:

$$s(t) = [A_c + A_m \cos(2\pi f_m t)] \cos(2\pi f_c t)$$

Using the trigonometric identity $\cos A \cos B = \frac{1}{2}[\cos(A + B) + \cos(A - B)]$, we rewrite the AM signal in its explicit spectral form:

$$s(t) = A_c \cos(2\pi f_c t) + \frac{A_m}{2} \cos(2\pi(f_c + f_m)t) + \frac{A_m}{2} \cos(2\pi(f_c - f_m)t)$$

- **AM Frequency Profile:** This produces a discrete carrier at f_c , a Lower Sideband (LSB) at $f_c - f_m$, and an Upper Sideband (Reflecting $f_c + f_m$).
- **Bandwidth Inefficiency:** Standard AM is inefficient because it requires roughly twice the baseband bandwidth ($2f_m$) to transmit a single-ended message, duplicating information symmetrically on both sides of the carrier.

Quadrature Modulation

Quadrature modulation resolves this inefficiency by using two orthogonal carriers simultaneously: $\cos(\omega t)$ for the In-phase (I) branch and $\sin(\omega t)$ for the Quadrature (Q) branch.

Because these two carrier waves are 90° apart in phase, they do not interfere with each other during transmission. This allows both independent signals to occupy the exact same frequency band simultaneously, doubling the spectral efficiency relative to standard AM.

7. I/Q Modulation and Demodulation Linearity

In real-world communications, I/Q modulation is more than a mathematical convenience: it is how physical transmitters and receivers operate.

Upconversion Time-Domain Form (Transmitter)

The transmitted hardware signal maps directly to the real and imaginary parts of the complex baseband signal $u(t) = g_1(t) + jg_2(t)$. It can be written as:

$$y(t) = g_1(t) \cos(2\pi f_c t) - g_2(t) \sin(2\pi f_c t)$$

Where:

- $g_1(t)$ represents the continuous **In-Phase (I) channel** stream.
- $g_2(t)$ represents the continuous **Quadrature (Q) channel** stream.
- f_c is the targeted carrier frequency.


Expressed compactly in complex notation, this matches our foundational framework:

$$y(t) = \text{Re}\{u(t)e^{j2\pi f_c t}\}$$

By changing $g_1(t)$ ($I(t)$) and $g_2(t)$ ($Q(t)$) over time, the resulting output carries amplitude changes, phase changes, or complex digital modulations (such as QAM or QPSK). For example, if I and Q are driven with equal static values, the resulting passband signal exhibits a predictable phase shift of exactly 45° .

I/Q Demodulation (Receiver Architecture)

At the receiver, the incoming signal $r(t)$ is split into two parallel analog processing branches. To decode the signal, it is multiplied by two synchronized local oscillators (0° cosine for the I branch, 90° sine for the Q branch) and passed through low-pass filters:



Mathematical Proof of Separation (Cross-Talk Elimination) Assuming a noise-free transmission where $r(t) = g_1(t) \cos(\omega_c t) - g_2(t) \sin(\omega_c t)$:

1. In-Phase Branch Multiplier:

$$r(t) \cos(\omega_c t) = g_1(t) \cos^2(\omega_c t) - g_2(t) \sin(\omega_c t) \cos(\omega_c t)$$

Applying standard trigonometric identities ($\cos^2 \theta = \frac{1+\cos 2\theta}{2}$ and $\sin \theta \cos \theta = \frac{\sin 2\theta}{2}$):

$$r(t) \cos(\omega_c t) = \frac{1}{2}g_1(t) + \frac{1}{2}g_1(t) \cos(2\omega_c t) - \frac{1}{2}g_2(t) \sin(2\omega_c t)$$

Passing this through an ideal Low-Pass Filter (LPF) completely attenuates the high-frequency $2\omega_c$ components, isolating the I data:

$$I(t) = \text{LPF}\{r(t) \cos(\omega_c t)\} = \frac{1}{2}g_1(t)$$

2. Quadrature Branch Multiplier: Multiply the incoming signal by $\sin(\omega_c t)$:

$$r(t) \sin(\omega_c t) = g_1(t) \cos(\omega_c t) \sin(\omega_c t) - g_2(t) \sin^2(\omega_c t)$$

Applying standard trigonometric identities ($\sin^2 \theta = \frac{1-\cos 2\theta}{2}$):

$$r(t) \sin(\omega_c t) = \frac{1}{2}g_1(t) \sin(2\omega_c t) - \frac{1}{2}g_2(t) + \frac{1}{2}g_2(t) \cos(2\omega_c t)$$

Passing this through the ideal Low-Pass Filter removes the $2\omega_c$ terms, leaving the scaled Q data:

$$Q(t) = \text{LPF}\{r(t) \sin(\omega_c t)\} = -\frac{1}{2}g_2(t)$$

(Note: Hardware paths flip the mixer sign or invert the digital array during processing to recover a positive $\frac{1}{2}g_2(t)$ stream).

Summary of Modulation vs. Demodulation

- **Modulation:** Adds information to a carrier wave, combining two separate streams (I and Q) to alter the overall amplitude and phase over time.
- **Demodulation:** Recovers information from the carrier wave, separating the I and Q components so amplitude and phase can be precisely measured.

8. Positive/Negative Frequency and Image Rejection

Because real physical waveforms have symmetric spectra, a naive, single-ended downconversion process (multiplying a passband signal by a single real local oscillator $\cos(\omega_c t)$) folds the negative frequency mirror image directly on top of the target baseband signal. This results in irreversible distortion unless blocked by expensive analog front-end RF filters.

Analytic I/Q representation prevents this image folding. By maintaining both the I and Q signal branches throughout the receiver, the system preserves the phase relationship between the channels. This phase difference creates constructive interference for the target signal and destructive interference for the image frequency, cancelling out the unwanted image without needing steep analog front-end filtering.

9. Practical Notes for DSP and Hardware

Complex Baseband Sampling (Nyquist Criteria)

If a real passband signal has an absolute bandwidth of B centered at a high carrier frequency f_c , a direct real sampling architecture requires an ADC sampling rate greater than twice the highest frequency component:

$$f_s > 2 \left(f_c + \frac{B}{2} \right)$$

By downconverting to complex baseband before digital conversion, the signal centers around 0 Hz and spans from $-\frac{B}{2}$ to $+\frac{B}{2}$. The Nyquist sampling rate for this complex bandwidth is reduced to:

$$f_s \geq B$$

This allows systems to use two lower-speed ADCs (sampling at $f_s \geq B$ for the separate I and Q paths) rather than a single, expensive ultra-high-speed ADC.

Filtering and Calibration

- **Filtering:** Digital low-pass filters are designed and applied to the decoupled I and Q streams to remove the high-frequency image components and residual mixer products shown in section 7.
- **Phase/Amplitude Calibration:** In physical systems, analog components introduce **I/Q imbalance** (gain differences or minor phase deviations from a perfect 90° shift). This imbalance causes image leakage, which modern communication front-ends estimate and correct using digital adaptive filters in DSP.

In quantum computing architectures, precise complex pulses (defined by their baseband envelope $u(t)$) are synthesized to shape the amplitude and phase of microwave lines. This control allows for exact quantum state manipulation.


During readout, the reflected or transmitted microwave signals are demodulated back into their complex baseband components (I and Q). This allows the system to extract subtle amplitude and phase shifts, which are then used to determine the state of the qubit.


Core Formulation Guide

Signal Metric	Mathematical Formulation
Real Carrier Modulation	$x(t) = A(t) \cos(\omega_c t + \phi(t)) = I(t) \cos(\omega_c t) - Q(t) \sin(\omega_c t)$
Complex Baseband Envelope	$u(t) = I(t) + jQ(t)$
Analytic Signal Derivation	$s_a(t) = s(t) + j\mathcal{H}\{s(t)\} = u(t)e^{j\omega_c t}$
Frequency Shift Theorem	$\mathcal{F}\{u(t)e^{j\omega_c t}\} = U(\omega - \omega_c)$
Ideal IQ Demodulation Output	$u(t) \approx 2 \cdot \text{LPF}\{r(t)e^{-j\omega_c t}\} = I(t) + jQ(t)$
Conjugate Spectrum Symmetry	$S(-\omega) = S^*(\omega)$ (For real-valued time waveforms)

- **Euler's Formula Expansion:** $e^{\pm j\omega t} = \cos(\omega t) \pm j \sin(\omega t)$
- **Sinusoidal Composition:** $\cos(\omega t) = \frac{e^{j\omega t} + e^{-j\omega t}}{2}$, $\sin(\omega t) = \frac{e^{j\omega t} - e^{-j\omega t}}{2j}$
- **Standard AM Waveform:** $s(t) = A_c \cos(2\pi f_c t) + \frac{A_m}{2} \cos(2\pi(f_c + f_m)t) + \frac{A_m}{2} \cos(2\pi(f_c - f_m)t)$
- **Quadrature Passband Signal:** $y(t) = g_1(t) \cos(2\pi f_c t) - g_2(t) \sin(2\pi f_c t) = \text{Re}\{u(t)e^{j2\pi f_c t}\}$

A basic Radio Frequency (RF) communication system is divided into three fundamental blocks: the **Transmitter (TX)**, the **Channel**, and the **Receiver (RX)**.





Transmitter (TX) Operations

1. **Baseband Signal Processing:** Raw digital data bits are mapped into complex symbols (I/Q channels) using digital processing. Filtering (e.g., root-raised cosine) shapes the pulses to limit spectral bandwidth.
2. **Upconversion (Modulation):** The low-frequency baseband signal is shifted up to a high-frequency carrier (ω_c). This positions the signal within its allocated spectral band.
3. **Power Amplification:** The modulated RF signal passes through a Power Amplifier (PA). The PA increases the signal amplitude to provide enough power to overcome transmission losses over the channel.

The Channel (The Qubit Context)

In standard wireless communications, the channel is free space, which introduces path loss, fading, and multipath interference.

In quantum computing electronics (superconducting qubit control), **the physical link is entirely wired using specialized coaxial cables**. These cables run down a dilution refrigerator from room temperature (300 K) to cryogenic stages (4 K, 100 mK, and finally the mixing chamber at ~ 10 mK).

- **Attenuators:** Thermal noise from room-temperature electronics travels down the lines and can dephase or excite the qubits. To prevent this, discrete cryogenic attenuators are placed at different temperature stages to reduce noise power.
- **Circulators and Isolators:** On the readout (receive) path, passive microwave devices called circulators allow signals to travel from the qubit to the receiver while blocking noise reflecting back down from the warm electronics.

Receiver (RX) Operations

1. **Low-Noise Amplification:** The incoming RF signal is weak. The first stage is a Low-Noise Amplifier (LNA). In quantum systems, this is a cryogenic High-Electron-Mobility Transistor (HEMT) amplifier or a quantum-limited Parametric Amplifier (e.g., TWPA) operating at 4 K or 10 mK. It boosts the signal while adding minimal thermal noise.
2. **Downconversion:** The high-frequency RF signal is translated back down to lower frequencies (either Intermediate Frequency or Baseband).

3. **Demodulation and Digital Processing:** The lower-frequency waveform is sampled by an Analog-to-Digital Converter (ADC). Digital Signal Processing (DSP) algorithms extract the original baseband data, evaluating changes in amplitude and phase over time.

2. RF Front-End Architectures

Translating a signal between baseband and RF frequencies can be done using different architectural configurations.

HETERODYNE (DIRECT CONVERSION / ZERO-IF)

[Baseband (DC)] <=====> [RF Frequency]
(Single Modulation Stage)

SUPERHETERODYNE

[Baseband (DC)] <=====> [Intermediate Frequency (IF)] <=====> [RF Frequency]
(Stage 1: Digital or Analog) (Stage 2: Analog Mixer)

Heterodyning (Direct Conversion / Zero-IF)

Heterodyning mixes a baseband signal directly up to the final RF carrier frequency in a single modulation stage (and vice versa at the receiver).

- **Mechanism:** The Local Oscillator (LO) frequency is set exactly equal to the target RF carrier frequency ($f_{LO} = f_{RF}$). When downconverting, the RF signal is mixed with f_{LO} , shifting the center frequency to 0 Hz (DC).
- **Advantages:** It eliminates the need for intermediate stages, reducing the component count by removing intermediate filters and mixers.
- **Disadvantages:** It is susceptible to **DC offsets** from LO leakage self-mixing, $1/f$ flicker noise at baseband, and I/Q amplitude/phase imbalances that degrade image rejection.

Superheterodyning

A superheterodyne architecture uses two distinct modulation stages to transition between baseband and RF. It introduces an **Intermediate Frequency (IF)** between the two domains.

- **Mechanism (Receiver Example):** 1. The incoming RF signal at f_{RF} passes through a bandpass filter and an LNA.
- 2. An analog mixer combines f_{RF} with a first local oscillator f_{LO1} to shift the signal to a fixed lower frequency, the Intermediate Frequency ($f_{IF} = |f_{RF} - f_{LO1}|$).
- 3. A highly selective **IF filter** (such as a SAW filter) isolates the target channel and rejects out-of-band interferers and the unwanted image frequency.

4. A second demodulation stage translates the clean signal from f_{IF} down to 0 Hz baseband for sampling.
- **Why it is used:** It is difficult to build high-gain, sharp, tunable filters at high RF frequencies. Shifting the signal to a fixed, lower IF allows high-selectivity filtering and amplification using stable, specialized components.

3. The Digital/Analog Interface & Sampling Rate Impact

The choice of radio architecture depends heavily on the capabilities of the digital-to-analog and analog-to-digital interfaces. The data converter sampling rate (f_s) determines the boundary between digital and analog processing.

The Nyquist Boundary

Digital processing can only manipulate the signal directly if the Nyquist-Shannon criteria is satisfied. To sample an RF signal digitally without aliasing distortion, the sampling rate must be greater than twice the highest frequency component of the target passband signal:

$$f_s > 2f_{RFmax}$$

- **If Met:** Modulation, upconversion, downconversion, and demodulation can be performed mathematically within the digital domain (DSP).
- **If Not Met:** Hardware processing is required. The system must use analog circuits (mixers, local oscillators, filters) to shift the signal down to a frequency range that the converters can handle.

4. Analysis of the Three Radio Architectures

Radio architectures are categorized by where the data converters (DAC/ADC) are placed along the processing chain.

1. DIRECT RF SAMPLING (All-Digital Up to RF)
[Digital DSP] -> [High-Speed DAC/ADC] -----> [RF Channel]
2. IF SAMPLING (Superheterodyne with Digital IF)
[Digital DSP] -> [DAC/ADC] -> [Analog Mixer & IF Filter] -> [Analog Mixer] ----> [RF Channel]
3. BASEBAND SAMPLING (Analog I/Q Quadrator)
[Digital DSP] -> [Low-Speed DAC/ADC] -> [Analog Quadrature Modulator (I/Q)] --> [RF Channel]

Architecture A: Direct RF Sampling (Almost All-Digital)

In this architecture, the DAC and ADC connect directly to the RF front-end components, such as the power amplifier or low-noise amplifier.

- **Operation:** All synthesis, filtering, upconversion, and downconversion are performed digitally using DSP algorithms. The DAC synthesizes the final RF carrier waveform directly, and the ADC samples the incoming RF carrier wave directly.
- **Hardware Requirements:** Requires ultra-high-speed data converters. For example, controlling a superconducting qubit with a transition frequency between 4 GHz and 8 GHz requires data converters sampling at $f_s > 16$ GSPS (Giga-Samples Per Second).

Architecture B: IF Sampling Radio

This configuration is a superheterodyne setup where the digital-to-analog boundary sits at an Intermediate Frequency.

- **Operation:** The digital processor handles modulation and shifts the baseband signal up to a digital IF. A moderately high-speed DAC generates this lower-frequency IF waveform. Analog components (mixers, local oscillators, and filters) then take over to perform the final upconversion from IF to the target RF band.
- **Hardware Requirements:** The data converters only need to sample fast enough to resolve the IF band ($f_s > 2f_{\text{IFmax}}$), which is lower than the final RF frequency. This allows the use of higher-resolution, lower-power converters.

Architecture C: Baseband-Sampling Radio

In this classic architecture, the data converters operate entirely at baseband processing rates, and the digital-to-analog boundary sits right at the output of the baseband processor.

- **Operation:** The DACs and ADCs process low-frequency baseband signals (I and Q channels separately). All modulation, downconversion, phase-shifting, and summation are performed in the analog domain using an analog quadrature modulator/demodulator.
- **Hardware Requirements:** Minimizes the required sampling rate for the data converters, which only need to match the instantaneous bandwidth of the modulated signal ($f_s \geq B$), rather than the carrier frequency.

5. Architectural Comparison

Parameter	Direct RF Sampling	IF Sampling	Baseband Sampling
Data Converter Speed	Ultra-High ($f_s > 2f_{\text{RF}}$)	Medium ($f_s > 2f_{\text{IF}}$)	Low ($f_s \geq$ Bandwidth)
Analog Component Count	Minimal (No mixers/IF stages)	Moderate (Mixers, LOs, IF filters)	High (Analog I/Q Modulator, LOs)
Susceptibility to Analog Imperfections	None (Digital domain)	Low (Isolated to analog stages)	High (Prone to I/Q imbalance, DC offset)
Flexibility / Reconfigurability	Maximum (Software-defined)	Moderate	Low (Fixed by analog hardware)

6. Advantages of Shifting to the Digital Domain

Moving the analog-digital boundary closer to the RF port offers several performance advantages:

- **Greater Operation Accuracy:** Digital operations are deterministic. Mathematical modulation avoids the signal degradation, harmonic generation, and non-linearities common to analog mixing stages.
- **Immunity to Environmental Factors:** Analog components change behavior due to temperature drift, component aging, and manufacturing tolerances. Digital logic circuits provide consistent performance over time regardless of these variations.
- **Simplified Bill of Materials (BOM) & Footprint:** Direct digital synthesis replaces physical components like discrete analog mixers, local oscillators, SAW filters, and splitters. This reduces the physical size of the PCB and simplifies hardware sourcing.
- **Power Savings:** While high-speed data converters draw significant current, eliminating multiple analog gain stages, ovens for local oscillators, and intermediate amplifiers can lower the overall power consumption of the system.
- **Hardware Flexibility and Reconfigurability:** In an all-digital radio architecture, changing the carrier frequency, modulation format (e.g., QPSK to 64-QAM), or pulse shape does not require modifying the physical hardware. It is updated by rewriting software parameters or modifying HDL code blocks.

Radio Frequency System on Chip (RFSoc)

An RFSoc integrates these digital and analog blocks onto a single semiconductor die. It combines high-performance FPGA programmable logic, multi-giga-sample

ADCs and DACs, and ARM processing cores within a single chip.

In quantum computing setups, RFSoc devices (such as the AMD Xilinx Zynq RFSoc series) are widely used. They allow developers to generate precise microwave control pulses and demodulate qubit readout signals directly on a single chip, avoiding the routing delays and calibration issues of discrete architectures.

Introduction to Radio Frequency System on Chip (RFSoc)

1. Challenges in Qubit Control & Readout Systems

Quantum computing architectures based on superconducting qubits impose demanding operational constraints on their supporting control and readout electronics.

- **Extremely High Timing and Signal Precision:** Qubit state manipulation relies on phase-coherent microwave pulses with nanosecond duration and precise amplitude shaping. Small timing jitters or phase drifts degrade quantum gate fidelities.
- **Dynamic Software and Firmware Ecosystem:** The firmware and high-level software tools required to interface with quantum processors are constantly evolving.
- **Co-Development Demands:** Because quantum hardware is experimental, research teams must act as co-developers. They frequently customize firmware to test new pulse-shaping or error-correction algorithms, placing heavy demands on engineering resources.
- **The Scalability Bottleneck:** As quantum processors scale from a few qubits to hundreds or thousands, traditional instrument racks become physically unmanageable, cost-prohibitive, and limited by inter-module synchronization latencies.

Current Commercial Solutions

The recent decade has seen the release of dedicated quantum computing products, for example from:

- **Qblox & Quantum Machines:** Specialized modular setups combining fast AWGs (Arbitrary Waveform Generators) and digital controllers.
- **Zurich Instruments & Keysight:** High-performance, low-latency lock-in amplifiers and microwave transceiver matrices.
- **Intermodulation Products:** Specialized multi-frequency signal analysis platforms.

While highly capable, these proprietary solutions are often expensive, locked into closed ecosystems, and require ongoing vendor development. This can limit their immediate utility for custom experimental configurations.


Advantages of Open-Source Frameworks

Despite the appeal of ready-to-run commercial products, there is significant value in exploring more accessible options. Open firmware and software for qubit control offer flexibility and reduce dependence on proprietary solutions, providing more accessible tools to the research community (e.g. **QICK: Quantum Instrumentation and Control Kit**).

This encourages broader contribution from the research community, potentially accelerating innovation and reducing costs.

2. What is an RFSoc?

An **RFSoc (Radio Frequency System on Chip)** integrates RF/analog, digital signal processing (DSP), and programmable logic (FPGA) on a single chip.



Historically, RF system designers used discrete multi-chip solutions:

1. **Three-Chip Solution:** Separate ICs for the Processing System (CPU), the Programmable Logic (FPGA), and the standalone data converters (DACs/ADCs).
2. **Two-Chip Solution:** A unified system-on-chip (CPU + FPGA) connected via a high-speed serial interface standard (like **JESD204B/C**) to an external RF data converter chip. This serial link introduces significant power consumption, routing complexity, and data transmission latency.

It offers a single-chip solution from digital to analog domain and vice versa (all digital solution).

Strategic Single-Chip Advantages

- **Integrated Front-End RF Interface:** Hardened **RF-DACs** and **RF-ADCs** on-chip remove the need for external front-end radio processing.
 - **Compact Physical Footprint:** Integrating these subsystems simplifies PCB layouts and significantly reduces the Bill of Materials (BOM).
 - **Minimized Latency:** Eliminating external serial transceiver links minimizes the propagation delay between the receive input and transmit output. This low latency is essential for fast, real-time feedback loops like **quantum error correction (QEC)** and active qubit reset.
 - **Hardware Acceleration:** The high-bandwidth, intra-chip connections between the processors and the FPGA fabric enable hardware acceleration of computationally intensive tasks.
-

3. The Zynq SoC Evolution

The RFSoc architecture builds directly on the evolution of AMD Xilinx's **Zynq System on Chip (SoC)** platform.

Historical Milestones

- **Zynq-7000 Family (2011):** First platform to combine a dual-core ARM Cortex-A9 processor with a 28 nm 7-Series FPGA fabric on a single die. This established the unified hardware/software paradigm.
- **Zynq UltraScale+ MPSoC:** Upgraded to a 16 nm FinFET architecture, incorporating a quad-core ARM application processor, a dual-core real-time processor, and a graphics processing unit (GPU).
- **Zynq UltraScale+ RFSoc (2018):** Added monolithic, multi-giga-sample data converters directly to the UltraScale+ MPSoC architecture.

The PS / PL Paradigm

A Zynq-based device is split into two primary functional domains:

1. **Processing System (PS):** The hardened, fixed-architecture processor core array. It handles software-level execution, runs operating systems, manages network interfaces, and coordinates high-level system parameters.
 2. **Programmable Logic (PL):** The reconfigurable FPGA fabric, composed of an array of Configurable Logic Blocks (CLBs). The PS views the PL as a flexible, low-latency programmable peripheral.
-

4. RFSoc Processing System (PS) Architecture

The Processing System of a modern UltraScale+ RFSoc contains specialized processing units designed for specific system tasks:

Application Processing Unit (APU)

The APU features a **Quad-core ARM Cortex-A53** processor optimized for high-level system management and user interaction.

- Contains a dedicated Floating-Point Unit (FPU) and NEON SIMD engine per core.
- Equipped with dedicated L1 caches (32 KB Instruction, 32 KB Data) and a shared 1 MB L2 cache.
- **Quantum Control Application:** The APU runs a fully featured Linux distribution (such as Ubuntu). It hosts network communication servers (e.g., Python-based TCP/IP sockets) that receive experimental pulse configurations from a user's PC, translating them into memory maps for the hardware layer.

Real-Time Processing Unit (RPU)

The RPU consists of a **Dual-core ARM Cortex-R5** array designed for deterministic, low-latency control loops.

- Operates in either split mode (two independent cores) or lock-step mode (redundant execution for safety-critical tasks).
- Features tightly coupled memory (TCM) to ensure deterministic execution times without cache-miss latencies.
- **Quantum Control Application:** The RPU handles highly time-sensitive control tasks, executing bare-metal code or running a lightweight Real-Time Operating System (RTOS) to coordinate precise trigger sequences.

Platform Management Unit (PMU)

A dedicated, isolated processor core that manages power gating, initialization sequencing, thermal monitoring, and device safety parameters independently of the APU and RPU.

5. The Programmable Logic (PL) Fabric

The PL fabric provides custom, parallel digital hardware processing, functioning as an on-chip FPGA.

Key Functional Hardened Blocks

- **Configurable Logic Blocks (CLBs):** The fundamental reconfigurable logic elements used to implement custom combinational and sequential digital circuits.
- **DSP48E2 Slices:** Specialized, high-speed hardened arithmetic blocks containing a 27×18 -bit multiplier, a 48-bit accumulator, and a pre-adder. These slices execute high-speed digital filtering, mixing, and real-time matrix operations at clock speeds up to several hundred MHz.
- **Block RAM (BRAM):** Dedicated, high-speed on-chip memory blocks. Each BRAM stores up to 36 Kb of data and can be configured as true dual-port RAM, ROM, or FIFO buffers. BRAM can be dynamically reshaped—for example, as a 4096×8 -bit array or an 8192×4 -bit array.
- **UltraRAM (URAM):** High-capacity, high-density on-chip memory cells designed to replace external storage. Each URAM block stores up to 288 Kb with a fixed configuration (4096×72 -bit, dual-port). Multiple URAM blocks can be cascaded together to build deep on-chip pulse playback buffers for storing complex quantum sequences.


Device Performance and Speed Grades


The performance of the logic elements in the PL fabric is designated by a device **Speed Grade** (e.g., 2E, 2I, 2LE, 2LI, 1E, 1I, 1M, 1LI).

- **Numerical Factors (2 vs. 1):** Higher numbers indicate faster transistor switching times and higher maximum internal clock frequencies.
- **Power Mode Indicator ('L'):** Denotes Low-Power variants that operate at reduced core voltages to minimize static power consumption.
- **Temperature Ratings:** Indicates the certified safe operating temperature bounds:
 - **E:** Extended (0°C to $+100^{\circ}\text{C}$)
 - **I:** Industrial (-40°C to $+100^{\circ}\text{C}$)
 - **M:** Military (-55°C to $+125^{\circ}\text{C}$)

6. PS-PL Interconnect Architecture

High-speed communication between the software domain (PS) and the hardware domain (PL) is built on the **ARM AMBA AXI4 (Advanced eXtensible Interface 4)** protocol framework.





Memory-Mapped I/O (MMIO)

The RFSoc uses a unified **Memory-Mapped I/O** architecture. The CPU does not use specialized input/output instructions; instead, it accesses hardware registers inside the PL fabric using standard memory instructions like **LOAD** and **STORE**.

Each custom hardware peripheral in the PL is assigned a unique base address within the main CPU system address map. Writing to that address updates a register in the FPGA fabric, allowing software to change hardware parameters directly.

The Three Functional AXI4 Variants

1. **AXI4 (Full):** A high-performance, memory-mapped interface that supports data burst transfers up to 256 beats. It is used for moving large data blocks between the PS system memory (DDR) and the PL.
2. **AXI4-Lite:** A simplified, single-transaction memory-mapped interface with low resource overhead. It is used to read and write to control status registers inside PL IP blocks, allowing the CPU to tune parameters like mixer frequencies or filter weights.
3. **AXI4-Stream:** A non-memory-mapped, unidirectional master-slave protocol designed for high-speed streaming data. It eliminates address overhead, allowing continuous data flow between the processing fabric and the data converters.

7. The RF Data Converter (RFDC) Subsystem

The defining feature of an RFSoc is the monolithic RF Data Converter (RFDC) block, which integrates multi-giga-sample ADCs and DACs directly into the digital system.

RFDC Architecture and Internal Subsystems

The data converters are grouped into structural units called **Tiles**. Each tile houses a local, low-noise **Phase-Locked Loop (PLL)** that generates high-speed sampling clocks from an external, low-jitter reference clock source.

- **Time-Interleaved ADC Architecture:** To achieve high sampling rates without reducing bit resolution, the ADCs use a time-interleaving architecture. A single high-speed channel combines multiple parallel sub-ADCs.

For example, a Gen 3 Dual-ADC tile uses 8 interleaved sub-ADCs per channel, while a Quad-ADC tile uses 4 sub-ADCs per channel.


- **Digital Upconverters (DUC) & Downconverters (DDC):** Hardened DSP blocks built directly into the converter tiles handle frequency translation.
- **DUC (Transmit):** Features a programmable interpolator, an inverse sinc compensation filter, and a complex digital mixer with a 48-bit **Numerically Controlled Oscillator (NCO)** to cleanly upconvert baseband signals to RF.
- **DDC (Receive):** Uses a complex digital mixer, an NCO, and a programmable decimation filter cascade to downconvert RF signals back to baseband.

8. Data Flow and PL-to-RFDC Interfacing

Data moves between the reconfigurable FPGA fabric (PL) and the high-speed data converters via dedicated AXI4-Stream links.

Transmit Path (RF-DAC Channel)


In the transmit path, the PL fabric acts as the AXI4-Stream **Manager** (source) and the RF-DAC channel operates as the **Subordinate** (sink).



When configured in a complex-to-real modulation scheme, the PL streams independent, digital complex baseband components (I and Q) directly into the DUC block of the RF-DAC tile. The internal mixer modulates these components onto a carrier wave, and the high-speed 14-bit DAC core synthesizes the final analog passband signal.

Receive Path (RF-ADC Channel)

In the receive path, the roles flip: the RF-ADC block acts as the AXI4-Stream **Manager** and streams digital data into the PL fabric, which acts as the **Subordinate**.



The incoming single-ended analog readout signal from the quantum hardware passes through an external **Balun** (a passive transformer that converts a single-

ended signal to a balanced differential pair) to match the differential inputs of the RFSoc.

The high-speed 14-bit ADC samples the differential waveform, the DDC filters and downconverts the signal into digital I/Q components, and the resulting data streams directly into the PL fabric for real-time state analysis.

9. Understanding Super Sample Rate (SSR)

A core engineering challenge in an RFSoc is bridging the speed gap between the high-speed data converters and the slower FPGA fabric.

The Speed Gap Problem

An RF-ADC may sample an incoming signal at $f_s = 5$ GSps. However, the programmable logic fabric (PL) typically operates at a maximum internal clock rate of a few hundred MHz (e.g., $f_{PL} = 625$ MHz). A standard sequential processing chain cannot handle data arriving at this rate because the required clock frequency is too high for the FPGA fabric.

The SSR Solution

The **Super Sample Rate (SSR)** method resolves this mismatch by deserializing the fast serial data stream into a wide parallel array of samples. Instead of transferring a single data sample on every clock cycle, the system packs multiple time-contiguous samples into a single wide AXI4-Stream bus word on each clock cycle of the slower PL clock.

HIGH-SPEED REAL-TIME STREAM (5 GSps Serial)

```
---[ Sample 0 ]-[ Sample 1 ]-[ Sample 2 ]-[ Sample 3 ]-[ Sample 4 ]-[ Sample 5 ]-[ Sample 6 ]
                                     |
                                     v  [Deserialized via SSR = 8]
                                     |
```

PARALLEL PL STREAM (625 MHz Parallel Clock Cycle)

```
+-----+
| AXI4-Stream Word: [Samp 7][Samp 6][Samp 5][Samp 4][Samp 3][Samp 2][Samp 1][Samp 0] (8 * 16) |
+-----+
```

The relationship is governed by the structural formula:

$$\text{Effective Data Sampling Rate} = f_{PL_Clock} \times \text{SSR}$$

For example, given an RF-ADC sampling at 5 GSps and an SSR factor of 8, the required internal PL interface clock frequency is reduced to a manageable rate:

$$f_{\text{PL_Clock}} = \frac{5 \text{ GSps}}{8} = 625 \text{ MHz}$$

Each 16-bit data sample is padded, meaning a single clock edge transfers $8 \times 16\text{-bit} = 128$ bits of parallel data. This approach allows the system to process high-bandwidth data in real time, but it requires highly parallel DSP architectures inside the PL to process multiple concurrent samples on every clock cycle.

10. Multi-GSps Sampling and Nyquist Zone Operations

The multi-giga-sample capabilities of the RFDC change how signals are synthesized and sampled across different frequency zones.

Direct Frequency Sampling (1st Nyquist Zone)

When an RF-ADC samples at $f_s = 5 \text{ GSps}$, its primary or **1st Nyquist Zone** spans from 0 Hz to the Nyquist frequency $\frac{f_s}{2} = 2.5 \text{ GHz}$.

Any analog signal within this 0 – 2.5 GHz window can be directly sampled and digitized without aliasing. To prevent high-frequency noise from folding back into this band, an analog low-pass filter should be placed before the ADC input.

Under-Sampling Operations (2nd Nyquist Zone)

The RFSoc can also capture higher frequencies by operating within its **2nd Nyquist Zone**, which spans from $\frac{f_s}{2}$ to f_s (2.5 GHz to 5.0 GHz for a 5 GSps sample rate).

This approach uses **under-sampling** to intentionally let signals in the higher band alias back into the primary baseband zone (0 – 2.5 GHz). To use this method effectively, an analog bandpass filter must be placed before the input to isolate the target frequency band and block interfering signals from other zones.

Comparison: RFSoc vs. Heterodyne Architectures

- **Traditional Systems:** Limited to lower sampling rates (10s to 100s of MHz). They require complex analog RF front-ends (including local oscillators, analog mixers, and image-rejection filters) to downconvert signals from microwave frequencies to a lower intermediate frequency (IF) before sampling.
 - **RFSoc Systems:** High sampling rates allow the data converters to digitize signals directly at RF or intermediate frequencies. This shifts frequency translation into software-defined digital mixers, improving system stability and flexibility.
-

11. Overview of RFSoc Development Boards

To simplify implementation, AMD Xilinx and academic partners offer development boards that package the RFSoc with essential supporting hardware.

ZCU208 Evaluation Board

A mainstream Gen 3 evaluation platform built around the **XCZU48DR** chip.

- **Resources:** Includes 8 RF-ADC channels (5 GSps) and 8 RF-DAC channels (9.85 GSps), both with 14-bit resolution.
- **Connectivity:** Uses high-density RFMC (RF Micro-Coaxial) connectors to route high-speed signals. External breakout cards are used to convert these signals to standard SMA connectors.

Supporting Hardware Add-On Cards

- **Balun Daughter Cards:** Convert differential signals from the chip's pins into single-ended SMA connections, matching standard laboratory equipment.
- **Loopback Cards (XM650):** Connect internal RF-DAC channels directly to corresponding RF-ADC channels on the board, simplifying diagnostic testing and firmware calibration loops.
- **Low-Noise Clock Cards:** Provide stable, low-jitter external reference clock sources to drive the on-chip PLLs, helping maintain phase coherence across channels.


Academic Boards: The RFSoc 4x2


The **RFSoc 4x2** is a compact development platform designed for academic labs and university coursework, replacing the earlier Gen 1 RFSoc 2x2 board.

- To reduce cost and board complexity, it only breaks out a subset of the chip's channels: **4 RF-ADC paths** and **2 RF-DAC paths**.
- Features integrated, on-board baluns that route directly to standard SMA connectors, removing the need for external breakout daughter cards.

12. Real-World Application: Two-Qubit Superconducting Architecture

The diagram below illustrates how an RFSoc development board integrates into a physical, two-qubit superconducting quantum control and readout system.





Signal Architecture Functions

1. **Qubit Coherent Drive Path:** Dedicated RF-DAC channels generate envelope-shaped microwave pulses centered near the qubit transition frequency ω_{01} ($\sim 4-6$ GHz). These pulses travel down the dilution refrigerator through coaxial lines with inline attenuators to suppress thermal noise.
2. **Qubit Multiplexed Readout Path:** A separate RF-DAC channel generates a continuous readout probe tone near the resonator frequency

ω_r ($\sim 6 - 7$ GHz). This tone passes through the readout line, where its phase is shifted depending on the qubit states.

3. **Isolation and Return Amplification:** The reflected readout signal passes through a series of passive circulators. These circulators allow the signal to travel upward while blocking thermal noise from traveling back down toward the sensitive qubits. The weak microwave signal is amplified by a cryogenic High-Electron-Mobility Transistor (HEMT) low-noise amplifier before returning to room temperature.
4. **Digitization and Phase Extraction:** The amplified analog signal enters the RF-ADC input channel. The on-chip DDC downconverts and extracts the digital I/Q baseband components, allowing real-time state discrimination algorithms running in the PL fabric to determine the final state of the qubits.

13. System Selection and Feature Matrix

Functional Capability	RFSoc Architectural Implementation
Baseband / Carrier Mix	On-Chip Digital Up/Downconverters (DUC/DDC) via 48b NCO blocks
Direct Sampling Bandwidth	Multi-GSps Data Converters covering up to 1st and 2nd Nyquist Zones directly
Data Stream Routing	AXI4-Stream channels matching internal hardware generation sources
PL Speed-Gap Management	Super Sample Rate (SSR) parallel deserialization buses
Software/Firmware Link	Memory-Mapped AXI4-Lite registers under Processing System (PS) execution
State Storage Capacity	Cascaded Hardened UltraRAM (URAM) true dual-port on-chip structures

14. RFSoc Family: Product Generations and Variants

Gen 1, Gen 2, Gen 3, and RFSoc DFE Overview

The Zynq UltraScale+ RFSoc product line has evolved through four distinct generations, each introducing architectural improvements and expanded channel counts.

- **Gen 1 (ZU2XDR):** The inaugural RFSoc generation released in 2018. Introduced dual RF-ADC tiles and quad RF-DAC tiles with baseline multi-GSps sampling capabilities.
- **Gen 2 (ZU3XDR):** Enhanced performance with updated clock distribution and tighter timing specifications. Maintains dual RF-ADC and quad RF-DAC tile architecture.
- **Gen 3 (ZU4XDR):** The latest mainstream variant offering both dual and quad RF-ADC tiles, along with flexible dual and quad RF-DAC tile configurations. Improved interleaving factors and higher maximum sampling rates.
- **RFSoc DFE (ZU6XDR):** A specialized variant targeting high-frequency communications applications, featuring dedicated digital front-end (DFE) processing subsystems.

RFDC Tile Architecture

The data converters in the RFSoc are organized into **Tiles**, with each tile containing a local **Phase-Locked Loop (PLL)** to generate high-speed sampling clocks.

- **RF-ADC Tiles:** Dual tiles contain 8 interleaved sub-ADCs per channel; Quad tiles contain 4 interleaved sub-ADCs per channel. Time-interleaving increases effective sampling rates without sacrificing resolution.
- **RF-DAC Tiles:** Quad tiles are standard across Gen 1 and Gen 2; Gen 3 adds flexible dual tile variants. Each tile hosts a dedicated DUC (Digital Up Converter) per channel for baseband-to-RF translation.

Device Speed Grades and Operating Temperature Ranges

Speed grades (e.g., 2E, 2I, 2LE, 2LI, 1E, 1I, 1M, 1LI) encode both performance tier and thermal rating:

- **Numerical Prefix (2 vs. 1):** Indicates relative switching speed and maximum internal clock frequency of logic elements in the PL.
- **Letter Suffix:**
 - **E (Extended):** 0°C to +100°C industrial operating range.
 - **I (Industrial):** –40°C to +100°C wide industrial range.
 - **M (Military):** –55°C to +125°C extended range for aerospace/defense.
 - **L (Low-Power):** Indicates reduced core voltage variants to minimize static power dissipation.

FPGA Device Specifications

The typical PL fabric frequency is limited to a few hundred MHz (e.g., 625 MHz), whereas RF data converter clock rates scale into the multi-GSps range. This fundamental speed mismatch necessitates techniques like SSR (Super Sample Rate) deserialization and careful architecture planning.

15. Processing System Architecture and Inter-Processor Coordination

The Processing System (PS) of a modern RFSoc integrates multiple specialized processing cores:

Application Processing Unit (APU)

The **APU** contains a Quad-core ARM Cortex-A53 processor optimized for user-level application code and high-level system orchestration.

- Runs a fully featured operating system such as **Ubuntu Linux**, providing a standard software development environment.

- Hosts network services (e.g., Python-based TCP/IP socket servers) that receive experimental pulse configurations from remote PCs.
- Manages translation between user-facing pulse specifications and low-level hardware memory maps.
- Per-core floating-point and NEON SIMD engines enable DSP-like computations at the application layer.
- Dedicated L1 caches (32 KB Instruction, 32 KB Data) and a shared 1 MB L2 cache reduce memory-access latencies.

Real-Time Processing Unit (RPU)

The **RPU** contains a Dual-core ARM Cortex-R5 array specifically designed for **deterministic, low-latency control loops**.

- Can operate in **split mode** (two independent controllers) or **lock-step mode** (redundant, safety-critical execution).
- Features **Tightly Coupled Memory (TCM)**, allowing cache-miss-free, deterministic execution times—a critical feature for quantum gate timing.
- Executes either bare-metal firmware or a lightweight Real-Time Operating System (RTOS) to coordinate precise hardware trigger sequences.

Platform Management Unit (PMU)

An isolated, dedicated processor core that manages power gating, thermal monitoring, device initialization, and safety parameters independently of the APU and RPU.

16. Programmable Logic Fabric: Detailed Component Overview

Configurable Logic Blocks (CLBs)

The fundamental reconfigurable logic elements that compose the FPGA fabric. CLBs implement combinational logic (multiplexers, lookup tables) and sequential state machines (flip-flops, registers).

Digital Signal Processing (DSP) Slices

DSP48E2 hardened blocks within the PL provide specialized arithmetic operations:

- **27 × 18-bit Multiplier:** High-speed multiply capability.
- **48-bit Accumulator:** Efficient sum-of-products operations for filtering and matrix computations.
- **Pre-Adder:** Enables efficient multiply-accumulate (MAC) and convolution operations.

- These operate at clock speeds up to several hundred MHz, enabling real-time digital filtering, complex mixing, and linear algebra operations for quantum state discrimination.

Block RAM (BRAM)

Dedicated, high-speed on-chip memory blocks optimized for dual-port access patterns.

- Each BRAM can store up to 36 Kb of data.
- Configurable as true dual-port RAM, single-port ROM, or FIFO buffers.
- Can be dynamically reshaped (e.g., 4096×8 -bit, 8192×4 -bit, or other configurations).
- Useful for small to medium pulse sequence buffers and intermediate signal buffering.

UltraRAM (URAM)

High-capacity, high-density on-chip memory cells designed to replace external DRAM for certain workloads.

- Each URAM block stores up to 288 Kb with a fixed configuration (4096×72 -bit, dual-port).
- Multiple URAM blocks can be cascaded to create deep on-chip pulse playback buffers for storing complex quantum sequences.
- Eliminates external memory latency penalties, critical for tight feedback loops in quantum readout.

17. Multi-Gigasample Bandwidth and Frequency Coverage

Direct Sampling in the 1st Nyquist Zone

When an RF-ADC operates at $f_s = 5$ GSps, its **1st Nyquist Zone** extends from 0 Hz to $\frac{f_s}{2} = 2.5$ GHz.

Any signal in this passband can be directly digitized without aliasing. A low-pass anti-aliasing filter should precede the ADC input to block out-of-band noise.

Example Application: Sampling qubit drive tones at ~ 4 GHz is not possible in the 1st Nyquist zone; the 2nd Nyquist zone must be used.

Under-Sampling in Higher Nyquist Zones

The **2nd Nyquist Zone** spans $\frac{f_s}{2}$ to f_s (i.e., 2.5 GHz to 5.0 GHz for 5 GSps sampling).

By placing a **bandpass filter** before the ADC input to isolate a high-frequency band (e.g., 4 – 5 GHz) and allowing **aliasing**, those signals can be digitized and downconverted into the baseband zone for processing.

Comparison to Traditional Systems:

- **Conventional Heterodyne Architectures:** Require local oscillators, analog mixers, and image-rejection filters to downconvert RF signals to an intermediate frequency (IF, typically tens to hundreds of MHz) before sampling at much lower rates.
 - **RFSoc Direct-Sampling Approach:** Digitizes RF or IF signals directly at multi-GSps rates, then performs all frequency translation in software-defined digital mixers. This is more flexible, reduces analog complexity, and improves dynamic range.
-

18. Development Boards and Their Practical Integration

Mainstream Development Boards: ZCU208

The **ZCU208** is a comprehensive evaluation platform built around the Gen 3 **XCZU48DR** RFSoc.

- **RF Channels:** 8 RF-ADC channels (up to 5 GSps) and 8 RF-DAC channels (up to 9.85 GSps), each with 14-bit resolution.
- **Connectivity:** High-speed RFMC connectors route signals; external breakout cards convert to standard SMA connectors.
- **Add-On Ecosystem:**
 - **Balun Cards:** Convert differential signals to single-ended for standard lab instruments.
 - **Loopback Cards (XM650):** Enable internal RF-DAC-to-RF-ADC connections for firmware testing and calibration.
 - **Low-Jitter Clock Cards:** Supply stable reference clocks to all on-chip PLLs.

Academic and Compact Boards: RFSoc 4x2

The **RFSoc 4x2** is purpose-built for academic labs, coursework, and resource-constrained applications.

- **Limited Channel Breakout:** 4 RF-ADC paths, 2 RF-DAC paths (a subset of the full die's capability).
- **Integrated Baluns:** On-board balun networks route directly to SMA connectors, eliminating the need for external breakout cards.
- **Cost and Complexity Reduction:** Significantly more affordable than the ZCU208, making it suitable for university research groups and student projects.

- **Direct Quantum Integration:** Ideal for small-scale quantum processor control with 2-4 qubits or parallel readout channels.
-

19. Real-World Quantum Control Setup Using RFSoc

The diagram below illustrates a complete two-qubit superconducting quantum processor control and measurement chain integrated with an RFSoc development board.

Signal Routing and Circuit Components

1. **Drive Path to Qubits:**
 - RF-DAC 0 generates envelope-shaped drive pulses centered near the qubit transition frequency ($\omega_{01} \approx 4 - 6$ GHz).
 - Attenuators in the cryogenic lines suppress thermal noise photons coupling back to the qubits.
 - Both qubits receive drive signals that can be individually tuned and scheduled.
2. **Readout Probe Path:**
 - RF-DAC 1 generates a continuous readout probe tone near the resonator frequency ($\omega_r \approx 6 - 7$ GHz).
 - The tone couples into the readout resonator, which is dipole-coupled to both qubits.
 - The qubit state modulates the phase and amplitude of the reflected probe signal.
3. **Signal Isolation and Amplification:**
 - **Circulators** act as passive one-way valves: they allow the reflected readout signal to travel upward while blocking thermal noise from traveling back down toward the qubits.
 - A **cryogenic HEMT Low-Noise Amplifier (LNA)** amplifies the weak (μ V-scale) reflected signal by ~ 40 dB, bringing it to measurable levels at room temperature.
4. **Signal Acquisition and Real-Time Discrimination:**
 - The amplified signal enters the RF-ADC 0 input.
 - The on-chip DDC downconverts the signal to digital I/Q components at baseband.
 - The PL fabric runs a real-time state discrimination algorithm, comparing the I/Q values against trained thresholds to determine the final qubit state(s).
 - Results are streamed back to the PS for logging and subsequent gate synthesis decisions.


System Advantages for Quantum Applications

- **Unified Control:** Single chip manages both qubit drive (DAC) and measurement (ADC) with inherent synchronization.
- **Low Latency:** On-chip interconnects minimize feedback delay, enabling fast error correction and active reset.
- **High Sampling Rate:** Direct RF sampling at multi-GSps provides rich signal information for state discrimination.
- **Flexibility:** Software-defined DUCs and DDCs allow rapid reconfiguration of carrier frequencies and filter parameters across many qubits.

Review of Digital Signal Processing

1. Analog-to-Digital Conversion (ADC)

Analog-to-Digital Conversion is the foundational process of translating a continuous physical waveform $z(t)$ into a discrete-time, discrete-amplitude digital equivalent suitable for computational processing. This conversion is achieved through two consecutive operations: **Sampling** and **Quantization**.



2. Sampling Dynamics

Sampling discretizes a continuous-time signal by measuring its instantaneous amplitude at uniform time intervals.

- **Sampling Period (t_s):** The constant time interval between consecutive measurements.
- **Sampling Rate (f_s):** The frequency at which the analog signal is measured, defined as the inverse of the sampling period:

$$f_s = \frac{1}{t_s}$$

The choice of f_s establishes the limits of the system's capture bandwidth and determines how closely the digital representation tracks the original physical signal.

The Effect of Sampling Rate Selection

Consider a baseline 100 Hz sinusoidal wave sampled at three different frequencies:

1. High Sampling Rate ($f_s = 2 \text{ kHz}$) Sampling at 2 kHz yields 20 samples per period ($\frac{2000 \text{ Hz}}{100 \text{ Hz}}$). The digitized points provide an accurate representation of the signal's true time-domain profile.

2. Moderate Sampling Rate ($f_s = 500 \text{ Hz}$) Sampling at 500 Hz yields 5 samples per period. This provides a balanced trade-off, preserving the essential signal characteristics while reducing the storage and processing burden on the digital hardware.

3. Low Sampling Rate ($f_s = 80 \text{ Hz}$) Sampling below the signal frequency (providing fewer than 2 samples per period) alters the perceived waveform. The sampled data points track a false lower frequency. This distortion is known as **Aliasing**.

In this under-sampled scenario, the apparent frequency of the aliased signal (f_{alias}) matches the difference between the true signal frequency and the sampling rate:

$$f_{\text{alias}} = |f_{\text{sine}} - f_s| = |100 \text{ Hz} - 80 \text{ Hz}| = 20 \text{ Hz}$$

3. Signal Characterization and Frequency Classification

Before applying sampling limits, signals are categorized by their spectral distribution:

- **Baseband Signals:** Signals whose spectral energy begins at or near 0 Hz and extends to a given cutoff frequency (e.g., raw audio or unmodulated data streams).
 - **Bandpass Signals:** Signals whose spectral content is entirely contained within a specific frequency band that does not border 0 Hz. This profile is typical of high-frequency modulated communication carriers or raw qubit readout responses.
 - **Bandlimited Signals:** A general classification for any signal whose total spectral energy is restricted to a defined range between an upper boundary (f_h) and a lower boundary (f_l).
-

4. The Nyquist-Shannon Sampling Theorem

To ensure a continuous bandlimited signal can be perfectly reconstructed from its discrete samples without information loss, the sampling rate must exceed

twice the maximum frequency or bandwidth of the signal.

Baseband Formulation

For standard baseband signals extending up to a maximum frequency $f_{\max}$, the minimum sampling criteria (the **Nyquist Rate**) is defined by:

$$f_s > 2f_{\max}$$

If f_s falls below this threshold, all spectral components exceeding $\frac{f_s}{2}$ cross into adjacent bands, causing irreversible aliasing distortion.

Bandpass Formulation

For a bandpass signal restricted between a low frequency f_l and a high frequency f_h , the criteria adapts to the absolute bandwidth of the signal ($\Delta f = f_h - f_l$):

$$f_s > 2(f_h - f_l)$$

This relationship allows the system to use a much lower sampling rate than direct baseband sampling would require. It relies on the technique of **Undersampling**, placing the signal within specific frequency zones called **Nyquist Zones**.

5. Nyquist Zones and Aliasing Mechanics

The sampling rate f_s divides the frequency spectrum into equal structural segments or intervals of width $\frac{f_s}{2}$. These intervals are called Nyquist Zones.

- **First Nyquist Zone:** Ranges from 0 to $\frac{f_s}{2}$. This is the primary window for direct sampling without aliasing.
- **Second Nyquist Zone:** Ranges from $\frac{f_s}{2}$ to f_s . Any signal component in this zone is folded or mirrored into the first zone upon sampling.
- **Third Nyquist Zone:** Ranges from f_s to $\frac{3f_s}{2}$. Signals fold forward into the first zone, maintaining their original spectral orientation.

General Spectral Mapping Rule

The relationship mapping any absolute physical frequency f to its alias location f_{alias} within the primary first zone ($0 \leq f_{\text{alias}} \leq \frac{f_s}{2}$) is governed by:

$$f_{\text{alias}} = |f - nf_s|$$


Where n is an integer chosen such that the resulting frequency falls within the bounds of the first Nyquist zone.

6. Undersampling Applications in Quantum Control

Superconducting qubits operate at microwave frequencies in the gigahertz range (4–8 GHz). Digitizing these signals using direct baseband sampling would require ultra-high-speed converters operating above 16 GSps, which adds substantial power, computational, and financial overhead.

Strategic Implementation

Undersampling addresses this challenge by intentionally operating the ADC at a lower sampling rate than the absolute carrier frequency, provided the total signal bandwidth remains narrower than $\frac{f_s}{2}$ and fits entirely within a single Nyquist zone.



Downconversion Example

Consider a modulated qubit readout response spanning 1.6 GHz to 1.7 GHz (Bandwidth = 100 MHz), sampled using an ADC running at $f_s = 1$ GHz. The Nyquist zones are 500 MHz wide, placing the signal entirely inside the 4th Nyquist zone (1.5 GHz to 2.0 GHz).

Applying the mapping rule:

- The 1.6 GHz edge maps to: $|1.6 \text{ GHz} - 2(1 \text{ GHz})| = 400 \text{ MHz}$
- The 1.7 GHz edge maps to: $|1.7 \text{ GHz} - 2(1 \text{ GHz})| = 300 \text{ MHz}$

The undersampling process shifts the high-frequency microwave band down to an intermediate baseband window (300 – 400 MHz) inside the first zone. This approach achieves direct digital downconversion, eliminating the need for an analog mixing stage.

Critical Constraint: The original signal spectrum must fit entirely within a single Nyquist zone. If a signal crosses a zone boundary, its components fold back on top of each other from opposite directions, corrupting the digitized data.

7. Sampling Jitter

In a real-world ADC, the sampling clock experiences tiny timing variations rather than maintaining a perfectly uniform period t_s . This variation is known as **Sampling Jitter**.

Jitter introduces amplitude noise when sampling a time-varying signal. If a sample is taken slightly early or late by an offset δt , the measured voltage deviates from its true value by an error δV . This amplitude error is directly proportional to the signal's rate of change (slew rate):

$$\delta V \approx \frac{dz(t)}{dt} \delta t$$

Because high-frequency signals change rapidly over time, even minor timing jitter can cause significant amplitude errors. To maintain clean phase coherence and minimize this noise floor in high-frequency applications like qubit control, systems require ultra-stable, low-phase-noise clock sources.

8. Quantization Principles

Quantization transforms a continuous analog voltage sample into a discrete digital value by mapping it to the nearest available amplitude step.

Bit Resolution and Step Size

An N -bit ADC provides 2^N discrete vertical quantization levels distributed across its full-scale input voltage range (V_{FS}). For a peak operating boundary from $-V_{max}$ to $+V_{max}$, the total full-scale range is $V_{FS} = 2V_{max}$. The minimum voltage increment or step size (q) is defined as:

$$q = \frac{V_{FS}}{2^{N-1}}$$

[Image comparing quantization step resolution between a 4-bit and a 6-bit converter]

Increasing the bit depth N provides a finer grid of quantization steps, allowing the digital system to approximate the continuous analog signal more accurately.

The Ideal Quantization Error Model

The difference between the continuous analog voltage and its assigned discrete step is the quantization error. The maximum error occurs when a sample falls exactly halfway between two steps, limiting the absolute error to $\pm \frac{q}{2}$.

For a dynamic, complex input signal, this error is typically modeled as a uniform random variable with a flat Probability Density Function (PDF) across the step interval:

$$p(e) = \begin{cases} \frac{1}{q}, & -\frac{q}{2} \leq e \leq \frac{q}{2} \\ 0, & \text{otherwise} \end{cases}$$

Integrating the squared error across this distribution yields the theoretical **Quantization Noise Power** (n_{adc}):

$$n_{\text{adc}} = \int_{-\frac{q}{2}}^{\frac{q}{2}} e^2 \cdot p(e) de = \int_{-\frac{q}{2}}^{\frac{q}{2}} \frac{e^2}{q} de = \frac{q^2}{12}$$

This noise floor is distributed evenly across the primary baseband from 0 to $\frac{f_s}{2}$. Increasing the converter's bit depth reduces the step size q , which lowers the quantization noise floor and prevents weak signals from being masked.

9. Dynamic Range and Spectral Purity

Dynamic Range (DR)

The Dynamic Range defines the ratio between the largest reliable signal an ADC can measure without clipping and its smallest detectable increment. Expressed in decibels (dB) for an N -bit system, it follows the linear relationship:

$$\text{DR} = 20 \log_{10} (2^N) = 20N \log_{10}(2) \approx 6.02N \text{ [dB]}$$

- A **10-bit converter** provides a dynamic range of ≈ 60.2 dB.
- A **14-bit converter** improves this range to ≈ 84.3 dB.

Frequency Spurs and SFDR

If the input waveform is perfectly periodic (such as a pure sine wave), the quantization errors repeat in a predictable pattern rather than acting as random noise. This periodicity generates discrete, unwanted harmonic tones called **Spurs** in the frequency domain.

These spurs are evaluated using the **Spurious-Free Dynamic Range (SFDR)** metric, which measures the ratio in dB between the fundamental signal power and the highest unwanted spur component.

10. Digital Hardware Data Representation

To minimize logic gate counts, power consumption, and processing latency, high-speed digital hardware elements (like FPGAs and ASICs) typically process fractional values using **Fixed-Point Arithmetic** rather than floating-point engines.

The Q-Format Configuration

Fixed-point numbers use a structural format designated as **Q $n.b$** :

- **n :** The number of bits assigned to represent the integer component.
- **b :** The number of bits assigned to represent the fractional component.
- **Total Word Length:** Defined by the sum $n + b$.

Bits are ordered from the Most Significant Bit (MSB) on the left to the Least Significant Bit (LSB) on the right, with each position scaled by a binary weight based on its distance from the radix point:

$$\text{Weight} = 2^{\text{index}}$$

Structural Formats

1. **Unsigned Fixed-Point:** All bit positions have a positive numerical scale factor. The maximum representable value is $2^n - 2^{-b}$.
2. **Two's Complement Fixed-Point:** The MSB acts as a sign indicator with a negative weight (a value of 1 indicates a negative number). For a **Q $n.b$** two's complement layout, the operating bounds are:

$$\text{Minimum Value} = -2^{n-1}, \quad \text{Maximum Value} = 2^{n-1} - 2^{-b}$$

11. Fixed-Point Arithmetic Math Operations

Addition and Subtraction

When adding or subtracting two fixed-point numbers with aligned binary points, the operations follow standard integer arithmetic logic. However, combining two values can produce an output that exceeds the available bit range. To prevent overflow, the word length must expand by one bit ($n + 1$).

Example Case Consider adding two unsigned 8-bit numbers configured in a Q5.3 layout (5 integer bits, 3 fractional bits).

- **Maximum possible value:** $2^5 - 2^{-3} = 31.875$
- **Worst-case sum:** $31.875 + 31.875 = 63.75$

To safely store this maximum sum without clipping, the hardware requires a 9-bit output word configured as Q6.3:

$$\begin{array}{r} 11111.111_2 \quad (31.875) \\ + \quad 11111.111_2 \quad (31.875) \\ \hline 11111.110_2 \quad (63.750) \end{array}$$

Multiplication

Multiplying two fixed-point numbers increases both the integer and fractional bit widths. When a value with a $Qn_1.b_1$ layout is multiplied by a value with a $Qn_2.b_2$ layout, the raw product format expands to:

$$\text{Output Format} = Q(n_1 + n_2).(b_1 + b_2)$$

Example Case Multiplying two independent 8-bit unsigned Q5.3 values yields a product with a total width of 16-bits, split into 10 integer bits and 6 fractional bits (Q10.6). In practical system designs, this raw product is often truncated or rounded downstream to match the input bit width of subsequent processing stages.

12. Digital-to-Analog Conversion (DAC) Mechanics

A Digital-to-Analog Converter performs the inverse operation of an ADC, converting a sequence of discrete digital codes into a continuous analog voltage waveform.

The Zero-Order Hold (ZOH) Effect

After converting a digital value to an analog level, a standard DAC maintains that voltage constant for an entire clock cycle until the next sample update arrives. This step-like output behavior is known as a **Zero-Order Hold (ZOH)** function.

In the time domain, the ZOH response acts as a rectangular pulse of duration t_s . In the frequency domain, this rectangular response applies a **Sinc** ($\frac{\sin x}{x}$) envelope to the output spectrum, scaling the amplitude of the synthesized signal:

$$H_{\text{zoh}}(f) = e^{-j\pi f t_s} \left(\frac{\sin(\pi f t_s)}{\pi f t_s} \right)$$

[Image showing a DAC output spectrum shaped by the characteristic sinc roll-off envelope]

This sinc envelope introduces two primary spectral effects:


1. **Aperture Distortion:** It causes a non-linear amplitude roll-off that dampens higher-frequency signals within the first Nyquist zone, even if they are below the Nyquist limit ($\frac{f_s}{2}$).
2. **Image Frequencies:** The sharp edges of the step-like time-domain waveform generate unwanted image frequencies that repeat across higher Nyquist zones. These images are centered around integer multiples of the sampling frequency ($f_s \pm f_{\text{fundamental}}$).

13. Reconstruction Filters and Oversampling Strategies

Reconstruction Filtering

To smooth out the step-like ZOH waveform and eliminate high-frequency image components from the upper Nyquist zones, an analog **Reconstruction Filter** (a low-pass or band-pass filter) is placed at the DAC output.

If a synthesized signal occupies the entire first Nyquist zone up to $\frac{f_s}{2}$, the image frequency appears immediately adjacent to it. Separating the two requires an expensive analog filter with an extremely sharp cutoff. Because of this practical limitation, the maximum usable output frequency for standard DAC configurations is typically restricted well below the theoretical Nyquist limit.



Mitigation Strategies

To simplify the analog filter design and minimize aperture distortion, systems often employ two main techniques:

- **Oversampling:** Running the DAC core at a higher sampling rate increases the width of the Nyquist zones. This pushes the unwanted image frequencies further away from the target baseband signal, allowing the use of a simpler reconstruction filter with a gradual cutoff.
 - **Inverse Sinc Digital Filtering:** A digital pre-distortion filter can be implemented in the FPGA fabric ahead of the DAC. This filter applies an inverse sinc gain profile ($H(f) \propto \frac{\dots}{\sin x}$) that counteracts the ZOH roll-off, flattening the overall frequency response across the first Nyquist zone.
-

14. Essential DSP Formula Reference Sheet

DSP System Parameter	Mathematical Formulation
Sampling Period to Rate Mapping	$f_s = \frac{1}{t_s}$
Baseband Nyquist Sampling Criteria	$f_s > 2f_{\max}$
Bandpass Nyquist Sampling Criteria	$f_s > 2(f_h - f_l)$
Nyquist Zone Aliasing Map Location	$f_{\text{alias}} = f - nf_s \quad (0 \leq f_{\text{alias}} \leq \frac{f_s}{2})$
Discrete Linear Quantization Step Size	$q = \frac{V_{\text{FS}}}{2^N}$
Theoretical Quantization Noise Power	$n_{\text{adc}} = \frac{q^2}{12}$
Converter Decibel Dynamic Range	$\text{DR} = 20 \log_{10}(2^N) \approx 6.02N \text{ [dB]}$
Two's Complement Q-Format Bounds	$\text{Range} = [-2^{n-1} \dots 2^{n-1} - 2^{-b}]$
Addition Expansion Word Length	$\text{Width}_{\text{sum}} = n + 1$
Multiplication Product Word Length	$\text{Format}_{\text{prod}} = Q(n_1 + n_2).(b_1 + b_2)$
DAC Zero-Order Hold Frequency Envelope	$H_{\text{zoh}}(f) = \frac{\sin(\pi f t_s)}{\pi f t_s}$

Multirate Operations

1. Limitations of Conventional ADC Systems

In a standard Analog-to-Digital Converter (ADC) operating within the first Nyquist Zone ($0 \leq f \leq \frac{f_s}{2}$), the input stage must be preceded by an analog anti-aliasing low-pass filter.

The Anti-Aliasing Constraint

- **Spectral Preservation:** The filter isolates the target frequency components located within the first Nyquist Zone while attenuating all out-of-band high-frequency components.
- **The Aliasing Threat:** Without this analog stage, any high-frequency noise or signal component above $\frac{f_s}{2}$ folds back directly into the primary band, corrupting the target signal.
- **The “Brick-Wall” Ideal:** Perfect spectral isolation requires a non-causal “brick-wall” filter with an instantaneous transition band. This response cannot be physically constructed in the analog domain, nor can it be realized computationally as a digital filter.

Real-World Filtering Compromises

Because physical analog components have finite roll-off slopes, hardware engineers must accept a design trade-off:

1. **In-Band Attenuation:** The filter’s cutoff can be placed below the upper limit of the first Nyquist zone, which unintentionally dampens high-frequency components of interest.

2. **Aliasing Overlap:** The transition band can be allowed to extend past $\frac{f_s}{2}$ into the second Nyquist zone, which introduces a controlled level of aliasing into the digitized output.

As a filter’s specifications become tighter (requiring narrower transition bands and higher stopband attenuation), the physical circuit becomes significantly more complex, sensitive to component tolerances, and expensive to manufacture.

2. The Multirate Engineering Solution

To bypass the limitations of sharp analog filters, modern high-speed architectures use **Oversampling** paired with **Multirate Digital Signal Processing**.

Oversampling

Operating the ADC core at a sampling frequency significantly higher than the Nyquist rate of the target signal widens the first Nyquist zone. This creates a wide spectral buffer between the signal of interest and the point where aliasing occurs.


As a result, the strict transition requirements of the analog front-end filter are greatly relaxed, allowing the use of simpler, lower-cost analog circuits. Any out-of-band noise or signals that pass through the relaxed analog filter can then be removed in the digital domain using sharp, highly stable digital filters.

The Computational Bottleneck

While oversampling simplifies the analog hardware, it presents a challenge for the digital domain: streaming data into the FPGA fabric at an unnecessarily high sampling rate increases the clock speed requirements and computational load for all subsequent processing stages.

Multirate Operations

Multirate processing resolves this bottleneck by dynamically altering the data sampling rate at different stages of the processing chain. This approach ensures that the analog interfaces (ADCs and DACs) operate at high rates to simplify filtering, while the core digital signal processing algorithms (such as qubit state discrimination) run at lower rates near the Nyquist minimum to optimize power and hardware efficiency.



3. Core Multirate Mechanisms and Scenarios

Multirate operations primarily change a system's sampling rate through integer scaling factors:

- **Decimation (Factor M):** Reduces the sampling frequency by an integer factor M (e.g., downsampling a 10 MHz input stream by a factor of $M = 2$ yields a 5 MHz output stream).
- **Interpolation (Factor L):** Increases the sampling frequency by an integer factor L (e.g., upsampling a 10 MHz baseband stream by a factor of $L = 4$ yields a 40 MHz output stream).

Common Multirate Scenarios

- **Asynchronous Path Rate Matching:** Aligns two digital signal streams operating at different native sampling rates so they can be combined or mixed. For instance, in an audio mixing console, an 8 kHz voice path is upsampled by a factor of 6 ($L = 6$) to match a 48 kHz music background path.
- **Converter Interface Matching:** Matches the lower internal processing rates of an FPGA-based waveform generator to the multi-giga-sample operating rates of an RF-DAC core.
- **Analog Filter Optimization:** Pairs with an oversampled ADC to handle the bulk of the anti-aliasing task via software-defined digital filters, keeping the physical analog hardware requirements minimal.


4. Oversampled ADC and Decimation Pipelines

Decimation Framework

Decimation is the two-step architectural process of reducing a signal's sampling rate from an initial frequency f_s down to a lower decimation frequency f_d :

$$f_d = \frac{f_s}{M}$$

To reduce the rate without losing information, the decimation pipeline must perform two sequential tasks:



1. **Digital Anti-Aliasing Filtering:** A digital low-pass filter (typically implemented as a Finite Impulse Response, or FIR, filter) attenuates all high-frequency components above $\frac{f_d}{2}$ before the sampling rate is altered.
2. **Downsampling:** The downsampling block reduces the sampling rate by keeping only every M^{th} incoming sample and discarding the $M - 1$ intermediate samples.

Frequency and Time Domain Decimation Example

Consider an oversampled ADC input stream $x[k]$ with a sampling frequency $f_s = 2$ GHz and a target signal bandwidth limited to 480 MHz.

1. The Conventional Approach If the system uses a standard 1 GHz ADC, any signal or noise in the 520 MHz to 1 GHz band will alias directly into the 0 – 480 MHz target band. Preventing this requires an analog filter with an extremely sharp transition band spanning just 40 MHz (from 480 MHz to 520 MHz), which is difficult and costly to build.

2. The Oversampled and Decimated Approach Operating the ADC at $f_s = 2$ GHz widens the primary Nyquist zone up to 1 GHz, allowing the use of a relaxed analog filter.

After digitization, a digital decimation pipeline with a factor of $M = 2$ reduces the sampling rate to 1 GHz.

- **The Filter Action:** The digital FIR filter attenuates the components in the upper half of the spectrum (500 MHz to 1 GHz).
- **The Downsampler Action:** The downsampler drops every second sample, increasing the time interval between the remaining samples by a factor of 2 ($t_d = 2t_s$).

In the time domain, the sample indices shift from the high-rate notation k to a lower-rate index n , reflecting the reduced data volume. In the frequency domain, the spectral images shift closer together. However, because the digital filter cleared the upper frequency band before downsampling, no aliasing occurs within the final 0 – 500 MHz band.

5. Interpolated DAC Pipelines


At the transmit interface, a similar dual-stage approach is used to simplify the analog reconstruction or image-rejection filters at the output of a Digital-to-Analog Converter (DAC).

Interpolation Framework

Interpolation increases the sampling rate of a digital signal from an initial baseband frequency f_s up to an oversampled frequency f_u by an integer factor L :

$$f_u = L \cdot f_s$$

The interpolation pipeline performs these operations in the reverse order of the decimation pipeline:



1. **Upsampling:** The upsampler inserts $L - 1$ zero-valued samples between each pair of original incoming samples, increasing the data rate by a factor of L .
2. **Digital Image Rejection Filtering:** This digital low-pass stage removes the unwanted high-frequency spectral images generated by the upsampling step. These images appear at integer multiples of the original sampling rate up to the new folding frequency $\frac{f_u}{2}$.

Time and Frequency Domain Interpolation Example

Consider a baseband system driving an output carrier at 10 MHz with a native digital input update rate $f_s = 30$ MSps.

1. Un-Interpolated Output Spectrum Without interpolation, the DAC synthesizes the primary 10 MHz tone alongside its first high-frequency spectral image located at:

$$f_{\text{image}} = f_s - f_{\text{out}} = 30 \text{ MHz} - 10 \text{ MHz} = 20 \text{ MHz}$$

To isolate the 10 MHz tone, the analog reconstruction filter must feature a very sharp transition band that starts immediately at 10 MHz and reaches full attenuation by 20 MHz.

2. Interpolated Output Spectrum ($L = 2$) Passing the digital signal through an interpolation pipeline with a factor of $L = 2$ increases the internal update rate to 60 MSps.

- **Time-Domain Insertion:** The upsampler inserts one zero-valued sample between each original sample.
- **Digital Filtering Action:** The digital low-pass filter processes these zero-filled streams, interpolating the intermediate points to smooth out the waveform. This fills in the zero-valued gaps with correct intermediate amplitudes.
- **Spectral Shift:** The digital filter suppresses the intermediate image at 20 MHz. The first physical image produced at the DAC output now appears at:

$$f_{\text{image_new}} = f_u - f_{\text{out}} = 60 \text{ MHz} - 10 \text{ MHz} = 50 \text{ MHz}$$

The analog reconstruction filter's transition window is now significantly wider, spanning from 10 MHz up to 50 MHz. This allows designers to use a simpler, lower-order analog filter while maintaining high signal purity.

6. Structural Comparison Matrix

Design Vector	Decimation Pipeline (ADC Side)	Interpolation Pipeline (DAC Side)
Primary Purpose	Reduces data rate to save computing power	Increases data rate to simplify
Integer Scaling Factor	M	L
Sampling Frequency Mapping	$f_d = \frac{f_s}{M}$	$f_u = L \cdot f_s$
Order of Internal Operations	1. Digital Low-Pass Filter 2. Downsampler (Keep 1 of M)	1. Zero-Insertion Upsampler 2. Digital Low-Pass Filter
Filter Functional Assignment	Anti-Aliasing (Clears upper bands)	Image Rejection (Smooths out)
Time Domain Impact	Discards unnecessary data samples	Interpolates intermediate values

RF Data Converters

1. RF Data Converters: Analogue to Digital

Analogue to Digital Converter (ADC) Fundamentals

The core process of Analog-to-Digital conversion consists of measuring a continuous physical signal at fixed time intervals (f_s) and mapping those continuous voltage values to a finite set of discrete values via **Quantization**.

- **Resolution Limits:** An n -bit ADC discretizes the operating voltage range into 2^n distinct levels. For example, Generation 3 RFSoc devices feature 14-bit precision, translating to $2^{14} = 16,384$ digital bins.
- **Quantization Step (Δ):** The minimum voltage difference between adjacent discrete steps is defined by:

$$\Delta = \frac{V_{\max} - V_{\min}}{2^n}$$

The measured analog voltage must be rounded to the closest available quantization step, which introduces structural quantization noise.

Superconducting Qubit Readout Context

In superconducting quantum computing architectures (e.g., Transmon Qubits), control and readout operations rely heavily on high-frequency microwave signals.

- **Readout Frequencies:** Qubit state readout is typically performed via an on-chip microwave resonator coupled directly to the qubit. These readout signals populate the microwave spectrum, usually ranging from **5 GHz to 8 GHz**.
- **Hardware Demands:** Directly processing these signals without bulky, drift-prone analog downconversion mixers requires specialized, ultra-high-speed RF-ADCs capable of running at multi-giga-samples per second (multi-GSps).

2. Multi-GHz Nyquist Zones and Direct Sampling

According to the Nyquist-Shannon criteria, direct, unambiguous sampling requires the sampling frequency f_s to exceed twice the maximum frequency component of the input signal.

First Nyquist Zone Direct Capture

The frequency window spanning from 0 Hz to $\frac{f_s}{2}$ constitutes the **First Nyquist Zone**.

- Because modern RF-ADCs operate at multi-GHz clock rates, their first Nyquist zone naturally spans several gigahertz.
- Any analog signal residing entirely within this first zone can be directly digitized by the ADC core without requiring external analog mixers, IF stages, or local oscillators.

Direct Baseband Example Consider an RF-ADC clock rate of $f_s = 4$ GHz. The primary Nyquist zone spans from 0 to 2 GHz. Any signal in this range is digitized directly. To prevent out-of-band high frequencies from folding down and corrupting the data, an analog anti-aliasing low-pass filter with a sharp cutoff near 2 GHz must precede the converter.

Second Nyquist Zone Undersampling

While standard systems use filters to avoid aliasing entirely, RF-ADCs can intentionally exploit aliasing to fold high-frequency signals from upper zones down into the first Nyquist zone. This technique is known as **Undersampling** or **Sub-sampling**.

- **Overcoming Clock Limits:** Generation 3 RF-ADCs feature a maximum native sampling rate of 5 GSps (implying a first Nyquist zone boundary of 2.5 GHz). However, qubit resonators typically operate at higher frequencies (5 – 8 GHz).
- **Direct RF Sampling:** By using a high-quality analog **bandpass filter** instead of a low-pass filter, the first Nyquist zone is cleared of baseband signals. This allows a signal from the Second Nyquist Zone ($\frac{f_s}{2}$ to f_s) to fold cleanly into the first zone without overlap or interference.

Spectral Inversion Warning: When a signal from an even Nyquist zone (like Zone 2 or Zone 4) folds down into the first zone, its spectral orientation is inverted (flipped left-to-right). This spectral inversion must be mathematically corrected downstream in the digital domain.

Practical Frequency Boundaries

While higher odd Nyquist zones (Zone 3, Zone 5, etc.) fold into the first zone with their original spectral orientation, the physical input stages of an RF-ADC impose practical frequency limits.

The internal analog input buffers, sample-and-hold circuits, and packaging track-lines attenuate signals at higher frequencies. The certified analog input bandwidth boundaries across AMD Xilinx RFSoc generations are:

- **Generation 1:** Characterized up to 4 GHz
- **Generation 2:** Characterized up to 5 GHz
- **Generation 3:** Characterized up to 6 GHz

Note: Custom boards like the RFSoc 4x2 utilize specialized input baluns capable of handling an extended bandwidth up to 10 GHz.

3. Analogue Front-End Filtering Strategies


Choosing an analog front-end filter design requires balancing noise performance against system flexibility:

Filter Architecture	Fixed Application Approach	Qubit Characterization Approach
Filter Classification	Narrowband RF Bandpass Filter	Full-Zone Low-Pass / Bandpass Filter
Passband Width	Matched tightly to the signal bandwidth	Spans the entire target Nyquist Zone
Noise Performance	Excellent out-of-band rejection and SNR	Higher broadband noise floor integration
System Flexibility	Rigid; cannot track frequency shifts	Highly adaptive for multi-qubit tuning

- **Zone 1 Configuration:** Uses a low-pass filter with a sharp cutoff near $\frac{f_s}{2}$ to block all higher-order image components from folding down.
- **Zone 2 Configuration:** Uses a dedicated bandpass filter spanning $\frac{f_s}{2}$ to f_s . This filter clears out baseband noise from Zone 1 and blocks higher-order content from Zone 3 and above, ensuring that only the target Zone 2 signal folds cleanly into the digitized spectrum.

4. Hardware Architecture: Interleaved Sub-ADC Tiles

To achieve multi-giga-sample capture rates without requiring ultra-fast, high-power single-core converters, RFSoc devices employ a hardware strategy called **Time-Interleaved ADCs**.



The Interleaving Principle

A high-speed RF-ADC channel is built from an array of m slow internal sub-ADC cores running in parallel. They share a single analog input line but are clocked at uniformly staggered time intervals.

By shifting the clock phase of each successive sub-ADC core, the effective sampling rate of the combined channel increases by a factor of M (the interleaving factor) relative to the clock rate of an individual sub-ADC core:

$$f_{s,\text{effective}} = m \times f_{\text{sub-clock}}$$

To ensure accurate sampling, the precise phase relationship for the n^{th} sub-ADC core must satisfy:

$$\theta_n = \frac{2\pi(n-1)}{m}, \quad \text{where } n = 1, 2, \dots, m$$

Tile Structures: Dual vs. Quad Layouts

The RFDC IP block organizes these converters into groups called **Tiles**. The maximum sampling speed of a tile depends directly on its sub-ADC interleaving factor:


- **Dual-Channel Tiles:** Each RF-ADC channel integrates **8 interleaved sub-ADCs**.
- **Quad-Channel Tiles:** Each RF-ADC channel integrates **4 interleaved sub-ADCs**.

Because Dual-Channel Tiles use twice as many interleaved sub-ADC cores per channel, they can run at twice the maximum sampling speed of Quad-Channel Tiles within the same hardware generation:

RFSoC Generation & Device	Channel Type	Total Internal Channels	Maximum Sampling Rate
Gen 1: ZU28DR	4× Dual Tiles	8 ADCs	4.096 GSps
Gen 2: ZU39DR	4× Quad Tiles	16 ADCs	2.220 GSps
Gen 3: ZU43DR	4× Single Tiles	4 ADCs	5.0 GSps
Gen 3: ZU46DR	Quad / Dual Mixed	Variable	Quad: 2.5 GSps / Dual: 5.0 GSps
Gen 3: ZU48DR	4× Dual Tiles	8 ADCs	5.0 GSps
Gen 3: ZU49DR	4× Quad Tiles	16 ADCs	2.5 GSps

5. The Digital Downconverter (DDC) Pipeline

Once the analog signal is digitized by the interleaved sub-ADC array, it enters a dedicated hardware processing pipeline. This pipeline shapes, downconverts, and filters the high-speed data stream before passing it to the Programmable Logic (PL) fabric.



Generation 3 Frontend Additions: The DSA

In Generation 3 devices, a programmable **Digital Step Attenuator (DSA)** is integrated directly ahead of the analog input buffers.

The DSA dynamically scales the amplitude of incoming RF signals to utilize the ADC's full dynamic range. This optimization helps prevent harmonic distortion, signal clipping, and unexpected over-voltage conditions.

Quadrature Modulation Correction (QMC)

When processing complex physical I/Q channels, minor imbalances in gain or phase between the twin paths can degrade signal integrity. The QMC block provides real-time digital calibration to correct for these phase and gain mismatches, ensuring high orthogonal balance.

Digital Complex Mixer and NCO Tuning

The digital complex mixer shifts the high-frequency input signal down to baseband by multiplying it with coherent sine and cosine waveforms. These local reference signals are generated by a highly tunable **48-bit Numerically Controlled Oscillator (NCO)** operating on heterodyne principles.

The mixer can be configured to run in three operational modes:

1. **Bypass Mode:** Disables the internal mixing logic completely. The raw, real-valued digital samples pass through unchanged.
2. **Coarse Mode:** Limits frequency translation to a fixed set of fractional carrier options ($\frac{f_s}{2}$, $\frac{f_s}{4}$, or $-\frac{f_s}{4}$). This hardware configuration bypasses the NCO, which saves power and reduces latency.
3. **Fine Mode:** Enables the full 48-bit NCO, allowing the mixer to translate any arbitrary center frequency between $-\frac{f_s}{2}$ and $+\frac{f_s}{2}$ down to baseband.

Correcting Inverted Spectra When sampling a signal from an even Nyquist zone, its folded spectrum is inverted within the first zone. To correct for this, engineers can configure the NCO with a **negative tuning frequency**. This mathematical adjustment automatically flips the spectrum back to its correct orientation during downconversion.

Programmable Decimation Engine

After downconversion, the signal is routed through a programmable decimation chain. This stage reduces the data sampling rate to match the signal's actual baseband bandwidth, lowering the processing burden on the FPGA logic.

The decimation chain is built from a cascade of digital half-band low-pass filters. Each active filter stage reduces the sample rate by half while attenuating any out-of-band noise or signals above the new Nyquist limit.

- **Gen 1 & Gen 2 Performance:** Supports integer decimation factors up to $8\times$ ($1\times, 2\times, 4\times, 8\times$).
 - **Gen 3 Performance:** Features an expanded filter cascade that supports decimation factors up to $40\times$ ($1\times, 2\times, 4\times, 6\times, 8\times, 10\times, 12\times, 20\times, 24\times, 40\times$).
-

6. RF Data Converters: Digital to Analogue

The role of an RF-DAC is to convert discrete-time digital samples from the programmable logic back into a continuous-time physical analog voltage wave.

Zero-Order Hold (ZOH) Distortions

To construct a continuous waveform from digital points, standard DACs use a **Zero-Order Hold (ZOH)** technique. The DAC latches and maintains the voltage level of each sample constant for the entire duration of the clock period ($t_s = \frac{1}{f_s}$).

This constant latching creates a sharp, step-like output waveform. In the frequency domain, this response can be modeled as a rectangular pulse, which applies a characteristic **Sinc** ($\frac{\sin x}{x}$) envelope to the output spectrum:

$$H_{\text{zoh}}(f) = \frac{\sin\left(\frac{\pi f}{f_s}\right)}{\frac{\pi f}{f_s}}$$

[Image comparing Normal ZOH mode roll-off with Mix-Mode frequency shaping]

The ZOH technique introduces two primary spectral challenges:

1. **High-Frequency Image Generation:** The sharp transitions of the step-like waveform generate unwanted image frequencies in upper Nyquist zones, centered at $n \cdot f_s \pm f_{\text{fundamental}}$.
2. **In-Band Aperture Roll-Off:** The sinc envelope causes a non-linear loss of amplitude across the first Nyquist zone. This roll-off results in a ≈ 4 **dB drop in gain** at the zone boundary ($\frac{f_s}{2}$), which represents a loss of more than 50% of the signal's output power.


Counteracting Amplitude Roll-Off

To counteract the amplitude drop caused by the sinc envelope, systems use two main approaches:

- **Oversampling:** Increasing the DAC's sampling rate relative to the target signal frequency keeps the signal close to DC, minimizing the effects of the sinc curve roll-off.
- **Inverse Sinc Digital Pre-Emphasis:** The digital data can be passed through an **Inverse Sinc Filter** prior to the DAC core. This filter boosts the gain of higher-frequency components to compensate for the ZOH roll-off, delivering a flat amplitude response across roughly 90% of the first Nyquist zone.

7. Dual-Zone DAC Modes: Normal vs. Mix-Mode

Just as an RF-ADC can sample high-frequency inputs via aliasing, an RF-DAC can exploit its high-frequency image components to synthesize output signals directly into upper Nyquist zones.



1. Normal ZOH Mode

In Normal Mode, the DAC holds each sample constant for the full duration of the clock cycle (t_s). This configuration concentrates the output energy within the **First Nyquist Zone**, making it ideal for baseband signal synthesis.

2. Mix-Mode

To optimize performance in the **Second Nyquist Zone** ($\frac{f_s}{2}$ to f_s), the DAC can be switched to Mix-Mode. In this configuration, the output pulse shape is inverted halfway through each clock period ($t_s/2$).

In the frequency domain, this phase inversion modifies the DAC's transfer function to follow a shifted sinc profile:

$$H_{\text{mix}}(f) = \frac{\sin^2\left(\frac{\pi f}{2f_s}\right)}{\frac{\pi f}{2f_s}}$$


This reshaping shifts the peak output energy out of the first zone and into the second Nyquist zone, providing a higher and flatter gain profile across that upper frequency band.

High-Frequency Synthesis Example Suppose a system needs to generate an 8 GHz pure sinusoidal wave using an RF-DAC running at $f_s = 10$ GSps. Direct synthesis within the first Nyquist zone is limited to 5 GHz.

However, synthesizing a 2 GHz tone inside the digital engine naturally creates a high-frequency image component at $f_s - f = 10 \text{ GHz} - 2 \text{ GHz} = 8 \text{ GHz}$ in the second zone. By operating the DAC tile in **Mix-Mode** and adding an external analog bandpass filter centered at 8 GHz, the 2 GHz primary tone is suppressed, and the 8 GHz image is isolated as a clean, high-frequency output.

8. The Digital Upconverter (DUC) Pipeline

The Generation 3 RF-DAC tiles feature an integrated, multi-stage Digital Upconverter (DUC) pipeline that prepares digital signals from the FPGA fabric for high-speed analog conversion.



Programmable Interpolation Engine

To match low-rate baseband signals from the FPGA to the high clock rates of the RF-DAC core, the pipeline uses a programmable interpolation engine. This engine can be configured to process either real-valued inputs or complex I/Q data pairs.

The interpolation process is handled by a cascade of four independent upsamplers and low-pass Finite Impulse Response (FIR) filters:

- **Stage 1 Selection:** Features three multiplexed, sharp-cutoff filters (**FIR1a**, **FIR1b**, **FIR1c**) providing native $2\times$, $3\times$, or $5\times$ interpolation rates. Only one filter in this initial block can be active at a time.

- **Downstream Stages:** Cascades through **FIR2**, **FIR3**, and **FIR4** to build higher overall interpolation factors.
- **Supported Interpolation Rates:** Gen 3 devices support a wide range of overall interpolation factors: $1\times$, $2\times$, $3\times$, $4\times$, $5\times$, $6\times$, $8\times$, $10\times$, $12\times$, $16\times$, $20\times$, $24\times$, and $40\times$.

Digital Complex Modulation Mixer

The digital complex mixer modulates baseband I/Q data up to a target intermediate frequency before it reaches the DAC core.

- Functions identically to the ADC mixer but operates in reverse to modulate signals rather than demodulate them.
- Integrates a **48-bit NCO** fine mixer alongside a low-power coarse mixer ($\frac{f_s}{2}, \pm \frac{f_s}{4}$).
- Supports multi-tile synchronization, allowing the NCO phase wheels across multiple independent tiles to align perfectly. This capability is critical for maintaining phase coherence across large multi-qubit control systems.

Generation 3 Image Rejection (IMR) Filter

Generation 3 architectures add an **Image Rejection (IMR) Filter** downstream from the core mixing and modulation stages.

- **Additional Interpolation:** Enabling the IMR filter introduces a final $2\times$ upsampling step, which can extend the system's maximum total interpolation factor from $40\times$ up to $80\times$.
- **Selectable Filter Profiles:** The IMR can be configured as a low-pass filter to retain the signal within the first Nyquist zone while suppressing the second-zone image, or as a high-pass filter to do the opposite.
- **Stopband Attenuation:** Both the low-pass and high-pass IMR profiles feature symmetric designs that deliver 60 dB of out-of-band stopband attenuation.

9. Comprehensive Architectural Layout Reference

Functional Capability	RF-ADC Processing Subsystem	RF-DAC Processing Subsystem
Native Resolution	14-bits	14-bits
Primary Operation	Downconversion / Decimation	Upconversion / Interpolation
Frontend Conditioning	Digital Step Attenuator (DSA) [Gen 3]	Variable Output Power (VO)
In-Band Error Calibration	Quadrature Modulation Correction (QMC)	Quadrature Modulation Cor
Frequency Shifting Engine	48-bit NCO / Coarse Heterodyne Mixer	48-bit NCO / Coarse Hetero
Rate-Change Factor Ranges	$1\times$ to $40\times$ Decimation [Gen 3]	$1\times$ to $80\times$ Interpolation (v
Target Nyquist Zones	Zone 1 (Direct) or Zone 2 (Aliased)	Zone 1 (Normal Mode) or Zo
High-Frequency Filter Step	Analog Front-End Anti-Alias Filter	Digital Image Rejection (IM

Semiconductor Qubit Electronics: The Elzerman Method

1. The Elzerman Device Architecture

The landmark experiment by Elzerman et al. (*Nature*, 2004) demonstrated the first single-shot readout of an individual electron spin trapped inside a semiconductor quantum dot. This architecture forms a foundational component for spin-based quantum computing processors.

Heterostructure and Electrostatic Confinement

- **The Material Platform:** The device is fabricated on a two-dimensional electron gas (2DEG) formed at the interface of a Gallium Arsenide/Aluminum Gallium Arsenide (GaAs/AlGaAs) heterostructure.
- **Quantum Dot (QDOT):** Surface metallic top-gates (M, R, T) apply negative electrostatic potentials to deplete the underlying 2DEG, isolating a sub-micron puddle of charge that acts as a artificial atom (Quantum Dot) capable of trapping individual electrons.
- **Plunger Gate (P):** A specialized electrode used to apply high-speed voltage pulses (ΔV_P). Shifting P directly translates the discrete chemical potential energy levels of the quantum dot up and down relative to the adjacent electron reservoir.
- **The Reservoir:** A large, continuous sea of electrons located to the left of the dot. It acts as an electron source and drain across a controllable tunnel barrier.

Charge Detection via Quantum Point Contact (QPC)

Because an individual electron carries a minute amount of charge, directly measuring the spin state requires a highly sensitive **spin-to-charge conversion** mechanism.

- **The QPC Sensor:** A narrow, nearby channel of electrons formed by separate electrostatic gates. The current flowing through the QPC (I_{QPC}) is extremely sensitive to changes in the local electrostatic environment due to capacitive coupling.
- **Operating Regime:** The QPC is biased below its first quantum of conductance ($G_0 = \frac{2e^2}{h}$). The highest sensitivity to charge displacement is achieved by biasing the channel halfway through its first conductance transition window.
- **Detection Signature:** When an electron leaves the quantum dot, the reduction in local negative charge causes an immediate, measurable increase in the current flowing through the QPC (ΔI_{QPC}).

2. Parameter Tuning and Zeeman Splitting

To operate the quantum dot as an isolated two-level spin system (a qubit), the energy degeneracy of the electron spin states must be lifted.

Magnetic Field and Zeeman Energy

The device is placed inside a dilution refrigerator and subjected to a strong in-plane magnetic field ($B = 10$ T). This magnetic field induces Zeeman splitting, separating the spin-up ($\uparrow$) and spin-down ($\downarrow$) energy levels by the Zeeman energy (ΔE_Z):

$$\Delta E_Z = g\mu_B B$$

Note: In GaAs, conduction electrons feature a negative effective g -factor ($g \approx -0.44$). This makes the spin-down state the excited state and the spin-up state the lower-energy ground state.

Thermal and Orbital Energy Relations

To ensure reliable operation without unwanted transitions, the system parameters must satisfy the following strict energy hierarchy:

$$\Delta E_{\text{orbital}} \gg \Delta E_Z \gg k_B T$$


- **Thermal Energy Counter-Selection:** At a dilution refrigerator temperature of $T = 0.3$ K, the thermal fluctuation energy is $k_B T \approx 25$ μeV . The Zeeman splitting at 10 T is $\Delta E_Z \approx 200$ μeV . Because $\Delta E_Z \gg k_B T$, thermal fluctuations cannot randomly excite an electron from the ground state to the excited state.
- **Orbital Cleanliness:** The spatial confinement energy spacing between orbital levels is $\Delta E_{\text{orbital}} \approx 1.1$ meV, and the charging energy is $\Delta E_C \approx 2.5$ meV. Because ΔE_Z is significantly smaller than these parameters, spin transitions occur entirely within the lowest orbital ground-state level.

3. The Two-Level Pulse Readout Protocol

The Elzerman protocol uses a multi-stage voltage pulse sequence applied to the plunger gate P to perform state preparation, initialization, and single-shot readout.

Plunger Voltage (V_P)

^



Stage 1: Empty the Dot

- **Action:** The plunger gate voltage is driven to a highly negative value.
- **Energy State:** Both the spin-up and spin-down energy levels of the dot are pushed well above the Fermi energy (E_F) of the adjacent electron reservoir.
- **Result:** Any electron currently occupying the quantum dot tunnels out into the reservoir, resetting the dot to a zero-electron ($N = 0$) state.

Stage 2: Inject and Wait

- **Action:** The plunger gate voltage is stepped upward to lower the dot's energy levels.
- **Energy State:** The potential well is lowered such that both the spin-up and spin-down levels drop below the reservoir's Fermi energy (E_F).
- **Result:** A single electron from the reservoir tunnels into the dot. Due to Coulomb blockade effects, the addition of this first electron shifts the energy cost for a second electron too high, preventing any further tunneling. The spin state ($\uparrow$ or $\downarrow$) of the injected electron is initially unknown.

Stage 3: Readout (Spin-to-Charge Conversion)

- **Action:** The plunger voltage is precisely adjusted to position the reservoir's Fermi level **exactly halfway** between the split spin-up and spin-down energy levels.
- **Case A: The electron is Spin Up ($\uparrow$ - Ground State)**
 - The spin-up energy level sits safely below the Fermi energy (E_F).
 - The electron does not have enough energy to leave the dot; it remains trapped.
 - *QPC Signature:* The QPC current (I_{QPC}) remains stable at a constant, lower baseline level.
- **Case B: The electron is Spin Down ($\downarrow$ - Excited State)**
 - The spin-down energy level sits above the Fermi energy (E_F).

- The electron tunnels out of the dot into the empty states of the reservoir, temporarily leaving the dot empty ($N = 0$).
- This empty state lowers the electrostatic barrier for entry, allowing a new electron with a spin-up state to tunnel into the dot from the reservoir.
- *QPC Signature:* The temporary $N = 0$ empty state causes a sharp, millisecond-scale pulse (a “bump”) in the QPC current (I_{QPC}) before dropping back to the baseline level once a spin-up electron enters.

4. Tuning the Read-Out Configuration

Because microfabrication defects and material variations create unpredictable localized potentials, the precise gate bias voltages must be determined experimentally before running the qubit protocol.

Mid-Gate (V_M) Bias vs. Plunger (V_P) Dynamics

- V_P (**Plunger Gate**): Handles high-speed AC perturbations, oscillating the dot’s internal energy levels up and down.
- V_M (**Mid Gate**): Sets the DC operating point (the bias window) where these energy level oscillations occur.

Experimental Calibration via Color-Map Analysis

To locate the optimal V_M bias point, engineers apply a test pulse sequence to V_P while sweeping V_M through an increasingly negative range (e.g., -1.12 mV to -1.13 mV). Observing the resulting change in QPC current (ΔI_{QPC}) over a 1 ms time window reveals four distinct operating regions:

Increasingly Negative V_M

```
|
|-- Zone 1: Dot is Always Full --> EF is too high; levels never lift above EF during empty
|
|-- Zone 2: Ideal Readout Zone --> Transition point ( $V_{M0}$ ) where spin-down tunnels but spin-up
|
|-- Zone 3: Dot is Always Empty --> EF is too low; both spin levels sit above EF, causing
|
|-- Zone 4: No Electron Entry --> EF is extremely low; the dot remains empty throughout
v
```

- **Zone 1 (Top Band):** ΔI_{QPC} simply mirrors the pulse voltage shape without showing any tunneling features. The Fermi level is too high relative to the dot’s levels, meaning the dot remains full ($N = 1$) across all pulse steps.

- **Zone 2 (Target Band):** Shows a clear 0.3 nA current step, signaling single-electron tunneling. Within this region lies the transition point (V_{M0}), where the spin-down level crosses the reservoir's Fermi level.
- **Zone 3 (Lower Band):** The QPC current shows that the dot empties during the readout stage regardless of the spin state. The Fermi level sits too low, causing both spin states to tunnel out and ruining the readout selection.
- **Zone 4 (Bottom Band):** The current behavior matches Zone 1, but the dot is completely empty. The Fermi level sits too low to allow an electron to enter the dot during the injection stage.

Finding the Optimal Operating Point

1. Identify the boundary voltage V_{M0} where the system transitions between Zone 2 and Zone 3. At this boundary, the spin-down level is perfectly aligned with the reservoir's Fermi energy (E_F).
2. Calculate the optimal operating bias by shifting V_{M0} by an amount corresponding to half the system's Zeeman energy ($\frac{\Delta E_Z}{2}$). This centers the Fermi energy directly halfway between the spin-up and spin-down levels during the readout stage.

5. Measurements and Spin Relaxation (T_1)

Statistically Confirming Spin State Detection

Because the tunneling process is governed by stochastic Poissonian statistics, the exact timing of the spin-down current bump varies from run to run. To verify that these current bumps represent true spin-down states rather than random noise, engineers measure the spin-down fraction while varying the delay time (t_{wait}) before readout.

As the waiting time increases, the detected spin-down fraction decreases exponentially. This decay occurs because the excited spin-down electrons naturally lose energy and relax into the spin-up ground state over time.

Spin-Orbit Coupling and the T_1 Lifetime

The characteristic timeline for this decay is called the **Spin Relaxation Time** (T_1). This lifetime is highly sensitive to the strength of the external magnetic field (B) due to spin-orbit interaction mechanisms:

- **At $B = 8$ T:** $T_1 \approx 0.55 \pm 0.07$ ms
- **At $B = 10$ T:** $T_1 \approx 0.85 \pm 0.11$ ms
- **At $B = 14$ T:** $T_1 \approx 0.12 \pm 0.03$ ms

A stronger magnetic field broadens the Zeeman gap, which can accelerate the relaxation rate and cause the spin-down fraction to decay more rapidly.


6. Readout Fidelity and Optimization Error Analysis

Readout Fidelity measures how accurately the digital measurement matches the true physical spin state of the electron.

Error Parameters (α and β)

Readout errors are categorized into two primary conditional probabilities:

$$\alpha = P(\text{"down"} \mid \text{spin-} \uparrow) \quad \text{and} \quad \beta = P(\text{"up"} \mid \text{spin-} \downarrow)$$



The α Error (False Positives) The system mistakenly flags a true spin-up ground-state electron as a spin-down state.

- **Thermal Activation:** High thermal energy can cause a spin-up electron to hop over the tunnel barrier into the reservoir.
- **Electrical Noise:** High-frequency electrical noise in the QPC lines can create a false spike that mimics a tunneling bump.

The β Error (False Negatives) The system mistakenly flags a true spin-down excited-state electron as a spin-up state.

- **Intra-Dot Relaxation (β_1):** The spin-down electron relaxes into a spin-up state before it has a chance to tunnel out of the dot.
- **Fast Double-Tunneling (β_2):** The spin-down electron tunnels out and a spin-up electron tunnels in so quickly that the QPC sensor's hardware interface misses the brief charge transition.

Minimizing Total Error

Optimizing readout fidelity requires choosing an intermediate current threshold that minimizes the sum of errors ($\alpha + \beta$):

- Setting the threshold **too low** causes the system to miss true spin-down bumps, increasing the β error.
- Setting the threshold **too high** causes random baseline noise to look like a spin-down transition, increasing the α error.

Hardware Improvement Strategies


- **To Reduce α :** Lower the electron temperature inside the dilution refrigerator to suppress thermally activated tunneling.
- **To Reduce β :** Increase the bandwidth of the QPC measurement electronics to capture ultra-fast double-tunneling events before relaxation occurs.

7. Advanced Verification: Randomized Benchmarking and QND

To scale individual single-shot measurements into multi-qubit quantum processors, systems require advanced verification and non-destructive readout protocols.

Randomized Benchmarking (RB)

While standard verification techniques can be skewed by State Preparation and Measurement (SPAM) errors, Randomized Benchmarking isolates gate errors by applying long sequences of random Clifford gates.



1. **Execution:** The system applies a random sequence of m Clifford gates, followed by a final calculated inversion gate designed to return the qubit to its initial starting state.
2. **Analysis:** The survival probability is plotted as a function of sequence length (m). Fitting this decay curve to an exponential model reveals the average gate fidelity, independent of errors in state preparation or readout. Evaluating this fit at $m = 0$ provides an accurate estimate of the baseline SPAM fidelity.

Quantum Non-Demolition (QND) Readout

The classic Elzerman method is inherently destructive: the target electron tunnels out of the dot during a spin-down readout event, destroying the qubit state.

To solve this problem, advanced architectures use **Quantum Non-Demolition (QND)** readout. This method preserves the target qubit's state by mapping it onto an adjacent ancilla qubit.

- **Mechanism:** A target spin qubit is coupled via an exchange interaction to a neighboring secondary (ancilla) spin qubit. The system applies a

conditional rotation (such as a controlled-NOT operation) to flip the ancilla qubit's spin only if the target qubit is in a specific state.

- **Result:** The ancilla qubit is read out using standard methods while the primary target qubit remains undisturbed inside its potential well. Repeating this non-destructive process multiple times boots the state measurement fidelity from a single-shot baseline of $\approx 80\%$ up to over 95%.

Coherent Manipulation and Characterization of Semiconductor Spin Qubits

Mariagrazia Graziano *mariagrazia.graziano@polito.it* Politecnico di Torino – VLSI Group

1. Core Qubit Operations: The Control Cycle

Scaling a semiconductor-based quantum processor requires three highly coordinated, repeatable steps executed via a tight digital-to-analogue hardware feedback loop:

1. **Initialize:** Place the two-level quantum system into a known, highly pure ground state (typically $|0\rangle$).
 2. **Manipulate:** Apply coherent electromagnetic or electrostatic perturbations to perform single- or multi-qubit gate operations.
 3. **Readout:** Determine the final projection state of the qubit (measuring the $|1\rangle$ state probability) using spin-to-charge conversion techniques like the Elzerman method.
 4. **Loop:** Reset and repeat the sequence (GOTO START) across thousands of identical runs to construct an accurate statistical quantum state tomography.
-

2. Electron Spin Resonance (ESR) and the Rabi Experiment

The **Rabi Experiment** is the fundamental diagnostic tool used to demonstrate the coherent time-evolution and oscillatory control of a two-level quantum system between its $|0\rangle$ and $|1\rangle$ states.

Electromagnetic Wave Generation

To drive transitions within a quantum dot, an alternating current (AC) is routed through an on-chip microwave transmission line placed immediately adjacent to the dot structure. This AC excitation generates a localized transverse oscillating magnetic field ($\vec{B}_{AC}$) perpendicular to the primary static sorting magnetic field ($\vec{B}_0$).

Magnetic Resonance Condition

According to the principles of magnetic resonance, the oscillating field $\vec{B}_{AC}$ along a single physical axis can be mathematically decomposed into two counter-rotating circular magnetic field vectors: one rotating clockwise and the other counter-clockwise on the Bloch sphere's equatorial plane.

- **Precession:** The electron spin precesses around the main static longitudinal axis ($\vec{B}_0$) at its native Larmor frequency, determined directly by its split Zeeman energy:

$$\Delta E_Z = h\nu_{\text{Larmor}}$$

- **Resonance:** When the microwave drive frequency matches ν_{Larmor} , the counter-rotating component that moves in the opposite direction averages out to zero over time. However, the component rotating in unison with the precessing spin becomes **stationary** relative to the qubit's rotating frame of reference.
- **The Frame Shift:** In this rotating frame, the static $\vec{B}_0$ vector is effectively canceled out, leaving only a stable, fixed magnetic field component pointing along a transverse axis. The spin vector immediately begins a secondary precession around this new axis, driving a deterministic angular rotation over the surface of the Bloch sphere.

The “Spin-Dance” Analogy

To visualize this dual-precession phenomenon, consider the following physical analogy:

The Coordinate Setup: Stand vertically to represent the longitudinal static magnetic field axis ($\vec{B}_0$). Keep your elbow tucked against your side and extend your hand outward at an angle; this arm represents the qubit's spin vector. **The Baseline Motion:** Spin around on your feet at a constant speed. This motion emulates the native Larmor precession of the qubit around the static field. **Introducing the Drive:** Extend your other arm horizontally to represent the transverse AC magnetic field component. If you spin your body at the exact same rate as this external drive field, your horizontal arm will look completely stationary from the perspective of your precessing arm. **The Double Precession:** Now, rotate your precessing arm up and down around that “stationary” horizontal arm while your entire body continues to spin on its feet. The resulting path traced by your hand through space forms a complex, overlapping **spiral pattern**.

3. Rabi Oscillations and Environmental Decoherence

When the resonant microwave pulse is applied for a controlled duration (called the **burst time**) and followed by an Elzerman readout, the probability of finding the qubit in the excited $|1\rangle$ state oscillates as a perfect sinusoidal function of time.

Controlling Rotation Speed

The frequency of this state-to-state oscillation—the **Rabi frequency** (Ω_{Rabi})—is directly proportional to the amplitude (power) of the applied AC control pulse. Higher drive power rotates the spin vector more rapidly across the Bloch sphere.

Decoherence and Environmental Noise


In real physical systems, these sinusoidal oscillations exhibit a characteristic exponential envelope decay. This signal damping is caused by **quantum decoherence**, where the qubit loses phase information due to interactions with its local environment.

- **The GaAs Bottleneck:** Gallium Arsenide (GaAs) heterostructures suffer from severe decoherence because every native isotope of Gallium and Arsenic possesses a non-zero nuclear spin. These surrounding nuclear spins flip randomly over time, creating a noisy, fluctuating local magnetic field (the Overhauser field). This fluctuation causes the qubit’s resonance frequency to drift stochastically, throwing the control pulses out of resonance.
- **Isotopic Purification via Silicon-28:** To suppress this primary decoherence path, architectures use **Silicon-28** (^{28}Si). Silicon can be isotopically purified to isolate the ^{28}Si isotope, which has an absolute **zero nuclear spin**. Eliminating the local nuclear magnetic background removes the main source of frequency drift, extending qubit coherence times by orders of magnitude.

Substrate Material	Nuclear Spin Profile	Experiment Type	Typical Coherence Time
GaAs / AlGaAs	High-density active nuclear spins	Ramsey Fringe	~ 10 ns
Purified Silicon-28	Zero nuclear spin background	Ramsey Fringe	~ 120 μs

4. Case Study: Cryo-CMOS Qubit Control via Intel Horse Ridge

The **Horse Ridge** cryogenic CMOS integrated circuit (developed by Intel and TU Delft) represents a major step forward in replacing bulky, room-temperature room instrumentation with a compact, silicon-based controller that operates directly inside the dilution refrigerator.



Core Architecture and Setup

- **The Device Under Test:** A double quantum dot defined inside a 10 nm thick $^{28}\text{Si}/\text{SiGe}$ quantum well.
- **Electrostatic Tuning:** Two main surface plunger gates—**LP (Left Plunger)** and **RP (Right Plunger)**—apply stable DC potentials to accumulate and isolate exactly one single electron underneath each gate electrode.
- **Magnetic Configuration:** An external static magnetic field of 380 mT is paired with an on-chip **micromagnet array** fabricated directly on top of the double-dot structure. This micromagnet creates a controlled magnetic field gradient, inducing a distinct Zeeman splitting for each qubit.

Characterizing Rabi Performance via Horse Ridge

To characterize the single-qubit gate capabilities of the Horse Ridge system, engineers use a simple test sequence:

1. **Initialize:** Tune the dot gates to settle the electron spin into the lower $|0\rangle$ ground state.
2. **Excite:** Apply a sharp, rectangular microwave control pulse generated directly by the SoC’s high-frequency output lines.
3. **Measure:** Read out the resulting spin projection state using the Elzerman method to evaluate the final $|1\rangle$ state probability.

By sweeping the duration of the microwave pulse, the system maps out clean Rabi oscillations. Testing across the chip’s dual RF channels yields precise operational metrics:

- **RF_{low} Output Channel:** Achieves a 1 MHz Rabi frequency operating at a carrier frequency of 13.4 GHz.
- **RF_{high} Output Channel:** Achieves a 400 kHz Rabi frequency operating at a carrier frequency of 17.5 GHz.


Improving Readout Visibility: The raw signal contrast (readout visibility) can be degraded by high-frequency electrical noise coupled onto the device’s control gates. To suppress this out-of-band noise and optimize state discrimination, a discrete analog **bandpass filter with a 2 GHz wide passband** is integrated directly into the chip’s output measurement path.

5. Dual-Axis Coherent Control: The Ramsey Experiment

While a standard Rabi sequence validates single-axis rotations, proving full, arbitrary control over a qubit requires demonstrating phase-coherent rotation capabilities across multiple orthogonal axes. This is achieved using a **Ramsey-style experiment**.

The Multi-Pulse Ramsey Sequence

The Ramsey protocol sequences multiple control pulses separated by a precise time interval to track phase evolution:



1. **First $\frac{\pi}{2}$ Pulse ($R_X(\frac{\pi}{2})$):** Applies a resonant rectangular microwave pulse along the X -axis. This turns the spin vector by exactly 90° , rotating it from the longitudinal pole down onto the Bloch sphere’s equatorial plane into a superposition state.
2. **Phase Rotation ($R_Z(\theta)$):** Rotates the spin vector around the longitudinal Z -axis by a variable angle θ , tracking the phase evolution between the qubit and the reference clock.
3. **Second $\frac{\pi}{2}$ Pulse ($R_X(\frac{\pi}{2})$):** Applies an identical X -axis rotation pulse. This second turn translates the accumulated phase angle back into a measurable longitudinal state population difference.

Physical Realization of Control Axes

- **X-Axis and Y-Axis Rotations:** Driven physically by pulsing a rectangular microwave burst (set to 13.7 GHz in this processor framework). The

duration of the microwave pulse is directly proportional to the intended rotation angle.

- **Z-Axis Rotations:** Rather than applying a physical pulse, Z-axis phase rotations are implemented via **Virtual Z-Gates**. The system updates the reference phase clock inside the controller’s digital signal processor with a precise phase offset (θ). This digital offset shifts the phase of all subsequent *X* and *Y* microwave pulses, tracking the qubit’s phase evolution without introducing hardware overhead or pulse calibration errors.

Coherent Validation Metrics

Plotting the measured $|1\rangle$ state probability against a Z-gate phase sweep from 0° to 360° maps out a clean, repeating **cosinusoidal fringe pattern**.

The close alignment between this experimental data and theoretical predictions confirms that the controller can manipulate the electron’s spin state with high phase coherence, validating its ability to execute any arbitrary single-qubit quantum gate.

Classical Electronics Interfaces for Quantum Processors

1. Controller Specifications across Qubit Modalities

Driving quantum processors requires high-performance classical electronic interfaces. The physical mechanisms used to isolate the quantum state dictate different architectural specifications for the controller’s analog front-end.

Transmon Qubits vs. Semiconductor Spin Qubits

The hardware control loop differs significantly between superconducting circuits and quantum dots:

Hardware Metric	Superconducting Transmons	Semiconductor Spin Qubits
Physical Core Composition	Shunt Capacitor SQUID Loop	Electrostatically confined 2DEG Puddles
Primary Control Drive	Phase-modulated Microwave Pulses	Baseband Voltage Pulses + Microwave
Primary Tuning Line	Flux-Bias Current Line	DC Plunger & Barrier Electrodes
Readout Mechanism	Dispersive Microwave Cavity Shift	Electrometer Impedance Modulation
Readout Output Path	RF Demodulation Engine	Cryogenic Baseband / RF-Reflectometry

Critical Signal Integrity Vectors

To maintain high gate fidelities and avoid degrading the qubit’s quantum coherence, driving signals must be tightly controlled across four primary domains:

DC Flux	Fast Flux	Combined Microwave
Frequency	Phase	Resonant Drive
Selection	Tuning	Channels
+-----+	+-----+	+-----+
	DILUTION REFRIGERATOR	
+-----+		+-----+

- **Longitudinal Frequency Tuning (Z-Axis):** Each transmon is coupled to a dedicated flux-bias line biased with a direct current (DC) source to set the qubit’s resonant frequency by modulating the SQUID loop’s magnetic flux. A separate **Arbitrary Waveform Generator (AWG)** generates fast flux pulses on each line for rapid frequency modulation during quantum operations.
- **Transverse Resonant Drive (X/Y-Axis):** A master AWG synthesizes intermediate frequency (IF) in-phase (*I*) and quadrature (*Q*) analog envelopes. These baseband channels drive an external **Vector Signal Generator (VSG)**, which performs I/Q modulation to upconvert the envelopes onto a microwave carrier frequency matching the qubit transition frequency. The resulting modulated carriers are combined through a multichannel combiner into a single coaxial feed before entering the dilution refrigerator.
- **Dispersive Readout Loop:** Readout tones are synthesized similarly using a VSG/AWG pair. The weak microwave signal reflected from the readout cavity passes through a multi-stage cryogenic amplification chain—including a **Josephson Parametric Amplifier (JPA)** at the base mixing-chamber stage (20 mK) and a cryogenic **Low-Noise Amplifier (LNA)** at the 4 K stage—to optimize the Signal-to-Noise Ratio (SNR) for short measurement times.

Semiconductor Spin Controller Architecture (e.g., 2-Qubit Silicon Rig)

A typical semiconductor setup comprises two quantum dots formed by metallic electrodes that electrostatically deplete the Two-Dimensional Electron Gas (2DEG). A Quantum Point Contact (QPC) located on the right side serves as the charge sensor for state readout. Electrons are loaded from a reservoir connected via ohmic contacts. Reference: Watson et al. ‘A programmable two-qubit quantum processor in silicon’. *Nature* 555.7698 (Feb. 2018).

The system requires highly synchronized baseband pulse delivery:

- **Gate Biasing & Initialization:** When no operation is performed, each quantum dot must contain a single electron at the same dot potential, and the tunnel barriers must be tuned to ensure negligible coupling between neighboring dots. These conditions are maintained using dedicated low-noise DC bias voltage generators connected to all gate electrodes (plunger

gates and barrier gates).

- **Single-Qubit ESR Control:** Resonant microwave control bursts for Electron Spin Resonance (ESR) are generated by a Vector Signal Generator (VSG) such as the Keysight E8267D. The microwave carrier is I/Q modulated using an AWG envelope (e.g., Tektronix 5014C) to produce the required pulse shapes. In cases where each qubit has a unique Larmor frequency induced by a localized micromagnet gradient, a single control line can control multiple qubits independently via **Frequency-Division Multiple Access (FDMA)** frequency multiplexing, thereby simplifying the system's wiring footprint.
- **Baseband Gate Pulsing:** Fast voltage pulses required for qubit initialization, readout tuning, and two-qubit exchange gates are generated by dedicated AWGs. These pulse streams are passed through low-pass filters and combined with the static DC bias lines using bias tees before connecting to the control gates. Distinct AWGs may be employed for two-qubit gates and readout, as their specifications can differ significantly.
- **Master Clock Synchronization:** To prevent timing skew across channels, an external master AWG distributes a hardware trigger line to synchronize the clocks of all sub-AWGs.
- **Electrometer Readout Signal Processing:** Readout relies on a Single-Electron Transistor (SET) or Quantum Point Contact (QPC) charge sensor biased at a fixed voltage. The qubit's spin state modulates the sensor's impedance, creating a small current variation. This direct readout measurement requires signal conditioning consisting of cryogenic low-noise transimpedance amplification, filtering, and integration before digitization by an analog-to-digital converter. The bandwidth and measurement speed are limited by capacitive parasitics in the wiring connecting the quantum device to the cryogenic amplifiers.

3. The Scaling Bottleneck and Cryogenic Solutions

Physical Constraints of Room-Temperature Control

Connecting thousands of individual physical qubits to room-temperature instrumentation creates a major hardware scaling bottleneck. Current quantum processor architectures require each qubit to have individual control lines, leading to a fundamental scaling limitation:

- **The Wiring Bottleneck:** The number of lines that can physically fit into a dilution refrigerator is fundamentally limited by the dilution fridge's volume and feedthrough constraints. Running independent coaxial lines for every single control and readout channel fills the available physical volume, making it impractical to scale to large qubit arrays.
- **Thermal Injection Load & System Complexity:** More connections cause the system to become more complex, expensive, and less reliable. The

heat conducted through each line from room-temperature stages (300 K) down to the colder cryogenic plates (4 K and 20 mK), together with the power dissipated in the various signal attenuators, adds to the thermal load of the dilution fridge, thus increasing requirements on its cooling power and operational costs.

- **Propagation Latency & Real-Time Control Limitations:** Long cables introduce significant signal propagation delays. A standard 1 meter coaxial cable results in a round-trip propagation time of approximately 10 ns, which is comparable to the duration of quantum operations (typically 10–100 ns). This latency limits the speed of active, real-time conditional error correction loops and feedback-based readout protocols, fundamentally constraining achievable gate fidelities.

Cryogenic Hardware Alternatives

To overcome these scaling limits, architectures must transition from individual, room-temperature instruments to localized cryogenic controllers. Two complementary strategies emerge:


- **Signal Multiplexing:** Using Frequency-Division Multiple Access (FDMA) or Time-Division Multiple Access (TDMA), a single physical wire can carry control and readout signals for multiple qubits, drastically reducing total cable count. The wiring from the 4 K stage to the qubit stage can be further reduced using electronics with limited complexity to implement multiplexing near the qubits.
- **Cryo-CMOS Integration:** Fabricating custom control circuits on silicon using standard CMOS processes allows the controller to operate directly inside the dilution refrigerator (typically at the 3 K or 4 K stage). This approach significantly reduces wiring complexity by moving the electronic interface closer to the qubits, shortening cable runs, eliminating bulk room-temperature instrumentation, and reducing the overall system footprint.

4. Deep Dive: Intel Horse Ridge Cryo-CMOS Controller

Intel’s **Horse Ridge** SoC is a custom cryogenic CMOS chip designed to sit at the 3 K stage of a dilution refrigerator. It acts as a wideband, frequency-multiplexed controller capable of driving both spin qubits and transmons.

Architecture Overview

Horse Ridge is a cryogenic controller implemented using frequency multiplexing and a shared interface to room-temperature equipment, conceived for both transmons and semiconductor quantum dots. The chip integrates a fully programmable digital back-end (FPGA fabric) with a wideband analog front-end modulator:



Digital Back-End Pulse Modulation

The digital core operates similarly to an FPGA, utilizing a hierarchical lookup table structure to schedule and synthesize microwave pulses. The overall memory organization comprises three linked hierarchical components:

1. **Command Loading & Instruction List:** An external computer compiles a high-level sequence of quantum gates (such as a Rabi sequence loop or two-qubit gate sequences) and transmits the list of “commands” to the chip’s core memory via USB or PCI, where they are stored sequentially.
2. **Instruction Scanning & Instruction Table:** The controller block manages interfaces with the external PC, the analog power amplifier, and other peripherals. It scans memory sequentially, executing commands one-by-one. Each command consists of a list of instructions to be executed. Each instruction has a reserved special place in a designated zone of memory: the **Instruction Table**, which stores a list of “things to do” (indexed as 0, 1, 2, 3, etc.). For each instruction, a pointer references another part in memory at a specific address range (e.g., “start 0, stop 999”).
3. **Envelope Memory & Parameter Retrieval:** The **Envelope Memory** is a list stored in the addressed memory zone, and each row contains numbers associated with amplitude, phase, pulse value, frequency point, etc. (e.g., 255, 0), represented within a minimum and maximum range (e.g., 0–255 for 8-bit resolution).
4. **Streaming Phase & Amplitude Vectors:** The digital back-end reads these phase and amplitude vectors out of the envelope memory in real time and streams the digital data directly to a **Digital-to-Analog Converter**

(DAC) array in the analog conversion stage.

Scripting an Automated Rabi Sequence Block

To execute an automated Rabi frequency characterization sweep (finding the Rabi frequency for a single qubit), the digital back-end executes a loop directly within its internal memory table, avoiding round-trip latency to the room-temperature master PC:

Example: Finding Rabi Frequency for a Single Qubit

Loop $i = 1 : \text{time-max}$

- Initialize (set qubit to ground state)
- Manipulate: send X-pulse for duration Time_i
- Read-out (measure final state)
- Increment i

This loop is implemented as:

Initialize Sequence Loop (Index $i = 1$ to max_time)
1. Read Instruction Table Address $\rightarrow$ Execute Qubit Initialization Pulse $\perp$ Fetch corresponding Envelope Memory segment (Address: Initialization)
2. Read Instruction Table Address $\rightarrow$ Fetch Envelope Segment $[\text{Address}_i]$ $\perp$ Stream X-Drive Pulse to DAC with precise Burst Duration (Time_i)
3. Read Instruction Table Address $\rightarrow$ Trigger Spin Readout Sequence $\perp$ Fetch corresponding Envelope Memory segment (Address: Readout)
4. Increment Index ($i = i + 1$) $\rightarrow$ Loop Until Max Timing Boundary is Reached

Each of the three main phases (Initialize, Manipulate, Read-out) corresponds to a configuration of a physical value (pulse timing, amplitude, frequency). The values are stored in a specific part of the memory called **Envelope Memory** at a specific address associated to the instruction, and are sent to the external analog modulation stages.

Analog Front-End Modulator and Package

- **Direct Waveform Synthesis:** The digital streams are sent to external stages that have as an initial block a multi-channel **Digital-to-Analog Converter (DAC)** array to generate raw baseband analog pulse envelopes. The time-domain and frequency-domain characteristics of these envelopes are carefully designed to minimize spectral leakage and sideband power.

- **Single Side-Band Modulation:** These analog envelopes drive an on-chip Single Side-Band (SSB) mixer and power amplifier engine at transistor level, modulating the pulses onto microwave carriers with minimal carrier leakage or image distortion. This provides efficient frequency upconversion for both transmon and semiconductor qubit control.
- **Advanced Multi-Core Die Realization:** The transmitter architecture is replicated across four independent on-chip pipelines (TX₁, TX₂, TX₃, TX₄) to support concurrent multi-qubit parallel drive routing.
- **Flip-Chip Packaging & Thermal Integration:** The silicon die is flip-chip bonded via Controlled Collapse Chip Connection (C₄) bumps onto a **324-pin Ball Grid Array (BGA)** package with impedance matching lines to the board and discrete decoupling capacitors for power supply noise reduction. This flip-chip configuration minimizes parasitic interconnect inductance relative to wire bonds, ensuring clean microwave transmission.
- **PCB and Thermal Shielding:** The BGA package is mounted onto a 6-layer Printed Circuit Board (PCB) with RF signal routing and enclosed into a gold-plated copper enclosure that acts as both an RF shield and a heat sink, connected directly to the 3 K plate of the dilution refrigerator. The board contains **40 DC lines (bias and supply), 10 high-speed digital interface lines, and 8 dedicated RF drive lines**. The chip is placed on the 3 K plate during operation, with an external FPGA serving as the master to synchronize the chip with other instruments used for qubit readout and initialization.

5. Comprehensive Modeling of a Semiconductor Spin-Qubit Processor

Performance Metrics

Horse Ridge demonstrates the feasibility of cryo-CMOS quantum controllers:

- **Wideband Frequency Multiplexing:** Supports simultaneous control of multiple qubits via FDMA across a broad RF bandwidth, reducing total wiring by an order of magnitude.
- **Integrated Multi-Channel DAC & Modulator:** Eliminates the need for external AWGs and VSGs for baseband and RF synthesis, reducing footprint and cost.
- **Reference:** J.V. Dijk et al. ‘A Scalable CryoCMOS Controller for the Wideband Frequency-Multiplexed Control of Spin Qubits and Transmons’. *IEEE Journal of Solid-State Circuits* 55.11 (Nov. 2020)

Generic Semiconductor Spin-Qubit Quantum Processor

A generic model for a semiconductor spin-qubit quantum processor comprises qubits encoded in the spin of electrons trapped in quantum dots and a charge

sensor (QPC or SET). The classical control electronics required for each line type includes:

- **Electron Spin Resonance (ESR) Lines:** For single-qubit rotations
- **Plunger Gate Lines:** For qubit potential tuning
- **Barrier Gate Lines:** For exchange coupling and two-qubit gates

Integrating these components yields a complete, modular semiconductor quantum processor architecture:

Core Node Routing Functions

- **DC Bias Sources:** Maintain the baseline potential of all barrier and plunger gates, stabilizing the electron confinement wells and fixing the static inter-dot tunnel barriers.
- **Resonant ESR Lines:** Deliver frequency-multiplexed microwave bursts generated by an AWG/LO modulation engine. If each qubit features a unique Larmor frequency (induced by a localized micromagnet gradient), a single physical ESR line can control multiple qubits independently using **Frequency-Division Multiple Access (FDMA)**, simplifying the system's wiring footprint.
- **Exchange Gate Controllers:** Dedicated fast AWGs supply voltage pulses to the barrier gates to modulate the inter-dot exchange coupling, executing two-qubit entangling operations (such as Swap or C-Phase gates).
- **Charge Sensor Processing Chain:** Measures current variations from the QPC or SET electrometer. The small analog signal is routed through a low-noise cryogenic transimpedance amplifier, filtered to remove out-of-band noise, integrated, and digitized to determine the final qubit state.


This guide offers an in-depth, structured explanation of the lecture notes from Professor Mariagrazia Graziano (Politecnico di Torino). It details the physics, architecture, and electronic control systems required to build and operate semiconductor-based quantum computers.

1. Qubit Encoding Physics: Spin vs. Charge

Semiconductor qubits are typically realized using **Quantum Dots (QDs)**—nano-scale semiconductor structures (often in Silicon or Gallium Arsenide) that trap individual electrons using electrostatic gates.

Spin Qubits

- **Encoding:** The logical states $|0\rangle$ and $|1\rangle$ map directly to the intrinsic magnetic dipole moment of a single electron trapped in a quantum dot: spin-up ($|\uparrow\rangle$) and spin-down ($|\downarrow\rangle$).



- **Coherence Time (T_2):** Relatively long ($\sim$ tens of μs in GaAs, extending to milliseconds or seconds in isotopically purified ^{28}Si). Spin states are well-isolated from environmental electric noise.
- **Manipulation Time:** Relatively slow ($\sim$ tens of ns), limited by the strength of oscillating magnetic fields or spin-orbit coupling available to drive transitions.

Charge Qubits

- **Encoding:** The qubit is encoded across a **Double Quantum Dot (DQD)** system. The state depends on the spatial position of a single electron:
- $|L\rangle$: Electron resides in the Left dot.
- $|R\rangle$: Electron resides in the Right dot.
- **Coherence Time (T_2):** Short ($\sim$ tens of ns). Because the qubit is based on charge position, it is highly susceptible to charge noise and voltage fluctuations in the surrounding substrate.
- **Manipulation Time:** Extremely fast ($\sim$ tens of ps). Transitions are driven using fast electric voltage pulses on electrostatic gates, which are much easier to generate dynamically than high-frequency magnetic fields.

Hybrid (Singlet-Triplet) Qubits

To balance these trade-offs, **hybrid qubits** use both position and spin states. For example, a **Singlet-Triplet (S - T_0) qubit** uses two electrons shared between two dots. The logical states are formed by the anti-symmetric Singlet state $|S\rangle$ and the symmetric, unpolarized Triplet state $|T_0\rangle$:

$$|S\rangle = \frac{1}{\sqrt{2}}(|\uparrow\downarrow\rangle - |\downarrow\uparrow\rangle)$$

$$|T_0\rangle = \frac{1}{\sqrt{2}}(|\uparrow\downarrow\rangle + |\downarrow\uparrow\rangle)$$

This configuration provides protection against global magnetic field fluctuations while allowing rapid, electrically controlled manipulation.

2. Mathematical Framework of the Charge Qubit

The behavior of an electron in a Double Quantum Dot (DQD) is governed by two main physical parameters:

1. **Detuning (ε):** The electrostatic potential energy difference between the left and right quantum dots ($\varepsilon = E_L - E_R$). This is directly modulated by applying voltages to the surface gates.
2. **Tunnel Coupling (t_0):** The quantum mechanical probability amplitude for the electron to tunnel through the potential barrier separating the two dots.

The Hamiltonian system in the localized position basis $\{|L\rangle, |R\rangle\}$ is written as:

$$H = \begin{pmatrix} \frac{\varepsilon}{2} & t_0 \\ t_0 & -\frac{\varepsilon}{2} \end{pmatrix}$$

Case 1: Zero Tunnel Coupling ($t_0 = 0$)

When the tunnel barrier is infinitely high, the states $|L\rangle$ and $|R\rangle$ are perfectly isolated eigenstates.

- The system energy eigenvalues plot as straight lines when mapped against detuning: $E = \pm \frac{\varepsilon}{2}$.
- At the degeneracy point ($\varepsilon = 0$), the energy levels cross. No quantum superposition can occur because the electron cannot cross the barrier.

Case 2: Finite Tunnel Coupling ($t_0 > 0$)

When the barrier is lowered, quantum tunneling mixes the localized states. Diagonalizing the Hamiltonian yields the hybridized energy eigenvalues:

$$E = \pm \sqrt{\left(\frac{\varepsilon}{2}\right)^2 + t_0^2}$$

- **Avoided Crossing:** At the zero-detuning point ($\varepsilon = 0$), the energy levels no longer cross. Instead, they split by a minimum energy gap equal to $2t_0$.
- **Eigenstates at $\varepsilon = 0$:** The new ground and excited states form symmetric and anti-symmetric molecular-like superpositions:
- Ground state: $|+\rangle = \frac{1}{\sqrt{2}}(|L\rangle + |R\rangle)$
- Excited state: $|-\rangle = \frac{1}{\sqrt{2}}(|L\rangle - |R\rangle)$

3. Two-Qubit Coupling Mechanism & Exchange Interaction (J)

When two spin qubits (Q0 and Q1) are placed in adjacent quantum dots, their physical proximity causes their wavefunctions to overlap, giving rise to the **Exchange Interaction (J)**.

Physical Principle

According to the Pauli Exclusion Principle, two electrons cannot occupy the exact same quantum state. If two adjacent electrons have **parallel spins** ($|\uparrow\uparrow\rangle$ or $|\downarrow\downarrow\rangle$), Coulomb repulsion forces them apart because spatial overlap is restricted. If they have **anti-parallel spins** ($|\uparrow\downarrow\rangle$ or $|\downarrow\uparrow\rangle$), they can tunnel into a shared orbital state (virtual tunneling).

This energy shift means the energy required to flip the spin of Q0 depends entirely on whether the spin of Q1 is pointing up or down.

Frequency Shifts

This state-dependent energy shift modifies the Electron Spin Resonance (ESR) frequencies. If you monitor the resonance frequency of Q_0 , you observe two distinct lines:

- f_1 : The resonance frequency of Q_0 when Q_1 is in state $|\downarrow\rangle$.
- f_2 : The resonance frequency of Q_0 when Q_1 is in state $|\uparrow\rangle$.

The Exchange Interaction energy parameter J quantified in terms of frequency is:

$$J = f_2 - f_1$$

By tuning the electrostatic gates (modifying detuning ε or lowering the inter-dot tunnel barrier), you turn J on and off, enabling controlled two-qubit logic operations.

4. Hardware Implementation: The Veldhorst Architecture

The slides reference a foundational 2015 experiment by Menno Veldhorst et al. (Nature), which demonstrated a two-qubit logic device in Silicon.

Device Control Topology

The device controls a two-dimensional electron gas using surface metallic gates (G_1, G_2, G_3) and a central coupling gate (G_c):

- **Plunger Gates (G_1, G_2):** Control the local chemical potential (detuning ε) of Qubit 1 and Qubit 2 respectively.
- **Coupling Gate (G_c):** Controls the electrostatic barrier between the dots, dynamically turning the exchange interaction J on or off.

Charge Stability & The “Diamond Diagram”

To initialize the system, operators use a **Charge Stability Diagram** (often called a honeycomb or diamond diagram). This plot maps out the stable electron charge configurations (N_1, N_2)—where N_1 and N_2 represent the exact number of electrons in Dot 1 and Dot 2—as a function of the gate voltages.

By precisely biasing the gates, the system is calibrated to the (1,1) regime, ensuring exactly one active electron occupies each dot.

Readout via Single-Electron Transistor (SET)

Readout uses a highly sensitive nearby **Single-Electron Transistor (SET)** coupled capacitively to the dots.

- **Spin-to-Charge Conversion:** Because an SET can only detect changes in local electrical charge (not magnetic spin), spin must be converted to charge position first.
- **Sequential Readout:**
- **Q2 Readout:** Gate voltages are pulsed toward the $(1, 1) \rightarrow (0, 2)$ charge transition boundary. If the spins form a singlet state, tunneling is allowed; if they form a triplet state, Pauli block prevents tunneling. The SET senses whether an electron moved.
- **Q1 Readout:** The system is shifted toward the $(0, 1) \rightarrow (0, 0)$ boundary, reading out the final electron state.

5. Timeline of a CNOT Gate Operation

The Controlled-NOT (CNOT) gate is the fundamental building block for universal quantum computation. The experimental process maps to the following timeline:

Phase 1: Prep	Phase 2: Init Rotation	Phase 3: Entangling	Phase 4: Conditional	Phase 5:
Align Q_0 to Reservoir (Elzerman)	Apply $\pi/2$ MW Pulse to Target Qubit (Q_0)	Shift Detuning to Induce Hybridization	Q_0 Rotates at distinct Rabi Rates based on Q_1	Apply 2π to control qubit

Phase 1: Initialization & Spin Alignment

- The target qubit (Q_0 , Blue) is initialized into a known pure spin state using the **Elzerman Method**.
- The quantum dot energy levels are aligned relative to the Fermi energy (E_F) of a nearby electron reservoir. A spin-down electron is energetically favored to stay in the dot, whereas a spin-up electron tunnels out into the reservoir and is instantly replaced by a spin-down electron. This initializes the system to $|\downarrow\rangle$.

Phase 2: Initial Target Rotation

- At this stage, the two qubits are kept decoupled ($J \approx 0$).
- An oscillating microwave (MW) pulse matching the native resonance frequency of Q_0 is applied. This drives a $\pi/2$ rotation (a square-root of NOT operation), placing the blue target qubit into a coherent superposition state:

$$\frac{1}{\sqrt{2}}(|\downarrow\rangle + |\uparrow\rangle)$$

Phase 3: Controlled Hybridization (The Entangling Step)

- The control electronics apply a rapid voltage pulse to alter the detuning ε , shifting the system's energy levels.
- This brings the state of the blue qubit (Q_0) into close resonance with a higher orbital level of the red control qubit (Q_1). The wavefunctions mix, turning on the exchange interaction ($J > 0$). The two qubits are now highly entangled.

Phase 4: Conditional Rabi Oscillations

- Because J is active, the resonance frequency of Q_0 splits into two paths depending on Q_1 's state:
- If $Q_1 = |\downarrow\rangle$, Q_0 precesses at frequency f_1 .
- If $Q_1 = |\uparrow\rangle$, Q_0 precesses at frequency f_2 .
- As time progresses, the probability wave of Q_0 oscillates (Rabi oscillations). The phase accumulation rate is noticeably different for both conditions.
- **The Critical Intercept ($\sim 0.5 \mu\text{s}$):** The system dwells in this interaction state until exactly $\Delta t = \frac{1}{2J}$ has elapsed. At this precise timestamp, the phase difference between the two configurations reaches exactly π radians (180°).

Phase 5: Target Finalization

- The detuning pulse is turned off, dropping J back to zero and stopping the conditional evolution.
- A final $\pi/2$ microwave pulse is applied to Q_0 .
- If Q_1 was $|\downarrow\rangle$, the total phase accumulated causes the final pulse to return Q_0 safely to its original state.
- If Q_1 was $|\uparrow\rangle$, the extra π phase shift causes the final pulse to flip Q_0 to the opposite spin state.

This sequence completes the logical target flip conditioned entirely on the control state, completing the **CNOT** gate operation.

Semiconductor Qubits: Materials, Fabrication, and Heterostructure Physics

1. Physical Encoding Across Quantum Dot Modalities

Semiconductor quantum dots can isolate discrete charges or spins to define a two-level quantum system (qubit). Depending on the physical architecture and

the number of confined electrons, quantum information can be encoded using different degrees of freedom.

Classification of Quantum Dots and Qubits

The mapping of the physical computational space varies depending on the hardware configuration:

Structure	Qubit Type	# Electrons	Physical Encoding Basis
Single Quantum Dot (SQD)	Spin Qubit	1	Electron spin orientation ($ \uparrow\rangle, \downarrow\rangle$)
Single Quantum Dot (SQD)	Charge Qubit	1	Discrete orbital energy levels
Double Quantum Dot (DQD)	Charge Qubit	1	Spatial electron position (Left $ 1\rangle$ vs. Right $ 2\rangle$)
Double Quantum Dot (DQD)	Spin Qubit	2	Two-electron spin states (Singlet $ S\rangle$, Triplet $ T\rangle$)

2. Strategic Motivations for the Semiconductor Platform

Solid-state semiconductor quantum dots are an attractive platform for building large-scale quantum processors due to several core advantages:

- **Leveraging Industrial Infrastructure:** Spin qubits utilize manufacturing techniques pioneered by the microelectronics industry. The ability to fabricate devices using advanced lithography, chemical vapor deposition, and etching standardizes yield and precision.
- **Classical-Quantum Monolithic Integration:** A primary bottleneck in quantum computing scaling is the routing of discrete high-frequency lines from room temperature down to sub-Kelvin environments. Because semiconductor quantum dots are structurally similar to classical Field-Effect Transistors (FETs), they can be co-integrated with cryogenic control and readout circuitry (such as Cryo-CMOS multiplexers, amplifiers, and DACs) on the same silicon die or package.
- **High Real Estate Scaling Density:** Compared to alternative modalities like superconducting qubits—which span hundreds of micrometers due to planar microwave resonators—semiconductor quantum dots have pitch requirements on the order of a few hundred nanometers. This small footprint allows millions of physical qubits to fit onto a single square centimeter of substrate.

3. High-Purity Atomic Layer Deposition: Molecular Beam Epitaxy (MBE)

To achieve the long electron mean-free paths needed for spin coherence, heterostructures require sharp atomic-layer transitions with extremely low back-

ground defect densities. This purity is achieved using **Molecular Beam Epitaxy (MBE)**.

Ultra-High Vacuum (UHV) Chamber Environment

The core growth process takes place inside a stainless steel chamber maintained at ultra-high vacuum pressures ($\sim 10^{-11}$ mbar). This extreme vacuum is sustained by continuous liquid nitrogen cryoshrouds, which condense residual gas molecules on the chamber walls to prevent impurities from incorporating into the crystal matrix.

Effusion Cells and Shutter Mechanics

Elements like Gallium (Ga), Arsenic (As), and Aluminum (Al) are housed in individual pyrolytic boron nitride crucibles within thermal **effusion cells**.

- **Sublimation:** Heating elements wound around the crucibles heat the source materials until they sublime, producing a directed, low-pressure beam of atoms or molecules traveling at thermal velocities toward the substrate.
- **Shutter Control:** Pneumatic mechanical shutters positioned in front of each cell open and close in fractions of a second. This quick action allows for monolayer-precise modulation of the incoming atomic flux, creating atomically sharp material transitions at the heterostructure interfaces.
- **In-Situ Metrology:** Growth rates and surface crystal structures are monitored in real-time using **Reflection High-Energy Electron Diffraction (RHEED)**. An electron gun fires a high-energy beam at a grazing angle relative to the substrate surface, generating a diffraction pattern on a fluorescent screen that tracks surface flatness and layer-by-layer accumulation. A quadrupole mass analyzer monitors the chemical purity of the ambient flux.

4. Physics of the AlGaAs/GaAs Heterostructure and 2DEG Formation

Layering materials with differing electronic properties creates a highly confined, two-dimensional channel of high-mobility electrons at the interface.

Heterostructure Composition and Energy Band Diagrams

An $\text{Al}_x\text{Ga}_{1-x}\text{As}$ alloy layer is grown epitaxially on top of a pure GaAs substrate. The bandgap (E_g) and electron affinity (χ) of the ternary alloy depend directly on the Aluminum mole fraction x :

$$E_g(x) = 1.422 \text{ eV} + 1.2475x \text{ eV}$$

$$\chi(x) = 4.07 - 1.1x \text{ eV}$$

For a standard composition of $x = 0.3$:

$$E_g(\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}) = 1.79 \text{ eV} \quad (\text{versus } 1.422 \text{ eV for GaAs})$$

$$\chi(\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}) = 3.74 \text{ eV} \quad (\text{versus } 4.07 \text{ eV for GaAs})$$

When these two materials are brought into contact, their Fermi levels (E_F) align at thermal equilibrium. Because GaAs has a larger electron affinity and a smaller bandgap, the conduction band edge (E_c) bends sharply at the interface. This bending forms a deep, triangular quantum well on the GaAs side of the junction.

Modulation Doping and Electrostatic Confinement

To populate this triangular well with carriers without introducing scattering defects, **modulation doping** is used:

1. Silicon (Si) donor impurities are embedded within the AlGaAs layer, separated from the interface by an undoped AlGaAs spacer layer.
2. The ionized electrons leave their parent Silicon donors in the wider-bandgap AlGaAs layer and drop into the lower energy states of the triangular quantum well on the GaAs side.
3. Because the mobile electrons are physically separated from the ionized Si^+ impurities by the spacer layer, remote impurity scattering is reduced. This confinement traps the carriers at the interface—typically buried roughly 50 nm beneath the physical wafer surface—forming a high-mobility **Two-Dimensional Electron Gas (2DEG)**.

Gate Deconstruction into Quantum Dots

To isolate a single electron from this continuous 2DEG, nanoscale metallic **Schottky top-gates** are patterned on the surface of the heterostructure using electron-beam lithography. Applying a negative voltage ($V_g < 0$) to these gates raises the local conduction band energy beneath them. This repels the underlying electrons and selectively depletes portions of the 2DEG, isolating a small electrostatic puddle of charge (the quantum dot) with discrete energy levels.

5. The Dephasing Bottleneck: Hyperfine Interaction

While AlGaAs/GaAs heterostructures provide high electron mobility and well-understood physics, they face a severe quantum dephasing bottleneck due to magnetic interactions with the host crystal lattice.

The Hyperfine Coupling Mechanism

The localized electron wavefunction confined within a GaAs quantum dot physically overlaps with approximately 10^4 to 10^6 host atomic nuclei. In these materials, every native isotope (^{69}Ga , ^{71}Ga , and ^{75}As) possesses a non-zero nuclear spin ($I = 3/2$).

This spatial overlap triggers a Fermi contact **hyperfine interaction**, which couples the electron's spin vector $\vec{S}$ to the surrounding collective nuclear spin bath $\vec{I}_k$:

$$H_{\text{hyperfine}} = A \sum_k |\psi(\vec{r}_k)|^2 \vec{S} \cdot \vec{I}_k$$

where A represents the atomic hyperfine coupling constant and $|\psi(\vec{r}_k)|^2$ is the electron probability density at the site of the k -th nucleus.

Overhauser Field Fluctuations and Dephasing

This collective coupling can be modeled as a local, quasi-static magnetic field called the **Overhauser Field** ($\vec{B}_N$) acting directly on the electron spin.

Because the nuclear spins fluctuate and flip randomly via slow thermal dipole-dipole interactions, the magnitude and direction of $\vec{B}_N$ drift stochastically over time. This fluctuating field shifts the electron's Zeeman splitting energy ($\Delta E_Z = g\mu_B|\vec{B}_0 + \vec{B}_N|$), causing the qubit's Larmor precession frequency to drift.

This phase randomization quickly dephases the qubit, limiting the uncorrected ensemble coherence time (T_2^*) in native GaAs devices to approximately 10 ns. To overcome this limit, modern architectures use alternative materials like Silicon-Silicon Germanium (SiGe/Si) or Silicon-on-Insulator (SOI) systems, where isotopic purification can eliminate nuclear spins to provide a magnetically quiet environment.

Energy Band Engineering of the $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}/\text{i-GaAs}$ Heterojunction

1. Isolated Material Parameters before Contact ($t = 0$)

To understand the interface physics of an $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}/\text{i-GaAs}$ **heterojunction** (commonly used to build high-electron-mobility transistors and semiconductor quantum dots), we evaluate the two bulk materials relative to a shared reference, the **Vacuum Level** (E_0).

Material Properties & Mathematical Models

The electronic properties of the ternary compound alloy $\text{Al}_x\text{Ga}_{1-x}\text{As}$ are determined by its Aluminum mole fraction x .


For an alloy composition of $x = 0.3$, the energy bandgap (E_g) and the electron affinity ($q\chi$) are modeled by:

$$E_g(x) = 1.42 \text{ eV} + 1.2475 \cdot x \implies E_g(0.3) = 1.79 \text{ eV}$$

$$q\chi(x) = 4.07 \text{ eV} - 1.1 \cdot x \implies q\chi(0.3) = 3.74 \text{ eV}$$

Bulk Band Alignment Comparison

- **Modulation-Doped Left Layer (m-Al_{0.3}Ga_{0.7}As):** Features a wide bandgap ($E_{g1} = 1.79 \text{ eV}$) and a smaller electron affinity ($q\chi_1 = 3.74 \text{ eV}$). It is doped n -type using Silicon (Si donor ions, $\oplus$), driving its Fermi Level (E_F) close to or above the conduction band edge (E_{C1}).
- **Intrinsic Right Layer (i-GaAs):** Features a narrower bandgap ($E_{g2} = 1.42 \text{ eV}$) and a larger electron affinity ($q\chi_2 = 4.07 \text{ eV}$). Being intrinsic, its Fermi level sits precisely at its intrinsic midpoint ($E_{F2} = E_{Fi}$).



2. Band Discontinuities at the Interface: The Electron Affinity Rule

When the materials form an interface, the difference in atomic composition causes abrupt steps in the conduction and valence bands. These steps are called **band offsets** or **discontinuities**.

According to **Anderson's Electron Affinity Rule**, the conduction band discontinuity (ΔE_C) is determined solely by the difference in electron affinities between the two materials:

$$\Delta E_C = q\chi_1 - q\chi_2 = 3.74 \text{ eV} - 4.07 \text{ eV} = -330 \text{ meV}$$

This negative value indicates an abrupt drop in conduction energy when crossing from the AlGaAs layer into the GaAs substrate.

To calculate the matching valence band discontinuity (ΔE_V), we evaluate the total difference in bandgaps (ΔE_g):

$$\Delta E_g = E_{g1} - E_{g2} = 1.79 \text{ eV} - 1.42 \text{ eV} = +370 \text{ meV}$$

By matching the physical boundaries across the material interface:

$$E_{C1} + \Delta E_C = E_{C2}$$

$$E_{V1} + \Delta E_V = E_{V2}$$

Subtracting these equations yields the relationship between band discontinuities and the bandgap difference:

$$(E_{C1} - E_{V1}) + \Delta E_C - \Delta E_V = E_{C2} - E_{V2}$$

$$E_{g1} + \Delta E_C - \Delta E_V = E_{g2} \implies \Delta E_V = \Delta E_g + \Delta E_C$$


Substituting the numerical parameters gives:

$$\Delta E_V = 370 \text{ meV} + (-330 \text{ meV}) = +40 \text{ meV}$$

The positive step confirms that the valence band steps upward when crossing from the wide-bandgap layer into the narrow-bandgap layer.

3. Electrostatic Equilibrium, Band Bending, and Poisson's Equation

At $t > 0$, contact is established. Because the Fermi levels of the two isolated materials are misaligned ($E_{F1} > E_{F2}$), electrons flow from the n -doped AlGaAs layer across the interface into the empty lower energy states of the intrinsic GaAs layer.



This migration breaks local charge neutrality and sets up a built-in electric field. At **thermal equilibrium**, the net current flows stop and the Fermi Level (E_F) becomes uniform across the entire system.

Electrostatic Modeling via Poisson's Equation

The spatial variation of the electrostatic potential (Φ) and the resulting bending of the energy bands are governed by **Poisson's Equation**:

$$\frac{\partial^2 \Phi}{\partial x^2} = -\frac{\rho(x)}{\varepsilon_s}$$

where $\rho(x)$ is the net volumetric charge density and ε_s is the dielectric permittivity of the semiconductor substrate.

- **On the Left (m-AlGaAs, $x < 0$):** Moving electrons leaves behind fixed, ionized dopant atoms (Si^+). This forms a space-charge depletion region with a uniform positive charge density:

$$\rho(x) = +qN_D$$

- **On the Right (i-GaAs, $x > 0$):** The injected electrons collect at the interface. This creates a high-density mobile negative charge profile that decays exponentially as it moves deeper into the semiconductor:

$$\rho(x) = -qN_0 \cdot e^{-\frac{x}{L_D}}$$

The spatial reach of this shielding screening cloud is defined by the **Intrinsic Debye Length** (L_{Di}):

$$L_{Di} = \sqrt{\frac{k_B \epsilon_s T}{2n_i q^2}}$$

The Resulting Equilibrium Band Diagram

This charge distribution distorts the electrostatic potential profile across the interface, bending the energy bands into their final shape at thermal equilibrium:

- **On the AlGaAs Side:** The bands bend upward as they approach the junction, reflecting the depletion of electrons and the presence of positive space charge.
- **On the GaAs Side:** The bands bend downward as they approach the junction due to the accumulation of electrons.
- **The Triangular Quantum Well and 2DEG:** This combination of downward band bending and the abrupt $\Delta E_C = -330$ meV conduction band step forms a **triangular quantum well** at the interface on the GaAs side. Because this well drops significantly below the flat Fermi level (E_F), electrons pool inside it. This confined layer of electrons forms a high-mobility **Two-Dimensional Electron Gas (2DEG)**, which provides the starting platform for isolating single electron spins in quantum dots.


4. Alternate Model Formulations: Metal-Semiconductor Analogies

The electrostatic band bending behavior at the heterojunction interface can be modeled using configurations that resemble metal-semiconductor junctions.

Case A: Degenerate n -type to p -type Homojunction

Consider an interface between highly doped, degenerate n -type GaAs (where E_F sits above E_C , mimicking a metal source) and standard p -type GaAs.

The difference in their work functions ($q\Phi_{sm}$ and $q\Phi_{sp}$) establishes a built-in potential barrier. When brought into contact, electrons migrate from the degenerate n -type side into the p -type side, causing the bands to bend upward near the interface until the Fermi levels align.



/

 INTERFACE
Case B: Metal-Like Degenerate Layer to Intrinsic Substrate

Alternatively, consider a highly doped degenerate layer in contact with an intrinsic semiconductor (*i*-GaAs). Because the intrinsic layer's initial Fermi level sits at midgap (E_{Fi}), a large energy difference exists relative to the degenerate side.

Upon contact, electrons pour into the intrinsic substrate. This accumulation creates a localized downward band bending profile near the interface, identical to the accumulation dynamics that form a 2DEG channel. # Exercises: Heterojunction Band Engineering and Electrostatic Modeling ### Exercise 1: Boundary Conditions and Energy Parameters of the Isolated Heterojunction

Context: Before a heterojunction is brought into contact, its individual material parameters must be calculated relative to the vacuum level E_0 . You are analyzing a modulation-doped $m\text{-Al}_{0.3}\text{Ga}_{0.7}\text{As}$ / *i*-GaAs heterostructure engineered for spin-qubit applications.

Given Data: * Aluminum mole fraction: $x = 0.3$

- Empirical bandgap formula for $\text{Al}_x\text{Ga}_{1-x}\text{As}$: $E_g(x) = 1.42 \text{ eV} + 1.2475 \cdot x$
- Empirical electron affinity formula for $\text{Al}_x\text{Ga}_{1-x}\text{As}$: $q\chi(x) = 4.07 \text{ eV} - 1.1 \cdot x$
- Properties of intrinsic Gallium Arsenide (*i*-GaAs): $E_{g2} = 1.42 \text{ eV}$, $q\chi_2 = 4.07 \text{ eV}$

Questions: 1. **Calculate** the bulk energy bandgap (E_{g1}) and the electron affinity ($q\chi_1$) for the $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}$ layer at $t = 0^+$.

2. **Sketch** the qualitative energy band diagram for both materials *before contact*, aligning them to a shared flat reference line representing the Vacuum Level (E_0). Label E_{C1} , E_{V1} , E_{F1} , E_{C2} , E_{V2} , and E_{F2} (E_{Fi}).

Exercise 2: Deriving Interface Band Discontinuities (The Affinity Rule)

Context: When the interface between $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}$ and *i*-GaAs is formed, the sudden change in material composition yields sharp energy steps at the conduction and valence band edges.

Questions: 1. Using **Anderson's Electron Affinity Rule**, calculate the conduction band discontinuity (ΔE_C) across the interface. State clearly whether

the energy steps *down* or *up* when moving from the wide-bandgap material to the narrow-bandgap material.

2. Prove analytically that the valence band discontinuity is given by $\Delta E_V = \Delta E_g + \Delta E_C$.
3. Calculate the numerical value of ΔE_V for this system. Does the valence band edge step up or down at the junction crossing?

Exercise 3: Charge Density and Electrostatic Potential (Poisson Modeling)

Context: At thermal equilibrium ($t \rightarrow \infty$), the Fermi level (E_F) collapses into a single, flat horizontal line across the entire structure. This creates a net charge distribution $\rho(x)$ that induces band bending near the junction ($x = 0$).

Questions: 1. Write down the fundamental differential equation (**Poisson's Equation**) that relates the spatial variation of the electrostatic potential (Φ) to the local volumetric charge density ($\rho(x)$).

2. The m -AlGaAs layer ($x < 0$) undergoes carrier depletion, leaving behind fixed ionized donor impurities (Si^+). State the mathematical expression for $\rho(x)$ in this region.
3. The intrinsic GaAs substrate ($x > 0$) accumulates free carriers at the interface. This accumulation decays exponentially deeper into the bulk as:

$$\rho(x) = -qN_0 \cdot e^{-\frac{x}{L_{Di}}}$$

Define the physical meaning of the parameter L_{Di} (**Intrinsic Debye Length**) and state how it limits the spatial reach of the screening charge cloud.

Exercise 4: Comprehensive Band Engineering and 2DEG Analysis

Context: The combined action of the conduction band offset (ΔE_C) and the electrostatic band bending creates the physical architecture needed to hold a semiconductor spin qubit.

Questions: 1. **Sketch** the complete, detailed energy band diagram of the $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As} / i\text{-GaAs}$ heterojunction at **thermal equilibrium**. Ensure your sketch accurately portrays: * A completely flat, continuous Fermi Level (E_F). * An upward band bending profile on the left ($x < 0$) indicating depletion. * The sharp offset step ($\Delta E_C = -330 \text{ meV}$) at $x = 0$. * A downward band bending profile on the right ($x > 0$) indicating accumulation. 2. **Identify and highlight** the geometric region on your sketch where a **Two-Dimensional Electron Gas**

(2DEG) forms. Explain what energy boundary conditions trap the electrons inside this region.

Exercise 5: Comparative Analysis of Metal-Semiconductor Analogs

Context: The electrostatic profiles of complex heterojunctions can often be modeled by comparing them to equivalent homojunction scenarios or metal-semiconductor junctions.

Questions: 1. **Case A:** Imagine a homojunction composed of a highly degenerate n -type GaAs layer joined to a standard p -type GaAs layer. Sketch its equilibrium band diagram. Why does the degenerate n -type side behave electrostatically like a metal contact?

2. **Case B:** Now consider a degenerate metal-like semiconductor layer in direct contact with an intrinsic semiconductor (i -GaAs). Show how the matching charge migration forces the conduction band edge (E_C) to dip below the Fermi level at the interface. Connect this behavior directly to the triangular quantum well formation observed in multi-material heterostructures.

Here are the solutions and detailed explanations for each of the five exercises.

Solutions & Detailed Explanations

Solution to Exercise 1: Boundary Conditions and Energy Parameters

1. Numerical Calculations To find the bulk energy properties of the isolated $\text{Al}_{0.3}\text{Ga}_{0.7}\text{As}$ layer at $t = 0^+$, substitute $x = 0.3$ into the given empirical formulas:

- **Energy Bandgap (E_{g1}):**

$$E_{g1} = 1.42 \text{ eV} + 1.2475 \cdot (0.3) = 1.42 \text{ eV} + 0.37425 \text{ eV} = 1.79425 \text{ eV} \approx 1.79 \text{ eV}$$

- **Electron Affinity ($q\chi_1$):**

$$q\chi_1 = 4.07 \text{ eV} - 1.1 \cdot (0.3) = 4.07 \text{ eV} - 0.33 \text{ eV} = 3.74 \text{ eV}$$

2. Qualitative Energy Band Diagram (Before Contact) Before contact, the two regions are separate systems with unaligned Fermi levels. They are drawn side-by-side using the common Vacuum Level (E_0) as a flat horizontal reference line:

- **Explanation of Alignment:** Because $q\chi_2 > q\chi_1$, the conduction band edge of GaAs (E_{C2}) sits lower in energy than that of AlGaAs (E_{C1}).

n-Al_{0.3}Ga_{0.7}As (Region 1)	vs	i-GaAs (Region 2)
Vacuum Level (E_0)	→	Vacuum Level (E_0)
Electron Affinity ($q\chi_1 = 3.74$ eV)		Electron Affinity ($q\chi_2 = 4.07$ eV)
Conduction Band (E_{C1})		Conduction Band (E_{C2})
Fermi Level (E_{F1} n-type)		Midgap Fermi Level ($E_{F2} = E_{Fi}$)
Bandgap ($E_{g1} = 1.79$ eV)		Bandgap ($E_{g2} = 1.42$ eV)
Valence Band (E_{V1})		Valence Band (E_{V2})

Figure 1: Energy band alignment properties prior to contact formation.

Because the left side is *n*-type modulated, its Fermi level (E_{F1}) is close to E_{C1} . The right side is intrinsic, meaning E_{F2} lies exactly at the midgap position (E_{Fi}).

Solution to Exercise 2: Interface Band Discontinuities

1. Conduction Band Discontinuity (ΔE_C) Using Anderson's Electron Affinity Rule, the energy step at the conduction band edge is given by:

$$\Delta E_C = q\chi_1 - q\chi_2 = 3.74 \text{ eV} - 4.07 \text{ eV} = -0.33 \text{ eV} = -330 \text{ meV}$$

- **Direction:** The negative sign dictates that the conduction band **steps down** by 330 meV at the interface boundary when moving from the wide-bandgap AlGaAs to the narrow-bandgap GaAs.

2. Analytical Proof for ΔE_V At the exact spatial coordinate of the junction interface, the physical energy boundaries must balance. We define the step transitions by setting:

1. $E_{C2} = E_{C1} + \Delta E_C$
2. $E_{V2} = E_{V1} + \Delta E_V$

Subtracting Equation (2) from Equation (1) yields:

$$(E_{C2} - E_{V2}) = (E_{C1} - E_{V1}) + \Delta E_C - \Delta E_V$$

By definition, the difference between the conduction and valence band edges equals the bandgap of that respective material ($E_C - E_V = E_g$). Substituting E_{g2} and E_{g1} into the expression:

$$E_{g2} = E_{g1} + \Delta E_C - \Delta E_V$$

Rearranging the terms to isolate the valence band offset (ΔE_V) gives:

$$\Delta E_V = (E_{g1} - E_{g2}) + \Delta E_C$$

$$\Delta E_V = \Delta E_g + \Delta E_C \quad \blacksquare$$

3. Numerical Calculation of ΔE_V First, find the total bandgap difference (ΔE_g):

$$\Delta E_g = E_{g1} - E_{g2} = 1.79 \text{ eV} - 1.42 \text{ eV} = +0.37 \text{ eV} = +370 \text{ meV}$$

Now, substitute ΔE_g and ΔE_C into our proven formula:

$$\Delta E_V = 370 \text{ meV} + (-330 \text{ meV}) = +40 \text{ meV}$$

- **Direction:** The positive sign indicates that the valence band edge **steps up** by 40 meV when crossing the junction from AlGaAs into GaAs.

Solution to Exercise 3: Charge Density and Poisson Modeling

1. Poisson's Equation The core electrostatic behavior governing band bending is expressed by the one-dimensional form of Poisson's Equation:

$$\frac{\partial^2 \Phi}{\partial x^2} = -\frac{\rho(x)}{\varepsilon_s}$$

2. Charge Density $\rho(x)$ in the Depletion Region ($x < 0$) As electrons migrate across the interface into the substrate, they leave behind uncompensated, static, positively charged dopant impurities. Assuming a sharp depletion profile, the net volumetric charge density is uniform and positive:

$$\rho(x) = +qN_D \quad (\text{for } -W_d < x < 0)$$

3. Intrinsic Debye Length (L_{Di}) Evaluation The parameter L_{Di} is defined mathematically as:

$$L_{Di} = \sqrt{\frac{k_B \varepsilon_s T}{2n_i q^2}}$$

- **Physical Meaning & Limits:** The Intrinsic Debye Length represents the characteristic screening distance in a semiconductor. It quantifies how far an uncompensated electrostatic perturbation or charge sheet can penetrate into the bulk material before the mobile carriers rearrange to screen it out. Because $\rho(x) \propto e^{-x/L_{Di}}$, the accumulation charge density drops significantly within a few multiples of L_{Di} , effectively restricting the excess electron cloud to a narrow channel right at the interface.

Solution to Exercise 4: Comprehensive Band Engineering and 2DEG Analysis

1. Equilibrium Energy Band Diagram Sketch At thermal equilibrium, the Fermi level (E_F) is completely flat and continuous across the entire junction. The vacuum level (E_0), conduction band (E_C), and valence band (E_V) bend in response to the internal electric field generated by the charge transfer:

Energy Band Element	n-AlGaAs Layer ($x < 0$)	i-GaAs Layer ($x > 0$)
Conduction Band Edge (E_C)	Starts at E_{C1} , bends down to 0.8 eV	Starts at 1.8 eV (well), bends up to
Interface Discontinuity	$\Delta E_C = -330$ meV (Drop)	$\Delta E_C = -330$ meV (Drop)
Quantum Well State	No confinement zone present	Triangular Well forms 2DEG at int
Fermi Level State (E_F)	Flat equilibrium line (2.5 eV)	Flat equilibrium line (2.5 eV)
Bending Potential (V_n)	Negligible baseline reference	Active potential upward shift toward
Valence Band Edge (E_V)	Starts at E_{V1} , bends down to -1.5 eV	Starts at -1.1 eV, bends up to E_{V2}
Interface Discontinuity	$\Delta E_V = +40$ meV (Jump)	$\Delta E_V = +40$ meV (Jump)

Figure 2: Side-by-side regional band diagram alignment profile.

2. Identifying the Two-Dimensional Electron Gas (2DEG)

- **Location:** The 2DEG forms on the *i*-GaAs side of the interface ($x > 0$), inside the narrow triangular notch highlighted in the diagram where the bent conduction band edge (E_{C2}) dips below the flat equilibrium Fermi level (E_F).
- **Energy Boundary Conditions:** Electrons are trapped inside this region by two distinct physical barriers:
 1. **Left Boundary ($x = 0$):** The large, abrupt quantum mechanical energy step of $\Delta E_C = 330$ meV prevents electrons from spilling backward into the AlGaAs layer.
 2. **Right Boundary ($x > 0$):** The electrostatic slope of the downward-bent conduction band acts as a barrier that prevents electrons from escaping

deeper into the bulk GaAs substrate. This tight quantum confinement restricts their motion to a two-dimensional plane parallel to the interface, creating the stable 2DEG channel used to isolate spin qubits.

Solution to Exercise 5: Comparative Analysis of Metal-Semiconductor Analogs

1. Case A: Degenerate n -type to p -type Homojunction When a semiconductor is degenerately n -doped, its active donor concentration is so high that the Fermi level (E_F) is forced up past the conduction band edge (E_C), populating the bottom of the conduction band with a high density of free electrons.

- **Metal-Like Electrostatic Behavior:** Because the sea of free electrons in a degenerate semiconductor is highly dense, it can shift charge rapidly with minimal changes to its internal energy bands. Like a bulk metal contact, its screening length is extremely short, causing it to act as a low-resistance source or drain that injects carriers into an adjacent channel while keeping its own potential stable.

2. Case B: Metal-Like Degenerate Layer to Intrinsic Substrate When this metal-like degenerate layer makes contact with an intrinsic substrate (i -GaAs), the initial energy alignment is highly unbalanced because the intrinsic substrate's Fermi level sits deep at midgap (E_{Fi}).

Energy Band Element	Degenerate Contact Layer ($x < 0$)	Intrinsic Substrate Layer
Conduction Band (E_C)	Flat horizontal line positioned at 2.0	Curves smoothly downward from interface
Fermi Level (E_F)	Aligned exactly with the Conduction Band edge	Stays flat at 2.0, deep inside valence band
Band Behavior Profile	Constant profile (No bending/degenerate)	Dynamic profile (Downward bending)

Figure 3: Side-by-side regional band diagram alignment profile for the degenerate interface.

- **Connection to Triangular Quantum Wells:** To balance this large energy difference and equate the Fermi levels, a massive wave of electrons rushes from the degenerate layer into the interface of the intrinsic substrate. This accumulation of negative charge forces the conduction band edge (E_C) of the intrinsic semiconductor to bend sharply downward until it dips beneath the Fermi level. This classical accumulation bending creates a sharp triangular potential profile at the interface. This mirrors the behavior of multi-material heterostructures, demonstrating how sharp changes in localized doping can be used to engineer quantum confinement channels without changing the underlying host crystal.

Lecture 3: The Physics of Semiconductor Quantum Dots

To understand why the transition from Gallium Arsenide (GaAs) to isotopically purified Silicon (^{28}Si) is such a massive leap for quantum computing, we have to look closely at the microscopic noise environments and solid-state physics governing these devices.

1. The Hyperfine Interaction and Nuclear Spin Bath

When an electron is confined within a semiconductor quantum dot, it is not perfectly isolated. The wave function of the electron extends over thousands of atoms in the host crystal lattice.

The “Spin Bath” Problem

The nuclei of these host atoms often possess their own intrinsic magnetic moments (nuclear spin). The interaction between the confined electron’s spin and the surrounding magnetic moments of the host nuclei is known as the **hyperfine interaction**.

- **GaAs (The Noisy Environment):** In Gallium Arsenide, every single natural isotope (^{69}Ga , ^{71}Ga , and ^{75}As) has a nuclear spin of $I = 3/2$. This creates a highly fluctuating, chaotic “nuclear spin bath.” This random magnetic field causes the electron spin to lose its quantum phase rapidly, leading to a very short free dephasing time (~ 10 ns).
 - **Natural Silicon (The Quieter Alternative):** Natural silicon consists mostly of ^{28}Si (92.2%) and ^{30}Si (3.1%), both of which have a nuclear spin of $I = 0$. However, it also contains about 4.7% of ^{29}Si , which has a nuclear spin of $I = 1/2$. Even this small percentage of magnetic nuclei limits dephasing to $\sim 1\text{--}10$ μs .
 - **Isotopically Purified ^{28}Si (The “Semiconductor Vacuum”):** By chemically purifying silicon to eliminate the ^{29}Si isotope, the host lattice becomes entirely composed of atoms with **zero nuclear spin**. Because there is no nuclear spin bath to interact with, the dominant mechanism for spin dephasing is wiped out, boosting coherence times into the hundreds of microseconds or even milliseconds.
-

2. Bandgap Engineering and Heterostructure Mechanics

To trap an electron in three dimensions to create a quantum dot, researchers use a combination of electrostatic gates and semiconductor heterostructures.

Bandgap Variation in $\text{Si}_{1-x}\text{Ge}_x$

By alloying Silicon with Germanium, the bandgap can be precisely tuned. The empirical formula provided by Professor Piccinini:

$$E_g(x) = 1.12 - 0.41x + 0.008x^2 \text{ eV}$$

shows that adding Germanium *narrows* the bandgap. For a typical 30% Germanium concentration ($x = 0.3$), the bandgap drops from 1.12 eV (pure Si) to 0.99 eV.

Type II Band Alignment

When a layer of pure Silicon is sandwiched between layers of Silicon-Germanium (SiGe/Si/SiGe), they form a **Type II heterostructure**. In a Type II alignment, the conduction band minimum and the valence band maximum occur in different material layers. Specifically, for a strained Si layer on a relaxed SiGe substrate:

- The **conduction band offset** forms a sharp potential energy well inside the Silicon layer.
- Electrons are strongly attracted to this well, restricting their movement vertically and forcing them into a two-dimensional electron gas (2DEG). Electrostatic surface gates are then used to slice this 2DEG into isolated, individual electron pockets—the quantum dots.

3. Strain Engineering and Effective Mass Modification

The physical dimensions of a quantum dot are tightly constrained by quantum mechanics. The spatial extent (size) of a quantum dot required to isolate a single electron is governed by the confinement energy, which is inversely proportional to the effective mass (m^*) of the charge carrier:

$$E_{\text{conf}} \propto \frac{1}{m^* \cdot L^2}$$

Where L is the characteristic size of the dot.

The Challenge of Pure Silicon

In its natural, unstrained state, Silicon has a relatively heavy longitudinal effective mass ($m_l^* \approx 0.98 m_0$) and six equivalent conduction band minima (valleys). This heavy mass means that to achieve distinct quantum energy levels, the physical size of the quantum dot must be incredibly small (often under 20 nm), which pushes the absolute limits of lithographic manufacturing.

The Role of Strain (LPCVD Deposition)

Using Low-Pressure Chemical Vapor Deposition (LPCVD) with gases like Silane (SiH_4) and Germane (GeH_4), a thin layer of pure Silicon is grown on top of a thicker, relaxed $\text{Si}_{0.7}\text{Ge}_{0.3}$ buffer layer.

Because Germanium atoms are larger than Silicon atoms, the underlying SiGe lattice has a larger lattice constant than pure Silicon. The thin Silicon layer is forced to stretch its atomic bonds horizontally to match the SiGe template, putting the Silicon under **tensile strain**.

This mechanical strain fundamentally alters the crystal symmetry:

- It splits the six-fold degenerate conduction band valleys, lowering two valleys (the Δ_2 valleys) relative to the other four.
- It modifies the curvature of the conduction band, effectively **reducing the effective mass** of the electrons moving in the plane of the layer.
- With a lighter effective mass, the electron wave function expands slightly. This means the physical dimensions (L) of the quantum dot can be made larger and more forgiving to fabricate using modern semiconductor lithography, while still maintaining excellent quantum confinement.

Heterostructure Band Engineering and Quantum Dot Confinement in Strained $\text{Si}_{1-x}\text{Ge}_x$ Systems

Gianluca Piccinini *gianluca.piccinini@polito.it* Politecnico di Torino – VLSI Group

1. Relaxation, Lattice Constants, and Misfit Strain Mechanics

Unlike lattice-matched material systems such as AlGaAs/GaAs, combining Silicon (Si) and Germanium (Ge) introduces distinct structural distortions due to a significant mismatch in their natural atomic spacings.

Bulk Material Relationships

The relaxed lattice constant (d) of a homogeneous $\text{Si}_{1-x}\text{Ge}_x$ alloy increases with the Germanium fraction x , interpolating between pure Silicon and pure Germanium:

- Pure Silicon ($x = 0$): $d_{\text{Si}} = 5.43 \text{ \AA}$
- Pure Germanium ($x = 1$): $d_{\text{Ge}} = 5.66 \text{ \AA}$

For an alloy composition of $x = 0.3$ (i.e., $\text{Si}_{0.7}\text{Ge}_{0.3}$), the relaxed lattice constant is approximately:

$$d(0.3) \approx 5.52 \text{ \AA}$$

The relaxed energy bandgap (E_g) and electron affinity ($q\chi$) for an unstrained $\text{Si}_{1-x}\text{Ge}_x$ layer vary non-linearly according to the empirical models:

$$E_g(x) = 1.12 - 0.41x + 0.008x^2 \implies E_g(0.3) = 0.99 \text{ eV}$$

$$q\chi(x) = 4.05 - 0.05x \implies q\chi(0.3) = 4.03 \text{ eV}$$

Epitaxial Strain Generation

When a thin semiconductor film is grown epitaxially on a much thicker crystal substrate, the atoms in the thin film are forced to alter their lateral spacing to match the substrate's native lattice constant. This locking mechanism generates mechanical stress within the active layer:

- **Tensile Stress (Structure A):** Occurs when a thin layer of pure Si ($d_{\text{Si}} = 5.43 \text{ \AA}$) is grown on top of a thick, relaxed $\text{Si}_{0.7}\text{Ge}_{0.3}$ virtual substrate ($d = 5.52 \text{ \AA}$). Because the substrate's atomic spacing is larger, the Silicon crystal lattice is pulled outward laterally. For a typical 1 GPa tensile load, the electronic properties shift to:

$$E_{g|\text{tensile}} \approx 1.096 \text{ eV}, \quad q\chi|\text{tensile} \approx 4.10 \text{ eV}$$

- **Compressive Stress (Structure B):** Occurs when a thin layer of pure Ge ($d_{\text{Ge}} = 5.66 \text{ \AA}$) is deposited on top of the same $\text{Si}_{0.7}\text{Ge}_{0.3}$ virtual substrate ($d = 5.52 \text{ \AA}$). Because the substrate's atomic spacing is smaller, the Germanium lattice is compressed inward laterally. Under a 1 GPa compressive load, its properties shift from its relaxed values ($E_{g0} = 0.66 \text{ eV}$, $q\chi_0 = 4.0 \text{ eV}$) to:

$$E_{g|\text{compress}} \approx 0.60 \text{ eV}, \quad q\chi|\text{compress} \approx 4.06 \text{ eV}$$

2. Band Structure Deformation and Effective Mass Alterations

Mechanical deformation alters the electronic band structures of group-IV semiconductors, breaking crystal symmetries and changing the curvature of the energy bands.

Conduction Band Flattening under Tensile Strain

In bulk, relaxed Silicon, the conduction band minimum consists of six degenerate valleys (Δ_6) along the $\langle 100 \rangle$ crystallographic directions. Applying biaxial tensile strain breaks this symmetry, splitting the valleys into two distinct sets:

1. Two lower-energy longitudinal valleys (Δ_2) oriented perpendicular to the interface plane.
2. Four higher-energy transverse valleys (Δ_4) oriented within the interface plane.

This splitting lowers the net conduction band energy edge and alters the parabolic curvature of the active energy states. According to the **Effective Mass Theorem**, the effective mass (m^*) of an electron inside a crystal depends directly on the curvature of its energy-momentum (E - k) dispersion relation:

$$m^* = \frac{\hbar^2}{\frac{\partial^2 E}{\partial k^2}}$$

Tensile strain sharpens the curvature of the active Δ_2 conduction band minima, causing a **decrease in effective mass** ($m^* \downarrow$). This reduction boosts the low-temperature carrier mobility (μ_n), which improves the switching speed of classical circuits and sharpens spin-orbital isolation in quantum dot gates:

$$\mu_n = \frac{q\tau}{m^*}$$

Valence Band Splitting under Compressive Strain

In relaxed substrates, the heavy-hole and light-hole bands are degenerate at the Γ -point maximum ($k = 0$). Biaxial compressive strain lifts this degeneracy by shifting the valence band upward, which reduces the total effective bandgap.

3. Structure A: Tensile Strained Silicon Quantum Wells

Structure A utilizes a thin, pure Silicon channel sandwiched between thick, relaxed cladding layers of $\text{Si}_{0.7}\text{Ge}_{0.3}$. The resulting tensile strain modifies the band edges to create a quantum confinement channel.

Mathematical Discontinuity Derivation

Using the Electron Affinity Rule and energy balance equations:

$$\text{Cladding Context (1): } q\chi_1 = 4.03 \text{ eV, } E_{g1} = 0.99 \text{ eV}$$

Strained Core (2): $q\chi_2 = 4.10$ eV, $E_{g2} = 1.096$ eV

- **Conduction Band Discontinuity (ΔE_C):**

$$\Delta E_C = q\chi_1 - q\chi_2 = 4.03 \text{ eV} - 4.10 \text{ eV} = -70 \text{ meV}$$

- **Bandgap Difference (ΔE_g):**

$$\Delta E_g = E_{g1} - E_{g2} = 0.99 \text{ eV} - 1.096 \text{ eV} = -106 \text{ meV}$$

- **Valence Band Discontinuity (ΔE_V):**

$$\Delta E_V = \Delta E_g + \Delta E_C = -106 \text{ meV} + (-70 \text{ meV}) = -176 \text{ meV}$$

Resulting Band Edge Alignment

Because ΔE_C is negative and ΔE_V is negative, both energy bands shift downward in the central Silicon region. This configuration forms a **Type-II (staggered) heterojunction alignment**:

Device Element / Feature	Value / Diagrammatic Behavior
Left Material Region	Relaxed Si _{0.7} Ge _{0.3} barrier layer
Center Material Region	Strained Silicon channel layer
Right Material Region	Relaxed Si _{0.7} Ge _{0.3} barrier layer
Vacuum Level Baseline (E_0)	Reference top horizontal line across all regions
Barrier Electron Affinity ($q\chi_1$)	4.03 eV (Both left and right relaxed layers)
Channel Electron Affinity ($q\chi_2$)	4.10 eV (Strained Silicon inner core)
Conduction Band Offset (ΔE_C)	Drop of -70 meV from E_{C1} down to E_{C2}
Valence Band Offset (ΔE_V)	Drop of -176 meV from E_{V1} down to E_{V2}
Fermi Level Alignment (E_F)	Uniform flat baseline crossing below the conduction bands
Carrier Confinement Zone	Quantum well for electrons formed within the Strained Si

Figure 4: Comprehensive energy band alignment parameters for a Strained Silicon heterostructure.

- **Confinement Core Impact:** The central Silicon layer forms a **70 meV deep quantum well for electrons**. This potential well traps negative charges within the pure Silicon core, creating the confinement channel needed to isolate single electron spins for spin qubits.

4. Structure B: Compressively Strained Germanium Quantum Wells

Structure B reverses the configuration by sandwiching a pure Germanium active core between the same $\text{Si}_{0.7}\text{Ge}_{0.3}$ virtual substrate layers, exposing the central core to compressive strain.

Mathematical Discontinuity Derivation

Using the matching parameters for the compressed Germanium state:

$$\text{Cladding Context (1): } q\chi_1 = 4.03 \text{ eV}, E_{g1} = 0.99 \text{ eV}$$

$$\text{Strained Core (2): } q\chi_2 = 4.06 \text{ eV}, E_{g2} = 0.60 \text{ eV}$$

- **Conduction Band Discontinuity (ΔE_C):**

$$\Delta E_C = q\chi_1 - q\chi_2 = 4.03 \text{ eV} - 4.06 \text{ eV} = -30 \text{ meV}$$

- **Bandgap Difference (ΔE_g):**

$$\Delta E_g = E_{g1} - E_{g2} = 0.99 \text{ eV} - 0.60 \text{ eV} = +390 \text{ meV}$$

- **Valence Band Discontinuity (ΔE_V):**

$$\Delta E_V = \Delta E_g + \Delta E_C = +390 \text{ meV} + (-30 \text{ meV}) = +360 \text{ meV}$$

Resulting Band Edge Alignment

The combination of a small negative conduction band step ($\Delta E_C = -30 \text{ meV}$) and a large positive valence band step ($\Delta E_V = +360 \text{ meV}$) confines both carriers within the central core. This configuration forms a **Type-I (straddling) heterojunction alignment**:

- **Confinement Core Impact:** While the conduction band displays a shallow 30 meV dip, the valence band edge climbs sharply to form a **360 meV deep quantum well for holes**. This deep valence well provides excellent confinement for hole-spin qubits, isolating the states from thermal leakage and high-frequency background interactions.

Device Layer / Feature	Value / Diagrammatic Behavior
Left Material Region	Relaxed Si _{0.7} Ge _{0.3} outer barrier layer
Center Material Region	Compressed Germanium channel core
Right Material Region	Relaxed Si _{0.7} Ge _{0.3} outer barrier layer
Vacuum Level Baseline (E_0)	Continuous flat top reference across all zones
Barrier Electron Affinity ($q\chi_1$)	4.03 eV (Both left and right relaxed layers)
Channel Electron Affinity ($q\chi_2$)	4.06 eV (Compressed Germanium core)
Conduction Band Offset (ΔE_C)	Drop of -30 meV from E_{C1} down to E_{C2}
Valence Band Offset (ΔE_V)	Large upward jump of $+360$ meV from E_{V1} up to E_{V2}
Fermi Level Alignment (E_F)	Uniform flat baseline crossing below the conduction bands
Carrier Confinement Zone	Deep quantum well for holes formed inside Compressed Ge

Figure 5: Comprehensive energy band alignment parameters for a Compressed Germanium heterostructure.

5. Analytical Practice Problems with Solutions

Problem 1: Discontinuity Sensitivity Analysis

An engineer modifies Structure A by shifting the cladding composition to Si_{0.8}Ge_{0.2}. The new relaxed cladding parameters are measured as $E_{g1} = 1.04$ eV and $q\chi_1 = 4.04$ eV. The properties of the strained Silicon core remain unchanged at $E_{g2} = 1.096$ eV and $q\chi_2 = 4.10$ eV.

Questions:

1. Calculate the new conduction band offset (ΔE_C) and valence band offset (ΔE_V).
2. Determine if this new configuration provides a shallower or deeper confinement well for electrons compared to the original Si_{0.7}Ge_{0.3} cladding.

Solution

1. Apply the Electron Affinity Rule and the valence band discontinuity relation:

$$\Delta E_C = q\chi_1 - q\chi_2 = 4.04 \text{ eV} - 4.10 \text{ eV} = -60 \text{ meV}$$

Next, calculate the new bandgap difference (ΔE_g):

$$\Delta E_g = E_{g1} - E_{g2} = 1.04 \text{ eV} - 1.096 \text{ eV} = -56 \text{ meV}$$

Substitute these values into the continuity equation to find ΔE_V :

$$\Delta E_V = \Delta E_g + \Delta E_C = -56 \text{ meV} + (-60 \text{ meV}) = -116 \text{ meV}$$

2. The original configuration with $\text{Si}_{0.7}\text{Ge}_{0.3}$ cladding produced an electron confinement well depth of $|\Delta E_C| = 70 \text{ meV}$. Modifying the cladding to $\text{Si}_{0.8}\text{Ge}_{0.2}$ reduces this depth to $|\Delta E_C| = 60 \text{ meV}$. Therefore, **the new configuration provides a shallower electron confinement well**, which reduces the electrostatic energy barrier holding the qubit.

Problem 2: Mechanical Strain Limit Check

A Germanium active core is epitaxially grown on a relaxed substrate. Due to structural constraints, the layer experiences a partial stress profile of 500 MPa, reducing the shift in its electronic properties by half relative to the relaxed state. Its intermediate parameters are $E_{g2} = 0.63 \text{ eV}$ and $q\chi_2 = 4.03 \text{ eV}$. The cladding layer remains standard $\text{Si}_{0.7}\text{Ge}_{0.3}$ ($E_{g1} = 0.99 \text{ eV}$, $q\chi_1 = 4.03 \text{ eV}$).

Questions:

1. Compute the band edge discontinuities ΔE_C and ΔE_V .
2. Analyze the resulting band alignment. Does the conduction band still step downward, or does it become perfectly flat across the heterojunction interface?

Solution

1. Calculate the band edge discontinuities:

$$\Delta E_C = q\chi_1 - q\chi_2 = 4.03 \text{ eV} - 4.03 \text{ eV} = 0 \text{ eV} = 0 \text{ meV}$$

Calculate the bandgap difference (ΔE_g):

$$\Delta E_g = E_{g1} - E_{g2} = 0.99 \text{ eV} - 0.63 \text{ eV} = +360 \text{ meV}$$

Substitute these values to find the valence band discontinuity (ΔE_V):

$$\Delta E_V = \Delta E_g + \Delta E_C = 360 \text{ meV} + 0 \text{ meV} = +360 \text{ meV}$$

2. Because $\Delta E_C = 0 \text{ meV}$, **the conduction band edge becomes perfectly flat across the heterojunction interface**, removing the step transition. The valence band retains a large step change of +360 meV, meaning the system provides robust spatial confinement for hole states while leaving conduction electrons free to move across the interface.

Lecture 5: The Physics and Material Challenges of MOS Qubits

Section 1: Deepening the Physical Concepts

1. The Multi-Dimensional Schrödinger Equation

To understand how Structure A confines a particle, we begin with the time-independent Schrödinger equation for a single particle with an effective mass m^* moving in a three-dimensional potential $U(x, y, z)$:

$$\left[-\frac{\hbar^2}{2m^*} \nabla^2 + U(x, y, z) \right] \psi(x, y, z) = E \psi(x, y, z)$$

Where $\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} + \frac{\partial^2}{\partial z^2}$. Because the structural potential is additive ($U = U_x + U_y + U_z$), the total Hamiltonian operator can be written as a sum of independent, 1D directional operators:

$$\hat{H} = \hat{H}_x + \hat{H}_y + \hat{H}_z$$

When we substitute the separation of variables ansatz $\psi(x, y, z) = \psi_x(x)\psi_y(y)\psi_z(z)$ into the equation and divide the entire expression by $\psi(x, y, z)$, the partial differential equation breaks down into three independent, ordinary differential equations. Each equation depends on only one spatial coordinate:

$$\left(-\frac{\hbar^2}{2m^*} \frac{d^2\psi_x}{dx^2} + U_x(x)\psi_x \right) \frac{1}{\psi_x} + \left(-\frac{\hbar^2}{2m^*} \frac{d^2\psi_y}{dy^2} + U_y(y)\psi_y \right) \frac{1}{\psi_y} + \left(-\frac{\hbar^2}{2m^*} \frac{d^2\psi_z}{dz^2} + U_z(z)\psi_z \right) \frac{1}{\psi_z} = E$$

For this equality to hold true for all possible independent values of x , y , and z , each bracketed term must equal a constant. These constants are the directional energies E_x , E_y , and E_z , proving that $E = E_x + E_y + E_z$.

2. Physical Reality of Boundary Conditions

In an ideal, infinitely deep potential well of width L , the potential energy jumps to infinity at the boundaries ($x = 0$ and $x = L$). Because a particle cannot possess infinite potential energy, the probability of finding it outside or exactly on the boundary must be zero ($\psi(0) = \psi(L) = 0$).

Inside the well, $U(x) = 0$, reducing the differential equation to a simple harmonic oscillator form:

$$\frac{d^2\psi_x}{dx^2} + k_x^2\psi_x = 0, \quad \text{where } k_x = \sqrt{\frac{2m^*E_x}{\hbar^2}}$$

The general solution is $\psi_x(x) = A \sin(k_x x) + B \cos(k_x x)$.

- Applying $\psi_x(0) = 0$ forces $B = 0$.
- Applying $\psi_x(L_x) = 0$ forces $\sin(k_x L_x) = 0$, which yields the quantization rule: $k_x L_x = m_x \pi$ for $m_x = 1, 2, 3, \dots$

Substituting this back into the definition of k_x yields the isolated spatial energies:

$$E_x = \frac{m_x^2 \pi^2 \hbar^2}{2m^* L_x^2}$$

Using the identity $\hbar = \frac{h}{2\pi}$, we can rewrite this in terms of Planck's constant (h):

$$E_x = \frac{m_x^2 \pi^2 \left(\frac{h}{2\pi}\right)^2}{2m^* L_x^2} = \frac{m_x^2 h^2}{8m^* L_x^2}$$

Section 2: Mathematical Derivations of the Exercises

Exercise 1: Energy Splitting Derivation (ΔE)

Problem Statement: Derive the exact subband splitting energy ΔE between the ground state and the first excited state for a perfectly cubical quantum dot ($L_x = L_y = L_z = L$).

Step-by-Step Solution:

1. Write out the general expression for total energy by summing all three dimensions:

$$E_{m_x, m_y, m_z} = \frac{h^2}{8m^*} \left(\frac{m_x^2}{L_x^2} + \frac{m_y^2}{L_y^2} + \frac{m_z^2}{L_z^2} \right)$$

2. Set $L_x = L_y = L_z = L$ to simplify the equation for a symmetrical cube:

$$E_{m_x, m_y, m_z} = \frac{h^2}{8m^* L^2} (m_x^2 + m_y^2 + m_z^2)$$

3. Calculate the Ground State Energy, which occurs at the lowest possible quantum numbers, $m_x = 1, m_y = 1, m_z = 1$:

$$E_{1,1,1} = \frac{h^2}{8m^* L^2} (1^2 + 1^2 + 1^2) = \frac{3h^2}{8m^* L^2}$$

4. Identify the First Excited State. This state requires incrementing one of the independent quantum numbers by one unit ($m_x = 1, m_y = 1, m_z = 2$):

$$E_{1,1,2} = \frac{h^2}{8m^*L^2} (1^2 + 1^2 + 2^2) = \frac{6h^2}{8m^*L^2}$$

5. Calculate the subband energy gap (ΔE):

$$\Delta E = E_{1,1,2} - E_{1,1,1} = \frac{6h^2}{8m^*L^2} - \frac{3h^2}{8m^*L^2} = \frac{3h^2}{8m^*L^2}$$

Exercise 2: Scaling Law Derivation (ΔE_{eV} vs L_{nm})

Problem Statement: Prove that the energy splitting equations can be consolidated into the practical engineering rule of thumb: $\Delta E = \frac{1.13}{(m^*/m_0) \cdot L_{\text{nm}}^2}$.

Constants and Base Units:

- Planck's Constant: $h = 6.626 \times 10^{-34} \text{ J} \cdot \text{s}$
- Rest Mass of an Electron: $m_0 = 9.11 \times 10^{-31} \text{ kg}$
- Electron Volt Conversion: $1 \text{ eV} = 1.602 \times 10^{-19} \text{ J}$
- Nanometer Scales: $L = L_{\text{nm}} \times 10^{-9} \text{ m}$

Step-by-Step Derivation:

1. Start with the raw SI unit equation for ΔE :

$$\Delta E_{\text{Joules}} = \frac{3h^2}{8m^*L^2}$$

2. Substitute the mass scaling factor $m^* = \left(\frac{m^*}{m_0}\right) m_0$ and the nanometer transformation for length L :

$$\Delta E_{\text{Joules}} = \frac{3 \cdot (6.626 \times 10^{-34})^2}{8 \cdot \left(\frac{m^*}{m_0}\right) (9.11 \times 10^{-31}) \cdot (L_{\text{nm}} \times 10^{-9})^2}$$

3. Isolate the numerical constants from the variable parameters:

$$\Delta E_{\text{Joules}} = \frac{3 \cdot (4.3904 \times 10^{-67})}{8 \cdot (9.11 \times 10^{-31}) \times 10^{-18}} \cdot \frac{1}{\left(\frac{m^*}{m_0}\right) \cdot L_{\text{nm}}^2}$$

$$\Delta E_{\text{Joules}} = \frac{1.3171 \times 10^{-66}}{7.288 \times 10^{-48}} \cdot \frac{1}{\left(\frac{m^*}{m_0}\right) \cdot L_{\text{nm}}^2} = 1.8072 \times 10^{-19} \cdot \frac{1}{\left(\frac{m^*}{m_0}\right) \cdot L_{\text{nm}}^2}$$

4. Convert the calculated energy from Joules into electron volts (eV) by dividing by the fundamental electronic charge (1.602×10^{-19}):

$$\Delta E_{\text{eV}} = \frac{1.8072 \times 10^{-19}}{1.602 \times 10^{-19}} \cdot \frac{1}{\left(\frac{m^*}{m_0}\right) \cdot L_{\text{nm}}^2}$$

$$\Delta E_{\text{eV}} = \frac{1.128}{\left(\frac{m^*}{m_0}\right) \cdot L_{\text{nm}}^2} \approx \frac{1.13}{\left(\frac{m^*}{m_0}\right) \cdot L_{\text{nm}}^2}$$

5. Rearrange this scaling law algebraically to solve for the maximum allowable nanometric dimension ($L_{\text{nm}}^{\text{MAX}}$) given a minimum energy threshold (ΔE_{min}):

$$L_{\text{nm}}^2 = \frac{1.13}{\left(\frac{m^*}{m_0}\right) \cdot \Delta E_{\text{min}}} \implies L_{\text{nm}}^{\text{MAX}} = \sqrt{\frac{1.13}{\left(\frac{m^*}{m_0}\right) \cdot \Delta E_{\text{min}}}}$$

Section 3: Material Implementation Exercises

To keep a qubit stable at an operating temperature of $T = 1$ K, thermal fluctuations must not accidentally excite the system. The operational design requirement dictates that the quantum energy gap must be at least ten times larger than the environmental thermal energy:

$$\Delta E_{\text{min}} > 10 \cdot k_B T$$

$$\Delta E_{\text{min}} = 10 \cdot (8.617 \times 10^{-5} \text{ eV/K} \cdot 1 \text{ K}) = 8.617 \times 10^{-4} \text{ eV} \approx 0.86 \text{ meV}$$

To ensure an operating margin safely above this limit against background electrical noise, Prof. Piccinini's notes establish a minimum threshold of:

$$\Delta E_{\text{min}} = 5 \text{ meV} = 5 \times 10^{-3} \text{ eV}$$

Exercise 3: Designing a Gallium Arsenide (GaAs) Qubit

Given Parameters:

- $\Delta E_{\text{min}} = 5 \times 10^{-3} \text{ eV}$
- Effective Mass Ratio ($\frac{m^*}{m_0}$) = 0.067

Calculation:

$$L_{\text{nm}}^{\text{MAX}} = \sqrt{\frac{1.13}{0.067 \cdot (5 \times 10^{-3})}}$$

$$L_{\text{nm}}^{\text{MAX}} = \sqrt{\frac{1.13}{3.35 \times 10^{-4}}} = \sqrt{3373.13} = 58.07 \text{ nm}$$

- **Engineering Constraint:** $L_{\text{GaAs}} \leq 60 \text{ nm}$
-

Exercise 4: Designing a Silicon (Si) Qubit

Given Parameters:

- $\Delta E_{\text{min}} = 5 \times 10^{-3} \text{ eV}$
- Effective Mass Ratio ($\frac{m^*}{m_0}$) = 0.26

Calculation:

$$L_{\text{nm}}^{\text{MAX}} = \sqrt{\frac{1.13}{0.26 \cdot (5 \times 10^{-3})}}$$

$$L_{\text{nm}}^{\text{MAX}} = \sqrt{\frac{1.13}{1.3 \times 10^{-3}}} = \sqrt{869.23} = 29.48 \text{ nm}$$

- **Engineering Constraint:** $L_{\text{Si}} \leq 30 \text{ nm}$

Mechanical Conclusion

Because electrons in Silicon behave as though they are heavier than those in Gallium Arsenide ($0.26 \cdot m_0$ vs $0.067 \cdot m_0$), their natural quantum wavelengths are significantly shorter.

Consequently, Silicon requires nearly **double the lithographic precision** ($\approx 30 \text{ nm}$ manufacturing structures vs $\approx 60 \text{ nm}$) to achieve identical quantum isolation characteristics. This sets the stage for the lithographic variability challenges outlined in Part 2 of the lecture notes.

1. The MOS Quantum Dot Architecture & Electrostatic Confinement

The cross-section on Page 9 outlines a device structurally identical to a nanoscale MOSFET but operated in a completely different regime.

Instead of driving a continuous current from Source to Drain, a MOS quantum dot uses localized electrostatic gating to isolate **individual electrons**.

- **The Potential Well:** Surface metallic or polysilicon gates are biased to create a local minimum in the conduction band edge (E_C). This forms a zero-dimensional potential well.

- **Energy Quantization:** Because the physical dimensions are nanometric (~ 50 nm), the continuous energy bands split into discrete, quantized electronic states ($E_0, E_1, \dots$).
- **Tunneling Regime:** The source and drain act as electron reservoirs. For single-electron operation, the electrochemical potential (μ) of the dot is precisely tuned via gate voltages to sit between the chemical potentials of the source (μ_S) and drain (μ_D), allowing only one electron at a time to occupy a localized state via quantum tunneling.

2. The Microscopic Nightmare: Interface Defects and Oxide Charges

Page 10 highlights the core material bottleneck of MOS qubits: **the Si/SiO₂ interface**. Unlike the pristine, atomically smooth interfaces of epitaxial heterostructures (like GaAs/AlGaAs or Si/SiGe), the thermal oxidation of Silicon to form SiO₂ is inherently amorphous and chemically imperfect.

```

[ Metal Gate ]
-----
      SiO2          --> Q_m  (Mobile Ionic Charge - e.g., Na+ migration)
                   --> Q_ot (Oxide-trapped Charge - broken bonds in bulk)
- - - - - --> Q_f  (Fixed-oxide Charge - unreacted Si near interface)
~~~~~ SiO_x ~~~~~ --> Q_it (Interface-trapped Charge - dangling bonds/Pb centers)
-----
[ Silicon ]

```

The Four Oxide Charges & Their Impact on Qubits:

1. **Interface-trapped Charge (Q_{it}):** These occur at the transition layer (SiO_x) due to structural defects like “dangling bonds” (specifically called P_b centers, where a surface Silicon atom is bonded to only three oxygen atoms). These traps create states *within* the Silicon bandgap. They can dynamically capture and emit electrons. When an electron traps/detraps near a qubit, it alters the local electrostatic potential, causing **charge noise** that rapidly dephases the qubit spin.
2. **Fixed-oxide Charge (Q_f):** Unreacted Silicon ions that remain near the interface after oxidation. They are immobile but create a permanent, random background Coulomb potential, distorting the intended shape of the quantum dot.
3. **Oxide-trapped Charge (Q_{ot}):** Defects inside the bulk oxide caused by ionizing radiation or high-field stress.
4. **Mobile Ionic Charge (Q_m):** Impurities like Na⁺ or K⁺ ions that drift under high electric fields, causing long-term calibration drift in the qubit’s operating frequencies.

3. Diagnostics via C-V (Capacitance-Voltage) Characterization

To engineer out these defects, researchers rely on Capacitance-Voltage (C-V) measurements of MOS capacitors, modeled using **Poisson-Boltzmann statistics**.

Poisson-Boltzmann Analysis (Page 12)

The semiconductor's internal potential (ϕ) as a function of depth (x) into the substrate is determined by solving the 1D Poisson equation:

$$\frac{d^2\phi}{dx^2} = -\frac{\rho(x)}{\epsilon_s}$$

Where the charge density $\rho(x)$ includes holes, electrons, and ionized acceptors/donors (N_a). Under thermal equilibrium, carrier distributions follow Boltzmann statistics:

$$\rho(\phi) = q \left(p_0 e^{-\frac{q\phi}{kT}} - n_0 e^{\frac{q\phi}{kT}} - N_a \right)$$

Integrating this equation yields the total semiconductor surface charge (Q_s) as a function of the surface potential (ϕ_s):

- **Accumulation** ($\phi_s < 0$): Majority carriers (holes) pile up at the interface exponentially: $\propto \exp\left(-\frac{q\phi_s}{2kT}\right)$.
- **Depletion** ($0 < \phi_s < 2\phi_F$): Holes are pushed away, leaving a fixed space-charge layer of ionized acceptors.
- **Inversion** ($\phi_s > 2\phi_F$): The surface potential is so high that minority carriers (electrons) are drawn to the interface exponentially: $\propto \exp\left(\frac{q\phi_s}{2kT}\right)$.

Reading Trap Density (N_{it}) from C-V Shifts (Page 13)

When interface traps (N_{it}) are present, they introduce an additional capacitance (C_{it}) in parallel with the semiconductor capacitance (C_s).

As the gate voltage (V_g) sweeps from accumulation to inversion, the fermi level sweeps through the bandgap, filling or emptying these interface states.

- **The “Stretch-Out” and Shift:** Because these traps must be charged/discharged, they resist changes in the surface potential. At low frequencies, this causes the C-V curve to stretch out and shift along the voltage axis.

- **Negative Bias Temperature Instability (NBTI):** The arrow on Page 13 points out that as trap density increases (often accelerated by stress or temperature), the voltage required to achieve inversion increases. For a qubit engineer, a high N_{it} means the C-V curve degrades, which translates directly to an unpredictable and noisy electrostatic environment for the quantum dot.

4. Scalability Barriers: Fabrication Variability & Control

Pages 14 and 15 detail why building a scalable quantum computer with millions of uniform MOS qubits is an immense engineering hurdle.

Fabrication Variability

In classical CMOS, a 1 nm variation in gate oxide thickness (t_{ox}) or gate width causes a minor variation in threshold voltage (ΔV_{th}) that can be easily averaged out or tolerated by digital logic noise margins. In quantum computing, a 1 nm geometry deviation fundamentally reshapes the quantum dot potential well. This changes:

- The **orbital energy spacing**, altering the required microwave frequencies used to manipulate the spin.
- The **tunnel coupling** to neighboring dots, which changes exponentially with distance.

Control, Readout, and Coupling

- **Spin-to-Charge Conversion:** Photons or magnetic fields directly manipulating a single electron spin generate signals too weak to detect directly. Instead, systems use spin-dependent tunneling (e.g., Pauli Spin Blockade) to convert the spin state (Up/Down) into a physical charge location (Left Dot/Right Dot).
- **SET/QPC Readout:** This spatial charge movement is then read out using an ultra-sensitive electrometer alongside the qubit, such as a **Single-Electron Transistor (SET)** or a **Quantum Point Contact (QPC)**. These sensors are highly susceptible to the same background charge noise described on Page 10.
- **Exchange Coupling:** Two-qubit logic gates rely on the Heisenberg exchange interaction (J), which depends on the physical overlap of the two electron wave functions:

$$H_{\text{exchange}} = J(t) \mathbf{S}_1 \cdot \mathbf{S}_2$$

Because this overlap is controlled via nanometer-scale gate voltages, even picometer-scale structural instabilities or minuscule electrical noise on the control


lines will cause gate errors, limiting the fidelity of the quantum processor.

Lecture 6: Qubit Fabrication & Device Architectures

By leveraging **FD-SOI (Fully Depleted Silicon-on-Insulator)** and **FinFET** technologies, researchers can use geometric confinement rather than complex electrostatic gating to isolate electrons, paving a realistic path toward mass-producible qubit arrays.

1. Planar Silicon Qubits: The Multi-Layer Gate Stack

The first section (Steps 1–5) outlines a classic, laboratory-grade planar qubit fabrication flow. It relies on a multi-layer overlapping gate architecture to create the tiny electrostatic potentials needed to trap a single electron.



The Multi-Step Process Demystified:

- **S1 (Polysilicon Channels):** Depositing and thinning a 40 nm polysilicon layer defines the primary physical channels where the quantum dots will reside.
- **S2 & S3 (Gate Definition & Isolation):** Atomic Layer Deposition (**ALD**) lays down a highly uniform 22 nm layer of Alumina (Al_2O_3). This acts as a high- κ dielectric layer. Aluminium (Al) is then patterned via lift-off to form **reservoir gates** (which supply electrons) and **plunger gates** (which shift the quantum dot's energy levels up and down).
- **S4 & S5 (The Overlapping Strategy):** To position control gates incredibly close to one another without short-circuiting them, the aluminum gates are oxidized using an oxygen plasma, creating a thin, insulating native oxide (AlO_x) skin. **Tunnel gates** are then deposited over this native oxide. This overlapping design allows researchers to tune the tunneling barriers between the dots and the reservoirs with nanometer-scale precision.

2. Advanced Planar Design: FD-SOI Qubit Technology

While overlapping gate stacks work well for few-qubit experiments, they are a nightmare to scale up to millions of qubits due to routing density and parasitic capacitances. Advanced industrial designs instead use **Ultra-Thin Body and Buried Oxide (UTBB) FD-SOI** technology.

Single vs. Double Quantum Dot Operating Conditions

Instead of forcing an electron into a tiny pocket using three or four overlapping surface gates, FD-SOI uses an ultra-thin silicon channel ($T_{ox} \approx 1$ nm) naturally bounded by an insulating oxide block on the bottom (**BOX** - Buried Oxide) and Shallow Trench Isolation (**STI**) on the sides.

Looking closely at the operating parameters provided in the lecture notes, we can map out how these voltage configurations isolate charge carriers:

Parameter	Single Quantum Dot Mode	Double Quantum Dot Mode	Physical Mechanism
V_P (Plunger Gate)	+1 V	+1 V	Attracts electrons into the channel, populating the quantum dot.
V_{TS} (Tunnel Source)	-0.2 V	-0.3 V	Pushes back against the source reservoir, thinning the electron channel to create a sharp tunnel barrier.
V_{TD} (Tunnel Drain)	-0.2 V	-0.3 V	Creates the corresponding tunnel barrier on the drain side.
V_T (Inter-dot Tunnel)	<i>N/A</i>	-0.3 V	Only in Double Dots: Slices the single large channel in half, pinching it in the middle to create two distinct, coupled quantum dots.

The Power of UTBB (Ultra-Thin Body and BOX)

By thinning the silicon body down to a few nanometers, the gate wraps tighter around the channel electrostatically. This yields:

- **Enhanced Gate Coupling & Control:** Lower operating voltages are needed to manipulate electron states.
 - **Reduced Parasitic Capacitance:** Eliminates the capacitive coupling found in overlapping gate architectures, increasing operational speed.
 - **The Back Gate Secret Weapon:** In UTBB, the entire silicon substrate beneath the Buried Oxide acts as a **back gate**. By applying a voltage to this back gate, engineers can globally shift the quantum dot potentials, dynamically altering their energy levels, tailoring carrier confinement, and fine-tuning electron occupancy without changing the top-gate configurations.
-

3. 3D Technologies: FinFET Qubit Architectures

When planar FD-SOI hits its physical scaling limits, the semiconductor industry pivots to **3D FinFET** structures—and qubit engineers are following suit.

Why FinFETs are Ideal for Large Qubit Arrays

In a FinFET architecture, the silicon channel is etched into a vertical, ultra-thin “fin.” The control gate is then draped over three sides of this 3D structure.

1. **Natural Quantum Dots:** In a classical FinFET operated at cryogenic temperatures (~ 100 mK), the corners of the silicon fin naturally create local potential minima. Electrons become tightly confined along the top and side walls of the fin, creating a “natural” quantum dot under the gate without requiring complex surface electrode geometries.
2. **Section along the Fin Direction:** When arrayed along the length of the fin, a sequence of parallel gates running perpendicular to the silicon fin can create a highly dense, linear register of quantum dots. One gate acts as a qubit plunger, the next acts as an inter-dot tunnel barrier, and so on.
3. **Massive Scalability & Industry Alignment:** Because commercial chip foundries have spent over a decade perfecting FinFET manufacturing, quantum computing architectures based on FinFETs can directly tap into existing commercial fabrication lines. This drastically reduces variability, improves yield, and opens a clear path toward integrating classical control electronics (Cryo-CMOS) right alongside the qubit core array.

This section of Professor Piccinini’s course explores the theoretical foundation required to control and read out spin qubits: **Electrostatic Modeling and Spin Physics**.

To move from a physical nanostructure to a working computer, we must mathematically model how electrons behave under strict electrostatic confinement (**Coulomb Blockade**), map out multi-electron configurations via **Stability Diagrams**, and utilize quantum mechanics (**Zeeman Effect** and **Pauli Spin Blockade**) for qubit initialization and readout.

1. The Single Quantum Dot & Coulomb Blockade Mechanics

When a quantum dot is structurally confined, adding a single electron to it requires a discrete amount of energy due to mutual electrostatic repulsion. This phenomenon is known as the **Coulomb Blockade**.

The Constant Interaction (Capacitive) Model

The total electrostatic energy $U(N)$ of a quantum dot containing N electrons is modeled by treating the gates, source, and drain as localized capacitors connected to a central island with total capacitance $C_\Sigma = C_S + C_D + C_G$.

The electrostatic potential energy is defined by:

$$U(N) = \frac{(eN - C_G V_G)^2}{2C_\Sigma}$$

The energy required to add the N -th electron to the dot is the **Chemical Potential** (μ_N):

$$\mu_N = U(N) - U(N-1) = \left(N - \frac{1}{2}\right) E_C - \frac{eC_G V_G}{C_\Sigma}$$

Where $E_C = \frac{e^2}{C_\Sigma}$ is the **Coulomb Charging Energy**.

- **The Blockade Condition:** If the electrochemical potentials of the source (μ_S) and drain (μ_D) flank an energy gap where no discrete μ_N level sits, an electron cannot enter or leave the dot. Current flow is completely blocked.
 - **Coulomb Diamonds:** Plotting differential conductance (dI/dV_{SD}) as a function of Source-Drain Bias (V_{SD}) and Gate Voltage (V_G) maps out diamond-shaped regions. Inside these diamonds, the electron number N is perfectly fixed and stable. Electron transport only occurs at the vertices where diamonds meet (Coulomb oscillation peaks).
-

2. SET Readout for a Quantum Dot

Reading out the subtle state of a single electron spin directly is impossible due to weak magnetic signals. Instead, we use a **Single-Electron Transistor (SET)**

as an ultra-sensitive electrometer located immediately adjacent to the quantum dot.

The Measurement Technique

1. A constant, small bias voltage (V_{SD}) is applied across the SET's source and drain electrodes.
 2. The SET gate voltage is tuned to the steep slope of a **Coulomb blockade oscillation** peak. In this regime, the current passing through the SET (I_{SET}) is extraordinarily sensitive to any nearby electric fields.
 3. When an electron tunnels into or out of the *neighboring quantum dot*, its physical movement shifts the local electrostatic environment. This acts like an additional virtual gate bias on the SET, instantly shifting the Coulomb oscillation peak and causing a dramatic, measurable jump in I_{SET} .
-

3. The Double Quantum Dot (DQD) Capacitive Model & Stability Diagram

When two quantum dots are placed close together, they form a Double Quantum Dot (DQD) system, modeled as two capacitive islands interconnected by an inter-dot capacitance (C_m) and cross-coupled to each other's gates (V_{G1}, V_{G2}).

The DQD Stability Diagram

Instead of single diamonds, a DQD system maps its charge configurations onto a two-dimensional plot of V_{G1} vs. V_{G2} .

- **Honeycomb Lattice:** The boundaries form a characteristic honeycomb network. Each cell corresponds to a stable charge configuration denoted as (N_1, N_2) , where N_1 is the electron count in Dot 1 and N_2 is the count in Dot 2.
 - **Triple Points:** The vertices where three charge states meet are called triple points. At these precise coordinates, electrochemical potentials align, enabling sequential tunneling of electrons through both dots, producing a current flow.
-

4. Spin Manipulation: The Zeeman Effect

To transform an isolated electron into a qubit, a static magnetic field (B) must be applied. This activates the **Zeeman Effect**, which breaks the energy degeneracy of the electron's spin-up ($|\uparrow\rangle$) and spin-down ($|\downarrow\rangle$) states.

The energy separation (ΔE) between these two states is linear with respect to the magnetic field strength:

$$\Delta E = g\mu_B B$$


Where:

- g is the Landé g-factor (in Silicon, $g \approx 2$, whereas in GaAs, $g \approx -0.44$).
- μ_B is the Bohr magneton ($\approx 57.88 \mu\text{eV/T}$).

As highlighted by Professor Piccinini's notes, at a standard laboratory field strength of $B = 1 \text{ T}$, the energy splitting equals $\Delta E = 116 \mu\text{eV}$. This energy separation defines the operational **Larmor frequency** ($\nu = \Delta E/h \approx 28 \text{ GHz}$) used to manipulate the qubit using microwave pulses.

5. DQD Operation Sequence & Pauli Spin Blockade (PSB)

The notes conclude with the foundational sequence used to initialize, manipulate, and read spin states in a DQD system using **Pauli Spin Blockade (PSB)**.



Step 1: Initialization

Using electrostatic gating, the DQD is loaded with exactly two electrons—one in each dot, putting the system into a stable (1,1) charge configuration.

Step 2: Tuning Energy Levels

The detuning parameter (ϵ), which measures the relative energy difference between the electrochemical levels of Dot 1 and Dot 2, is swept using gate voltages. The levels are adjusted to favor a transition from the (1,1) state to a (0,2) state (where both electrons occupy Dot 2).

Step 3: Spin-Selective Tunneling (The Blockade)

According to the Pauli Exclusion Principle, two electrons can only occupy the exact same spatial orbital (the ground state of Dot 2) if their spins are anti-parallel, forming a **Singlet state**: $S = \frac{1}{\sqrt{2}}(|\uparrow\downarrow\rangle - |\downarrow\uparrow\rangle)$.

- If the two electrons form a **Triplet state** (e.g., both spin-up: $|\uparrow\uparrow\rangle$), the second electron is forbidden from tunneling into the ground state of Dot 2.

It would be forced into a much higher, energetically unfavorable excited orbital.

- Consequently, if the system is in a Triplet spin configuration, electron tunneling is completely **suppressed (blocked)**.

Step 4: Manipulation and Readout

This blockade translates a quantum mechanical property (Spin) directly into a macroscopic physical state (Charge placement).

- **Readout:** The adjacent SET senses whether the system successfully transitioned to the $(0, 2)$ charge state or remained blocked in the $(1, 1)$ state.
- If the SET registers a charge shift corresponding to $(0, 2)$, the qubit was a **Singlet**.
- If no charge moves and the system stays stuck at $(1, 1)$, the qubit was a **Triplet**. This completes the high-fidelity spin-to-charge conversion loop required for quantum readout.

These lecture notes from Professor Gianluca Piccinini provide a rigorous mathematical and circuit-level derivation of the **Constant Interaction Model** for a Single Quantum Dot (SQD) and a Double Quantum Dot (DQD).

The following sections unpack the math, circuit analysis, and underlying physics behind these equations to show exactly how the stability diagrams and readout mechanisms are formed.

1. Electrostatic Derivation of the Dot Potential (V_{DOT})

To find the electrostatic potential inside the quantum dot (V_{DOT}) before any extra electron tunnels in, we model the system as a node connected to three capacitors: the gate capacitor (C_G), the source capacitor (C_S), and the drain capacitor (C_D).

Using the principle of **superposition** on the circuit node (Pages 19–20):

1. **Case 1** ($V_{DS} = 0$): Treat the drain and source as ground. The gate voltage V_{GS} couples to the dot via a capacitive voltage divider:

$$V'_{\text{DOT}} = V_{GS} \cdot \frac{C_G}{C_G + C_S + C_D} = V_{GS} \frac{C_G}{C_\Sigma}$$

2. **Case 2** ($V_{GS} = 0$): Treat the gate and source as ground. The drain bias V_{DS} couples to the dot via the drain capacitance:

$$V''_{\text{DOT}} = V_{DS} \cdot \frac{C_D}{C_G + C_S + C_D} = V_{DS} \frac{C_D}{C_\Sigma}$$

Adding these together gives the total electrostatic potential of the dot:

$$V_{\text{DOT}} = V_{GS} \frac{C_G}{C_\Sigma} + V_{DS} \frac{C_D}{C_\Sigma}$$

Where $C_\Sigma = C_S + C_D + C_G$ is the **total dot capacitance**.

The electrostatic energy contribution to the dot is $U_{\text{DOT}} = -qV_{\text{DOT}}$. Using the realistic values provided on Pages 17–18 ($C_G \approx 2$ aF, $C_S \approx 1$ aF, $C_D \approx 1$ aF, yielding $C_\Sigma = 4$ aF):

$$U_{\text{DOT}} = -qV_{GS} \left(\frac{2}{4} \right) - qV_{DS} \left(\frac{1}{4} \right) = -0.5qV_{GS} - 0.25qV_{DS}$$

2. The Charging Energy Component

When a single electron tunnels onto the dot, it injects a discrete charge $\Delta q = -q$. This alters the dot potential abruptly (Page 21):

$$\Delta V_{\text{DOT}} = \frac{-q}{C_\Sigma}$$

Because the electron is fighting the negative potential of electrons already trapped on the dot, the physical potential energy increases by the **Coulomb Charging Energy** (E_C):

$$\Delta U_{\text{DOT}} = -q(\Delta V_{\text{DOT}}) = \frac{q^2}{C_\Sigma}$$

With $C_\Sigma = 4$ aF, this charging energy calculation yields:

$$E_C = \frac{(1.6 \times 10^{-19} \text{ C})^2}{4 \times 10^{-18} \text{ F}} = 6.4 \times 10^{-21} \text{ J} \approx 40 \text{ meV}$$

This means that at room temperature ($k_B T \approx 26$ meV), thermal fluctuations could easily overcome this barrier, causing electrons to hop randomly on and off the dot. To stabilize individual charge states for qubits, the system must be cooled to cryogenic temperatures where $k_B T \ll 40$ meV.

3. Deriving Tunneling Conditions (T_1 and T_2)

An electron will only tunnel if there is an available energy state for it to occupy. Pages 22–27 derive the boundary constraints for electron movement.

Condition T_1 : Source $\rightarrow$ Dot Tunneling

For an electron to jump from the source reservoir into the dot, the total electrochemical potential of the dot with $N + 1$ electrons must be lower than the Fermi level of the source (E_{FS}):

$$E_m + U_{\text{DOT}} + E_C(N - N_0) + E_C < E_{FS}$$

Substituting U_{DOT} and $E_C = \frac{q^2}{C_\Sigma}$:

$$E_m - qV_{GS}\frac{C_G}{C_\Sigma} - qV_{DS}\frac{C_D}{C_\Sigma} + \frac{q^2}{C_\Sigma}(N - N_0 + 1) < E_{FS}$$

Isolating the voltage parameters when $N - N_0 = 0$ and setting the boundary threshold to an equality gives the lines that form the Coulomb diamonds:

- **When $V_{DS} = 0$:** The required gate voltage to open this tunneling channel is:

$$V_{GS} = \frac{C_\Sigma \Delta E_m}{qC_G} + \frac{q}{C_G}$$

- **When $V_{GS} = 0$:** The required drain voltage is:

$$V_{DS} = \frac{C_\Sigma \Delta E_m}{qC_D} + \frac{q}{C_D}$$

As N increases ($N - N_0 = 1, 2, \dots$), these voltage thresholds shift by discrete periods of $\frac{q}{C_G}$ and $\frac{q}{C_D}$ respectively, setting the repetitive spacing of the diamond pattern along the axes.

Condition T_2 : Dot $\rightarrow$ Drain Tunneling

For an electron to exit the dot into the lower-energy drain reservoir, the energy of the state it leaves behind must be higher than the drain's Fermi level ($E_{FD} = E_{FS} - qV_{DS}$):

$$E_m - qV_{GS}\frac{C_G}{C_\Sigma} - qV_{DS}\frac{C_D}{C_\Sigma} + \frac{q^2}{C_\Sigma}(N - N_0 - 1) > E_{FS} - qV_{DS}$$

Rearranging this inequality shows that the T_2 boundary lines have a positive slope because V_{DS} appears on both sides of the inequality, altering the net electrostatic acceleration.

4. The Coulomb Diamond Diagram Explained

When you plot the boundaries where T_1 and T_2 turn from closed to open as a function of V_{DS} vs. V_{GS} , they form the **Coulomb Diamond Diagram** (Page 28).

- **Inside the Diamonds:** Both T_1 and T_2 conditions are violated. No electrons can enter or exit. The number of electrons on the dot is fixed as a stable integer state (e.g., state ① or ②).
 - **The Diamond Perimeter:** These lines represent single-electron tunneling thresholds. Conduction spikes up exactly along these lines.
 - **Resonant Transport:** Outside the diamonds, sequential tunneling is unlocked (T_1 followed by T_2), allowing a steady, measurable current to flow through the device.
-

5. High-Sensitivity Qubit Readout via Cross-Capacitive Shifting

Page 29 illustrates how this exact diamond framework is utilized to read out a nearby qubit dot using an adjacent Single-Electron Transistor (SET).

- **Capacitive Cross-Coupling (C_C):** The SET dot and the Qubit dot are placed close together. When the Qubit dot's electron state changes (e.g., moving an electron out of the dot during a spin-to-charge conversion sequence), it acts as an abrupt change in voltage (ΔV_{qubit}) seen by the SET.
 - **Diamond Shift:** This extra voltage shifts the entire Coulomb diamond pattern of the SET along its gate voltage axis (the yellow diamonds vs. the black diamonds on Page 29).
 - **The Sweet Spot ($\left. \frac{\Delta G}{\Delta V_{GS}} \right|_{\text{MAX}}$):** By biasing the SET's drain voltage (V_{DS_bias}) and sweeping its gate to the sharpest, steepest edge of a conductance peak, the SET current becomes exceptionally vulnerable to small changes. A tiny shift in the qubit's charge state moves the peak directly over the operating point, triggering an immediate drop or spike in measured conductance (G).
-

6. Double Quantum Dot (DQD) Stability Contours

Finally, Pages 30–32 expand this concept to a two-dot system where the mutual inter-dot capacitance is initially treated as zero ($C_{m0} = 0$).

Without capacitive cross-talk between the dots, their electrical behaviors are completely independent:

- **Dot 1 Confinement:** Depends purely on its own gate voltage V_{G1} . The source-tunneling boundaries (T_S) form perfectly straight, parallel **vertical lines** spaced by $\frac{q}{C_{G1}}$.
- **Dot 2 Confinement:** Depends purely on its gate voltage V_{G2} . The drain-tunneling boundaries (T_D) form perfectly straight, parallel **horizontal lines** spaced by $\frac{q}{C_{G2}}$.

When these two orthogonal sets of lines overlap, they create a **pristine rectangular grid** (Page 32). Each block in this grid defines a highly stable, dual-electron state index configuration (N_1, N_2) . For example, the state $(1, 1)$ represents exactly one electron isolated in Dot 1 and one electron isolated in Dot 2—the foundational starting configuration for initialization and two-qubit exchange gates.